

BAYESIAN NONPARAMETRIC GLOTTAL INVERSE FILTERING

by

MARNIX VAN SOOM

A DISSERTATION SUBMITTED IN PARTIAL FULFILLMENT OF
THE REQUIREMENTS FOR THE DEGREE OF
DOCTOR OF PHILOSOPHY

in

THE FACULTY OF GRADUATE AND POSTDOCTORAL STUDIES
(Physics)

THE UNIVERSITY OF BRITISH COLUMBIA
(Vancouver)

February 2026

© Marnix Van Soom, 2026

The following individuals certify that they have read, and recommend to the Faculty of Graduate and Postdoctoral Studies for acceptance, the dissertation entitled:

Bayesian Nonparametric Glottal Inverse Filtering

submitted by **Marnix Van Soom** in partial fulfilment of the requirements for the degree of Doctor of Philosophy in Physics.

Examining Committee:

John Doe, Research Scientist, Physical Sciences Division, TRIUMF
Research Co-supervisor

Jane Doe, Professor, Department of Chemistry, UBC
Academic Co-supervisor

Abstract

A reliable and accurate white-box model of speech is still something of a holy grail in phonetics, forensics, and clinical medicine. The difficulty is that fitting such models to microphone recordings is a blind identification problem, known to practitioners as glottal inverse filtering, which requires inferring both source signal and filter simultaneously from a single waveform. Because the problem is severely underdetermined, we must lean hard on prior information to constrain the set of mathematically possible solutions: they must adhere to what we already know to be plausible or true. I will show how nonparametric Gaussian processes let us bake such prior knowledge directly into the model by (1) deriving a nonstationary kernel that can model the glottal cycle, (2) specializing that kernel into a surrogate model learned from simulated source signals via interdomain features, and (3) proposing a new and general framework for specifying arbitrary prior information about the filter; and how this improves reliability and accuracy at inference time.

Lay summary

In the past decade, we have made enormous strides in recognizing and synthesizing speech using large “black-box” models trained on massive datasets. Today, these models have learnt to speak and understand major world languages such as English, Chinese, and Spanish with impressive accuracy. Remarkably, they remain largely agnostic to the physical mechanisms underlying speech, relying instead on large-scale pattern recognition and vast quantities of data.

In contrast, “white-box” models grounded in accumulated scientific knowledge do focus explicitly on these physical mechanisms. While they are not remotely competitive with black-box models in speech recognition or synthesis, they do offer the valuable ability to extract biologically meaningful information about the speaker from the signal itself.

This is important. A growing body of evidence shows that quantitative features of voice and speech can act as digital biomarkers of neurological and clinical conditions. In Parkinson’s disease, for example, changes in speech and language have been reported to precede diagnosis by as much as a decade, and acoustic measures such as voice quality, articulation rate, and prosody have shown promise for early detection and monitoring. Similarly, alterations in specific acoustic parameters have been linked to mood disorders, suggesting that speech analysis could complement clinical assessment of depression.

Nevertheless, white-box models often lack the flexibility of black-box models needed to deal with real-world speech. I therefore propose a “grey-box” approach that attempts to combine the best of both worlds. By integrating physical insight with data-driven modeling, I derive a class of grey-box models that allows for far more accurate analysis and decomposition of speech into separate contributions from the vocal folds and the vocal tract. These models are also capable of indicating how certain they are of their measurements, while running close to real-time speed on modern hardware.

This was achieved by combining new theoretical insights in glottal flow and vocal tract models with customized machine learning models. I hope that this work can contribute to a more reliable and holistic use of speech analysis in clinical practice, forensic work, and the scientific study of vocalization as a whole.

Preface

Write your preface here. This must include a description of your role in the work as presented, and the role of any collaborators if applicable. Additionally, include a statement on the use of Generative AI tools in producing this work.

Table of contents

Abstract	iii
Lay summary	iv
Preface	v
Table of contents	vi
Acknowledgments	viii
1 Overview	1
1.1 Introduction	1
1.2 Problem statement	2
1.3 Source-filter theory	4
1.4 Contributions	9
2 Infinite kernel linear prediction	14
2.1 Linear prediction	14
2.2 Kernel linear prediction	15
2.3 Infinite kernel linear prediction	16
2.4 Inference	17
2.5 Woodbury-Cholesky reformulation	19
2.6 Prior reparameterizations	21
2.7 Summary	22
3 Priors for autoregressive filters	24
3.1 The isotropic prior	24
3.2 The Monahan prior	26
3.3 Spectral shaping priors	28
3.4 Modeling or expanding?	32
3.5 Summary	34
4 From parametric to nonparametric glottal flow models	35
4.1 The Liljencrants-Fant model	36
4.2 Classic piecewise polynomial models	40
4.3 Parametric piecewise polynomial models	41
4.4 Nonparametric piecewise polynomial models	47
4.5 Summary	53
5 The periodic and quasiperiodic arc cosine kernel	54
5.1 The periodic arc cosine kernel	55
5.2 The quasiperiodic arc cosine kernel	60
5.3 Gated variants	62
5.4 Registration with EGIFA	64
5.5 Calibration on EGIFA	66
5.6 Summary	70
6 Surrogate glottal flow models	71
6.1 Refinement	71
6.2 Phase warping	72

6.3	Generative models of voiced speech	74
6.4	Summary	75
7	Bayesian nonparametric glottal inverse filtering	76
7.1	Recap of all the ingredients	76
7.2	Current state-of-the-art GIF algorithms	76
7.3	Learning priors from synthetic data	77
7.4	EGIFA	78
7.5	Evaluation on EGIFA	78
7.6	Evaluation on VTRFormants	79
7.7	Summary	80
8	Conclusion	81
8.1	Future work	82
	Datasets	84
1.	EGIFA	84
2.	APLAWD	88
3.	CommonVoice	92
4.	VTRFormants	93
	Algorithms	96
1.	LBGID	96
2.	PRISM	100
3.	t-PRISM	106
4.	q-GPLVM	112
	Bibliography	123
	Appendices	135
A.	Regularization and priors	135
B.	Parametric, nonparametric, semiparametric?	137
C.	JAX implementation of the Liljencrants-Fant model	139
D.	More comments on the nonparametric glottal flow model	144
E.	On derivative time scales	146
F.	Predicting period-to-period jitter	148
G.	Gauge invariance in glottal inverse filtering evaluation	151
H.	Related contributions	155

Acknowledgments

Write your acknowledgments here. This section is optional, feel free to remove it if you do not wish to include it. Nonetheless, this is the place to thank professional colleagues and people who have given you the most support during the course of your work.

1 Overview

We introduce Bayesian nonparametric glottal inverse filtering (BNGIF) and provide a high-level overview of the main concepts involved, followed by a summary of the contributions made in this thesis.

1.1 Introduction

In this thesis we present a new approach to the problem of *glottal inverse filtering* (GIF) called *Bayesian nonparametric glottal inverse filtering* (BNGIF). At the heart of this new approach lies a Bayesian nonparametric *generative model of voiced speech* that reflects general constraints on the voice production process. Though conceived for GIF, the generative model can be conditioned on both raw speech data and high level speech features, which makes it useful for other related tasks such as augmenting or denoising speech data.

Speech can be thought of as the output of a linear system in which an upstream source signal (for example, the rapid opening and closing of the vocal folds) is modulated by a downstream filter (for example, the oral cavity and lips). This linear approximation to speech production, called source-filter theory (Fant, 1960), still underlies most of what is known today as acoustic phonetics (Maurer, 2016), and both source and filter are comparably well understood.

The GIF problem refers to source-filter separation of speech; that is, simultaneously inferring both the glottal source signal and (some of) the vocal tract filter characteristics from the speech signal alone (Figure 1). In the BNGIF approach we systematically take advantage of known constraints on the voice production process to home in on scientifically plausible solutions to the GIF problem. That is the main theme of this thesis.

GIF is an important and difficult problem in acoustic phonetics (Auvinen et al., 2014; Kadiri et al., 2021) and various approaches have been developed over the last 60 years. Paraphrasing Alku (2011), applications mainly come in three broad categories: (i) fundamental research about voice production (for example, quantifying voice quality), (ii) medical applications (such as analysis of pathological voices), and (iii) speech technology applications (for example, voice modification). To this we would add (iv) recent forensic and biometric applications such as evidence-based speaker identification (Bonastre et al., 2015).

Our approach to the GIF problem is mostly relevant to the first category (i) above because we try to model the process of scientific inference – make assumptions, model the data, fit the data, evaluate the model, reiterate – as a whole rather than ‘just’ modeling the data alone (MacKay, 2005, Ch. 28). It has been shown repeatedly that the unique way to do this is

FIGURE HERE

Figure 1 | **Overview of source-filter separation.** The speech waveform (left) is modeled as a glottal source signal (top right) passing through a series of vocal tract filters (bottom left). The task is to find such plausible source-filter pairs (colored samples on the right) that together produce a given speech waveform.

the probabilistic or *Bayesian* framework (Knuth & Skilling, 2012). Within this mathematical framework our generative model of voiced speech ultimately expresses nothing more than our assumptions about voice production in a logically coherent way. Because we insist on modeling scientific inference itself the accuracy of those assumptions will translate directly to the accuracy of our measurements, providing a much-needed means to calibrate uncertainty for use in (i) fundamental research about voice production (Kent & Vorperian, 2018). This is the main benefit of and motivation for using the Bayesian framework for BNGIF, despite its considerable computational expense compared to standard GIF algorithms which are based on traditional signal processing methods.

Another important benefit is the standardized form of Bayesian generative models: posterior inference is mechanical once the prior and likelihood are specified, which allows for a wide variety of generic algorithms to fit the data in true plug-and-play fashion (Murphy, 2022). This standardized form also ensures that generative models are easily modifiable and extendable for more specific applications such as (ii-iv) mentioned above.¹

Another distinguishing feature of BNGIF is the *nonparametric* nature of our generative model of speech. More precisely, the model contains both nonparametric and parametric components. Figure 1 shows these two neatly divided components on the right hand side: the glottal source signal (top) is modeled nonparametrically with a Gaussian process (Rasmussen & Williams, 2006), while the vocal tract filter (bottom) is expanded into a familiar rational transfer function that is parameterized by an overall gain factor and a vector of poles and zeros.

Nonparametric in this context does not refer to ‘distribution-free’ methods used in descriptive statistics or nonparametric spectral analysis (Stoica & Moses, 2005); rather it means that the glottal source signal is modeled directly through an infinite-dimensional distribution over functions whose qualitative behavior is specified by a just few global properties like roughness and wiggleness. This is convenient because this is exactly the kind of broad information we possess about the glottal source signal, unlike the vocal tract filter which in source-filter theory is naturally specified in terms of resonances of the vocal tract (Stevens, 2000).

1.2 Problem statement

The central problem of this thesis is boxed in Figure 2.

¹For example it is expected that our model of human voice production can be useful for mammal vocal communication research (Taylor & Reby, 2010) by simply tweaking some of its prior assumptions. An example of extending the model to downstream tasks is forensic phonetics, since “formant features [...] are widely accepted features used in forensic acoustic-phonetic speaker verification” (Becker et al., 2008, p. 1) and “the Bayesian decision framework [...] is now the standard paradigm for forensic speaker comparison” (Bonastre et al., 2015, p. 1). This Bayesian decision framework depends on likelihood ratios (Harrison, 2013) that can be calculated with the same computational tools used in this thesis.

Glottal inverse filtering (GIF)

Assume that the source-filter model of speech production (Fant, 1960) holds such that the speech signal at time $t \in \mathbb{R}$ is described by convolving three real-valued functions:

$$s(t) = u(t) * h(t) * r(t). \quad (1)$$

From left to right $s(t)$ is the speech pressure waveform, $u(t)$ is the glottal flow waveform, $h(t)$ is the impulse response of the vocal tract and $r(t)$ is the impulse response of the radiation characteristic, assumed known. Source-filter theory asserts that $s(t)$ can be thought of as the result of the source signal $u(t)$ passing through the filters $h(t)$ and $r(t)$.

The data consists of N noisy samples $\mathbf{d} = d_{1:N}$ of the speech pressure waveform $s(t)$ at times $\mathbf{t} = t_{1:N}$. We assume additive errors $\mathbf{e} = e_{1:N}$ such that

$$\mathbf{d} = \mathbf{s} + \mathbf{e} \implies d_n = s(t_n) + e_n \quad (n = 1 \dots N). \quad (2)$$

The error components e_n are modeled as white noise with expected power σ_n^2 which means that $e_n \sim \text{Normal}(0, \sigma_n^2)$ for $n = 1 \dots N$.

The data $(\mathbf{t}, \mathbf{d})$ is possibly augmented with noisy statistics of $s(t)$ such as a list of estimated glottal closing instants (GCIs).

The GIF problem is then to *find plausible estimates of the source $u(t)$ and the filter $h(t)$* that are consistent with the given data.

Figure 2 | **The GIF problem** which we attack in this thesis.

Source-filter theory as expressed by equation (1) is the most important assumption in this thesis because it provides a definite handle on the GIF problem. Even just stating the problem would be difficult without it. In particular, note the four functions (or waveforms) in (1); each of these play an important role in BNGIF. Before reviewing them in Section 1.3, we first consider how the *plausible estimates* mentioned in the statement above translate within the Bayesian framework adopted in this thesis.

Like most inverse problems, the main difficulty with GIF lies in the fact that it is vastly under-determined: even in the noiseless case ($\mathbf{d} = \mathbf{s}$), many solutions (that is, pairs of functions $u(t)$ and $h(t)$) are compatible with given data $\mathbf{d}$ (Landau, 1983). It thus appears that next to the data $\mathbf{d}$ itself we need additional criteria to choose between possible solutions of (1);, and this is where the idea of regularization comes in.

Appendix A discusses how in the Bayesian framework regularization is imposed by assigning priors to the unknown functions $u(t)$ and $h(t)$ in such a way that these priors express known constraints on the voice production process. The main idea is that a probability in the Bayesian framework describes a state of knowledge, not necessarily related to an observed frequency in the real world, so we can conveniently use prior probabilities to express constraints on the voice production process.

Even with that regularization in place, we will see that the problem is still underdetermined in the sense that we cannot in general expect to find unique solutions: the best we can hope for are *plausible estimates in the form of posterior samples of $u(t)$ and $h(t)$* . These are plausible in the sense that they are consistent with both the speech data $\mathbf{d}$ and whatever prior information about voice production we have expressed through the priors for $u(t)$ and $h(t)$.

1.3 Source-filter theory

FIGURE HERE

Figure 3 | Schematical representation of the source-filter model (1) with u, h, r, s and their Fourier transforms.

Source-filter theory² models the speech production process as a standard *linear time-invariant* (LTI) system (Antsaklis & Michel, 2006) which is essentially a linear approximation to the underlying nonlinear dynamics of voice production that is roughly valid up to frequencies of 4 or 5 kHz (Doval et al., 2006). The LTI system (1) used in this thesis is illustrated in Figure 3. We proceed with a brief sketch of the important concepts involved.

The *input* to the system is the source $u(t)$. The system has two *transfer functions* which determine how the amplitudes and phase of the frequencies present in the input $\tilde{u}(f)$ are modified or filtered by the system to produce the *output*. That frequency decomposition

$$\tilde{u}(f) = \mathcal{F}[u(t)](f) = \int_{-\infty}^{\infty} u(t) \exp\{-i2\pi ft\} dt \quad (3)$$

is obtained by the Fourier transform $\mathcal{F}$ of $u(t)$. Note that the system's transfer function is assumed to be independent from its input: source-filter theory assumes source and filter to be decoupled.

The vocal tract transfer function $\tilde{h}(f)$ describes the resonances and antiresonances of the vocal tract (comprising the throat, mouth and nasal cavity), while the radiation characteristic $\tilde{r}(f)$ describes the lip radiation effect. These transfer functions can be described equivalently in the time domain by the *impulse responses*

$$\begin{aligned} h(t) &= \mathcal{F}^{-1}[\tilde{h}(f)](t) = \int_{-\infty}^{\infty} \tilde{h}(f) \exp\{i2\pi ft\} df \\ r(t) &= \mathcal{F}^{-1}[\tilde{r}(f)](t) = \int_{-\infty}^{\infty} \tilde{r}(f) \exp\{i2\pi ft\} df \end{aligned} \quad (4)$$

²Currently the dominant paradigm in acoustic phonetics (Maurer, 2016), the source-filter theory of voice production is commonly attributed to (Fant, 1960) but appears to date back to the early 19th century (Chen, 2016). Its strength lies in its extremely useful and compact description of the acoustics of phonation, summarized by (Fant, 1960, p. 15) as “the speech wave is the response of the vocal tract filter systems to one or more sound sources.” A modern perspective on source-filter theory is given by Svec et al. (2021).

using the inverse Fourier transform $\mathcal{F}^{-1}$. Unlike $u(t)$ and $h(t)$, $r(t)$ is assumed fixed and known: the lip radiation effect is modeled in the GIF literature as a standard 6 dB/octave high pass filter (Stevens, 2000, p. 128) which has transfer function $\tilde{r}(f) = i2\pi x$. This corresponds to differentiation in the time domain as (4) yields $r(t) = \delta'(t)$.³

The *output* of the system is the speech pressure signal $s(t)$

$$\begin{aligned} s(t) &= u(t) * h(t) * r(t) && \text{(time domain)} \\ \Rightarrow \tilde{s}(f) &= \tilde{u}(f) \tilde{h}(f) \tilde{r}(f) && \text{(frequency domain)} \end{aligned} \quad (5)$$

assumed to be measured in the free field (that is, sufficiently far from the speaker's head) (Stevens, 2000). Filtering in LTI systems happens through multiplication in the frequency domain (5), which is equivalent to convolution in the time domain (1).

The meaning of these equations is best understood in the frequency domain. During speech production, the source drives the air in the vocal tract such that it vibrates with some initial distribution over frequencies $\tilde{u}(f)$. Excited, the vocal tract filter $\tilde{h}(f)$ reacts: frequencies in $\tilde{u}(f)$ close to its resonance frequencies are emphasized,⁴ while frequencies in $\tilde{u}(f)$ close to its antiresonance frequencies are (often brutally) suppressed. The result of this spectral imprinting is then radiated from the lips into the free field. This radiation effect is contained within $\tilde{r}(f)$: high frequencies are emphasized, essentially because the lips act like small (passive) tweeters (Schroeder, 1999). We are able to *hear* the spectral imprint of the speaker's vocal tract on the initial energy distribution (Klatt, 1986), because it is what we use to tell the difference between different speech sounds, and ultimately to understand the speaker's message.

Since $r(t)$ is assumed known, we can further simplify (1) by substituting $r(t) = \delta'(t)$:

$$s(t) = u(t) * h(t) * \delta'(t) = [u(t) * \delta'(t)] * h(t) = u(t) * [h(t) * \delta'(t)]. \quad (6)$$

We can differentiate either $u(t)$ or $h(t)$ in (1) because convolution is associative; we choose to differentiate $u(t)$ as is customary in the literature (Doval et al., 2006). This finally yields

$s(t) = u'(t) * h(t)$

(7)

where $u'(t)$ is the *glottal flow derivative* (DGF) which now acts as the source driving the single filter in (7). The original source $u(t)$ can be recovered through integration:

³Convolution of a test function $f(f)$ with the derivative of the Dirac delta function is equivalent to differentiation: $f(t) * \delta'(t) = f'(t)$. This can be justified by distribution theory (Lighthill, 1958).

⁴The vocal tract filter is *passive*, meaning that $|\tilde{h}(f)| < 1$ for all frequencies x (Flanagan, 1965). The resonances of the vocal tract do not actively amplify frequencies; frequencies close to the resonance frequencies are just less dampened relative to neighbouring frequency bands.

$$u(t) = \int_{-\infty}^t u'(\tau) d\tau. \quad (8)$$

Equation (7) is the basis for the nonparametric generative model of $s(t)$ which we use to separate speech data [noisy samples of $s(t)$] into source $u'(t)$ [equivalently $u(t)$ via (8)] and filter $h(t)$ [equivalently $\tilde{h}(f)$ via (4)].

With the general picture laid out, we now discuss the $u(t)$, $h(t)$ and $s(t)$ waveforms in further detail, and present a high-level overview of the proposed priors for each of them. Section 1.2 argued that effective regularization of the GIF problem requires realistic priors on $u(t)$ and $h(t)$ and we can show some of their realistic properties already by just sampling them. Note that these priors are nothing but generative models of $u(t)$, $h(t)$ and $s(t)$ – they are ‘priors’ because we can generate samples from these models according to a probability density that represents our prior information about the speech production process.

1.3.1 The glottal flow waveform $u(t)$

In source-filter theory the glottal flow $u(t)$ is the source component which provides the initial acoustic energy that is subsequently modified by the ‘downstream’ filter component (comprised of the vocal tract and the lips). This definition is quite general and the actual physical mechanism underlying $u(t)$ depends on the type of phonation employed.

In the GIF literature the type of phonation assumed is *voiced speech*, during which the vocal folds vibrate at the fundamental frequency F_0 while modulating a stream of air expelled from the lungs. The $u(t)$ waveform measures the flow rate (volume velocity, in ml/sec) of this airflow which looks like a quasiperiodic string of pulses, each pulse representing a “puff of air” passing through the glottis (Schroeder, 1999). The quasiperiodicity of $u(t)$ directly transfers to the voiced speech signal $s(t)$ because the vocal articulators move much more slowly than the vibrating vocal folds.

The most widely used model for the DGF $u'(t)$ in voiced speech is the Liljencrants-Fant model of Fant et al. (1985). We use this in Chapter 4 as the basis for a parametric prior for $u(t)$ denoted formally as

$$u(t) \sim \pi_{\text{LF}}(u(t)) \quad (9)$$

This parametric prior serves as the basis for a more expressive and nonparametric Gaussian process prior derived in Chapter 5:

$$u(t) \sim \pi_{\text{GP}}(u(t)) \quad (10)$$

The GP underlying this prior has a Matérn kernel that is calibrated to π_{LF} using Bayesian transfer learning (Xuan et al., 2021).

FIGURE HERE

Figure 4 | DGF and GF samples from π_{LF} and π_{GP} .

Samples from π_{LF} and π_{GP} are shown in Figure 4. Due to its nonparametric nature the latter is capable of representing a wide variety of $u(t)$ curves both in terms of roughness and overall pulse shape. Therefore, the nonparametric π_{GP} prior can be expected to yield a more realistic model of the glottal flow in real voiced speech compared to the parametric π_{LF} prior.

1.3.2 The impulse response of the vocal tract $h(t)$

Source-filter theory asserts that the physical resonances and antiresonances of the vocal tract show up as local maxima and minima in the envelope of the power spectrum $|\tilde{h}(f)|^2$ of the vocal tract transfer function. These extrema are known as *formants* and *antiformants*, respectively, and their amplitudes, frequencies and bandwidths give a compact description of the resonant and damping (‘antiresonant’) behavior of the vocal tract. During speech production, we manipulate these (anti)formant frequencies on the fly by skillfully changing the vocal tract configuration (such as rounding the lips, moving the tongue or closing the mouth) in just the right way as to produce the desired vowel or consonant. Humans really are master musicians of the vocal instrument.

The standard LTI assumption and canonical approach in source-filter theory is that $\tilde{h}(f)$ is a *rational transfer function* whose *poles* and *zeros* are identified with the aforementioned formants and antiformants, respectively (Stevens, 2000). We adopt this assumption in this thesis but without the one-to-one correspondence between (formants and pole pairs) and (antiformants and zero pairs). Instead, a single formant may be modeled by one or more poles while some ‘spectrum shaping’ poles need not correspond to formants at all; and similarly for antiformants and zeros. For this reason we prefer to think of $\tilde{h}(f)$ as a parametric rational *expansion* of the true underlying vocal tract transfer function. The parameters controlling this Padé expansion are its order, an overall gain factor and a vector of poles and zeros. All of these can be inferred from speech data.

A (strictly proper) rational transfer function $\tilde{h}(f)$ in the frequency domain implies through (4) that the impulse response $h(t)$ in the time domain consists of a real-valued superposition of exponentially decaying sinusoids. This real-valued superposition is more convenient mathematically for our purposes than a complex rational function, so in Chapter 3 we express our two priors for $h(t)$ in the time domain.⁵

FIGURE HERE

Figure 5 | Samples from $h(t)$ prior in time and spectral domain.

The first one is a parametric prior over pole-zero expansions of the true underlying vocal tract transfer function of a given order:

$$h(t) \sim \pi_{\text{PZ}(h(t))} \quad (11)$$

⁵This might seem counterintuitive at first, but it is actually a common approach in Bayesian spectrum analysis (Bretthorst, 1988; Turner & Sahani, 2014).

Likewise, the second prior parametrizes all-pole expansions, which, lacking zeros, are a special case of pole-zero expansions.

$$h(t) \sim \pi_{\text{AP}}(h(t)) \quad (12)$$

Samples from π_{PZ} and π_{AP} both in the time domain and the frequency domain are shown in Figure Figure 5. The qualitative differences between their power spectra in Figure Figure 5 are due to the fact that pole-zero expansions are more expressive than all-pole expansions. That is, samples from π_{AP} are more constrained because they are essentially samples from π_{PZ} for which the zeros are set up to cancel out exactly.

Without zeros π_{AP} cannot easily model antiformants of the vocal tract, which for example occur with nasal consonants [now sing]. This is often used as an argument against the use of all-pole transfer functions in speech processing (Mehta et al., 2012). Its samples also have an expected spectral tilt that is too steep to be realistic. In contrast π_{PZ} can model a more flexible class of spectra with a wide range of (anti)resonant behavior and possible spectral tilts, and is to be preferred to π_{AP} on these grounds, but at the cost of a roughly doubled amount of parameters.

Nevertheless, in practice all-pole expansions such as the ones represented by π_{AP} are generally thought to be sufficiently expressive to represent the vocal tract filter (Flanagan, 1965; Fulop, 2011) and have been computationally the preferred choice since the very beginning of digital acoustic phonetics (Atal & Hanauer, 1971). Next to theoretical grounds (Titze, 2000) this is mainly because they admit an efficient representation known as linear prediction (Markel & Gray, 1976), which is the traditional workhorse of acoustic analysis.

1.3.3 The speech pressure waveform $s(t)$

FIGURE HERE

Figure 6 | Samples from π_{VS} .

The speech pressure waveform

$$s(t) = u'(t) * h(t) \quad (13)$$

is the output of the LTI system that describes speech production according to source-filter theory. The priors for $u(t)$ and $h(t)$ introduced in ((10), (11), (12)) induce a prior for $s(t)$ through (7), which is derived analytically in Chapter 2. This prior is the aforementioned nonparametric generative model of voiced speech that underlies BNGIF:

$$s(t) \sim \pi_{\text{VS}}(s(t)) \quad (14)$$

Samples from π_{VS} are shown in Figure Figure 6. The quasiperiodicity exhibited by these waveforms is necessary for speech with a voiced sound and is induced by the quasiperiodicity of samples from π_{GP} as can be seen in the right panel of Figure Figure 4.

The π_{VS} prior is nonparametric because it is simply another GP whose behavior is governed by a learnable kernel that describes the intricate covariance structure of voiced speech. Accordingly, the samples shown in Figure 6 are nothing but draws from a high-dimensional multivariate normal distribution. BNGIF infers the source $u'(t)$ [equivalently $u(t)$ via (8)] and filter $h(t)$ [equivalently $\tilde{h}(f)$ via (4)] from (7) by learning the kernel that best describes the given speech data. This is a challenging but ultimately standard GP inference task that can be done in a variety of ways (Frigola, 2015; Simpson et al., 2020).

1.4 Contributions

The relatively long history of research into GIF naturally produced several lines along which the problem may be approached, but, following two recent reviews (Drugman et al., 2019a; Kadiri et al., 2021), it is possible to separate these approaches loosely into two broad classes: inverse filtering and joint source-filter optimization.⁶ Before reviewing the specific contributions made by BNGIF to the GIF literature, and what they are hoped to accomplish at what cost, we first situate the approach taken in this thesis within these two classes.

Inverse filtering Miller (1959) founded the basis for the first and larger class of GIF methods, which holds most of the popular GIF methods used today such as closed phase analysis (Wong et al., 1979) and iterative adaptive inverse filtering (Alku, 1992). The inverse filtering approach aims to recover the glottal source signal $u(t)$ by extracting from the speech data $\mathbf{d}$ an estimate of the vocal tract filter $\hat{h}(f)$ and then digitally filtering $\mathbf{d}$ by $1/\hat{h}(f)$, the inverse of that (hopefully nonzero) estimate. These methods are based on efficient digital signal processing algorithms. Most of them rely on LP analysis and differ mainly in how the all-pole vocal tract transfer function $\tilde{h}(f)$ is estimated (Kadiri et al., 2021). No explicit parametric form of $u(t)$ is assumed because $u(t)$ is recovered by applying the digital filter $1/\hat{h}(f)$ to a vector of speech samples $\mathbf{d}$.

Joint source-filter optimization Milenkovic (1986) proposed a different approach for the second class of GIF methods, a logical consequence of the increased computing power available in the 1980s. With joint source-filter optimization methods the fit of a parametric model to the speech data $\mathbf{d}$ is optimized to estimate both source $u(t)$ and filter $\tilde{h}(f)$ simultaneously (as opposed to consecutively as in the first class). These methods involve the nonconvex optimization of an objective function to determine the parameters θ of some assumed functional form of the DGF $u'(t) = f(t; \theta)$, typically an approximation of the LF model (Alzamendi & Schlotthauer, 2017; Degottex, 2010; Fu & Murphy, 2006; Schleusing, 2012). This optimization is computationally costly compared to the inverse filtering methods in the first class, and a variety of tricks are used to reduce runtime and improve robustness, such as two-stage optimization (Fu & Murphy, 2006) and regressing θ to a single parameter (Degottex, 2010).

⁶For the sake of the argument we ignore a here smaller third class which uses mixed-phase decomposition methods (Degottex, 2010; Drugman et al., 2019a; Kadiri et al., 2021) such as complex cepstrum decomposition (CCD) (Drugman et al., 2011).

Although we are not aware of any study that compares the performance of these two approaches systematically, it is possible to make a few general statements. Inverse filtering methods are generally efficient and nonparametric, while joint source-filter optimization methods have the distinct advantage of cleanly separating $u(t)$ and $\tilde{h}(f)$ from the outset. In doing so they avoid having to derive $u(t)$ directly from a filter estimate $\hat{\tilde{h}}(f)$ that inevitably contains pronounced effects from the former (Auvinen et al., 2014), to which Schroeder (1999, p. 94) refers as a “suspiciously circular sounding ‘bootstrap’ method.”

So, where does that leave BNGIF? The main motivation for and innovation of BNGIF is that it mixes something of both classes: it is a joint source-filter optimization method capable of the flexibility of inverse filtering methods through its nonparametric model for $u(t)$. In addition, the Bayesian framework adopted by BNGIF handles uncertainty in a straightforward way through sampling the posterior, which can be seen as generalization of the nonconvex optimization that characterizes (and plagues) the class of joint source-filter optimization methods it belongs to. In contrast, inverse filtering methods pay for their efficiency by having no principled way of dealing with uncertainty, other than operationalising it as input variability.

1.4.1 Contributions and merits of BNGIF

Nonparametric glottal flow model BNGIF utilizes a nonparametric glottal flow model. That is, the unknown function $u(t)$ is modeled nonparametrically with the GP prior (10);, calibrated to the widely used LF model for $u'(t)$. As explained above, this is intended to strike a balance between the nonparametric flexibility of inverse filtering methods and the parametric control of source-filter optimization methods.

Conversely, BNGIF also avoids two particular shortcomings of currently used GIF methods. Inverse filtering methods typically model (a part of) $u(t)$ with only a few poles, either implicitly (such as the CP method mentioned above) or explicitly (such as IAIF, also mentioned above), which has since long been known to be “incapable of representing realistic physiologic glottal volume velocity or area waveforms” (Deller, 1983, p. 62). Source-filter optimization methods, on the other hand, use analytical models for the glottal flow derivative $u'(t)$ (Doval et al., 2006) which are convenient but necessarily “limited in their ability to capture the behavior of the glottal source in natural speech” (Kadiri et al., 2021, p. 1926).

Uncertainty quantification It is commonly agreed that uncertainty analysis and quantification is important for fundamental scientific research and forensic or biometric applications of acoustic phonetics, for example when characterizing quasi-vowels of great apes (Ekström et al., 2023) or calculating the evidence for a proposition in the context of forensic speaker comparison (Bonastre et al., 2015). However, in the specific case of GIF, “a difficult blind separation problem since neither the vocal tract response nor the glottal flow contribution are actually observable” (Degottex et al., 2014, p. 962), we would argue that uncertainty quantification takes on a nearly essential quality, if only as a rough measure of the GIF algorithm’s trust in its own output.

As the reader probably anticipates by now, the natural framework for uncertainty quantification is Bayesian probability theory. Section 1.1 touched upon the reason for this, which is plain and simple: Bayesian theory prescribes how to reason consistently and quantitatively under uncertainty (that is, how to conduct inference mathematically), and it does so uniquely, as first demonstrated by (Cox, 1946). Uncertainty analysis is automatically incorporated in a Bayesian approach by sampling the posterior for the given problem, and that is of course what BNGIF does,⁷ as explained in Section 1.2 and illustrated in Figure 1. Uncertainty quantification of various GIF-derived statistics such as the open quotient (OQ) and first formant frequency (F_1) then reduces to calculating the empirical mean $\pm$ standard deviation of these statistics directly from posterior samples of $u(t)$ and $\tilde{h}(f)$.

Despite their appeal for uncertainty quantification, Bayesian approaches to GIF are rare in the literature, and next to BNGIF we know of only three other probabilistic GIF methods: (Auvinen et al., 2014; Rao & Ghosh, 2018; Wang et al., 2016). (Bleyer et al., 2017) confirm that the first Bayesian inversion approach to GIF is described in (Auvinen et al., 2014), who call it MCMC-GIF. Surprisingly, these methods all produce point estimates of the glottal flow and vocal tract filter and do not seem to use posterior samples for uncertainty analysis.

Adaptive GCI estimation Contrary to the sliding windows approach used by conventional LP analysis for formant estimation, GIF methods are *pitch-synchronous* in nature and require knowledge of the glottal closure instants (GCIs) in order to work properly (Drugman et al., 2019b). This means that the user has to run an independent GCI detection algorithm first before GIF methods can be used.⁸ A prominent (and according to (Chien et al., 2017), unique) exception to this is the aforementioned IAIF algorithm (Alku, 1992), which implements heuristic GCI estimation as part of its iterative structure.

On the other hand, BNGIF infers GCIs alongside and simultaneously with the source $u(t)$ and filter $\tilde{h}(f)$, which allows newly discovered information (during inference) about either one of these to constrain the others. This adaptive strategy helps to ‘bootstrap’ the GCIs from the speech signal, since GCIs are instrumental in recovering $u(t)$, but ultimately need to be derived from (and are defined in terms of) the latter. Another adaptive aspect of BNGIF’s built-in GCI estimation is that noisy GCI estimates may *optionally* be supplied by the user a priori; BNGIF will then use this information to speed up inference, and yield refined estimates as part of its output. These adaptive abilities are quite unique in the GIF literature as far as we are aware.

Modeling inertia and nonstationarity The inertia of the vocal apparatus and articulators imposes constraints on the speed at which the speech signal $s(t)$ can be produced and altered. It is often the case that $s(t)$ appears relatively ‘stable’ or stationary during short time intervals of about 10 to 50 msec, an observation which underlies the ubiquitous utility and adoption

⁷Or rather: tries to do, as sampling the posterior is intractable for most real-world problems, and efficient approximations must be used instead.

⁸Note that robust GCI detection “is a very difficult task in practice” (Alzamendi & Schlotthauer, 2017, p. 14) for which many algorithms have been proposed (Drugman et al., 2019b, Sec. 3.3).

of the short-time Fourier transform (STFT) in acoustic phonetics (Little, 2011). In the case of voiced speech, the providence of GIF, these inertial constraints produce correlations across *dozens* of pitch periods. The source $u(t)$ and filter $\tilde{h}(f)$ constituents of voiced speech are both inherently nonstationary yet strongly correlated at the level of neighbouring pitch periods. We can actually *hear* this fact: pitch and timbre are perceptual renditions of these statistical regularities (Schnupp et al., 2011).

BNGIF can model nonstationary voiced speech which exhibits these inertial correlations, and takes advantage of the latter to speed up inference (since any correlation means more predictability, which often leads to faster exploration of plausible hypotheses). Figures Figure 4 and Figure 6 illustrate the local inertia and nonstationarity of $u(t)$ and $s(t)$ respectively. But BNGIF also takes into account ‘filter correlations’; that is, the constrained variation of $\tilde{h}(f)$ from pitch period to pitch period. A pleasant consequence of this approach is probabilistic (anti)formant tracking, which we discuss just below.

The approach of BNGIF contrasts with many other GIF methods, which work one pitch period at a time and do not model inertial correlations between the pitch periods explicitly. For example, none of the algorithms considered in a recent review by (Chien et al., 2017) of the most important and representative inverse filtering methods use inter-pitch period correlations, and we are not aware of any joint source-filter optimization methods exploiting these correlations either.

Probabilistic (anti)formant tracking Formant tracking refers to estimating from speech data the (smooth) trajectories along which formants move over time, which in turn reflect the movement of the vocal articulators during phonation. Because of their efficiency as low-dimensional ‘information bearers’, formant tracks are highly important in acoustic phonetics and many algorithms exist for measuring them, mostly based on LP analysis (Van Soom & de Boer, 2020). These algorithms usually operate on a sliding window approach, estimating formants within each window and then using a post-processing step to join these point estimates into smooth trajectories (Whalen et al., 2022). A prominent example of this strategy is the widely used program (Boersma, 2001), whose default tracking function implements the Viterbi algorithm (Viterbi, 1967) to post-process the formant candidates into smooth trajectories.

In the case of BNGIF, formant tracking happens during inference and is achieved by exploiting the inertial correlations of the filter $\tilde{h}(f)$, as discussed above. That is, pole pairs are subjected to inter-pitch period correlations, and this takes care of ‘trajectory smoothing’ automatically while avoiding an ad-hoc post-processing step.

The same is true for the zero pairs: antiformant tracking is implicitly enabled in the case of a pole-zero expansion of the rational transfer function $\tilde{h}(f)$. We are not aware of other GIF methods with this ability, since GIF methods almost invariably assume all-pole transfer functions (Alku, 2011; Bleyer et al., 2017; Kadirri et al., 2021), even though spectral zeros occur

during the open phase of the glottis due to coupling to the subglottal cavity (Ananthapadmanabha & Fant, 1982; Plumpe et al., 1999).

There are a handful of formant tracking algorithms, however, that implement the same principle – that is, smoothing of formant tracks during inference rather than as a post-processing step – of which the Kalman-based autoregressive moving average (Mehta et al., 2012) method is a relatively recent example. BNGIF and KARMA also have in common the ability to optionally use noisy a priori estimates of the formant trajectories to speed up inference.

Avoiding harmonic attraction Higher values of the fundamental frequency F_0 pose challenges for formant tracking algorithms based on LP analysis due to a phenomenon (Whalen et al., 2022) call “harmonic attraction.” This refers to the fact that, as a result of spectral smoothing, the peaks of the estimated envelope obtained through LP are highly biased toward the harmonic partials of a speech spectrum (Yoshii & Goto, 2013), such that the estimated formant frequencies $\mathbf{F}$ tend to align with multiples of F_0 . For higher values of F_0 , this alignment causes the systematic correlations between F_0 and $\mathbf{F}$ known as F_0 bias. In the case that $F_0 \approx F_1$, known as $F_0 - F_1$ congruence (Kent & Vorperian, 2018), harmonic attraction can be especially challenging, since LP analysis, as a spectral smoothing device, lacks a conceptual model for separating the two – in contrast to, for example, the auditory processing regions of our brains (Schnupp et al., 2011).

Next to formant tracking, harmonic attraction equally affects GIF methods based on inverse filtering since the F_0 -biased formant estimates result in insufficient cancelation of the true vocal tract resonances, which in turn distorts inference of the glottal flow $u(t)$ (Auvinen et al., 2014). This problem is well known, and there is extensive literature on the degradation of inverse filtering methods under increased values of F_0 . In fact, ‘increased values of F_0 ’ often means ‘not a low-pitched male voice’ (Airaksinen et al., 2014), meaning that harmonic attraction renders a large majority of cases potentially problematic for inverse filtering methods! We also note that the $F_0 - F_1$ congruence mentioned above is especially problematic for GIF in the case of close rounded vowels, which remains “a major factor that limits the applicability of inverse filtering algorithms to accurate glottal flow estimation from continuous speech” (Chien et al., 2017, p. 17).

BNGIF is a joint-source filter optimization method and thereby automatically avoids harmonic attraction and the problems associated with $F_0 - F_1$ congruence. In addition, the probabilistic approach of BNGIF (as opposed to most GIF methods) is known to improve performance and reduced sensitivity to prior estimation of GCIs (Auvinen et al., 2014; Bleyer et al., 2017).

2 Infinite kernel linear prediction

The inference engine at the heart of BNGIF is a nonparametric Bayesian extension of linear prediction due to Yoshii & Goto (2013). The paper, presented over a decade ago at ICASSP 2013, has attracted fewer citations than its conceptual depth warrants — which, for our purposes, is fortunate, because it provides exactly the principled Bayesian foundation we need to replace the DSP heuristics of traditional glottal inverse filtering with a framework of testable probabilistic hypotheses. This chapter develops the model in full.

We proceed in three steps that mirror the paper’s own structure: a classical probabilistic formulation of LP, a kernelized generalization in which the excitation is replaced by a Gaussian process, and the full Bayesian nonparametric limit in which the kernel is itself inferred from the data. To these we add two contributions of our own: a Woodbury-Cholesky reformulation of the inference equations that avoids all explicit $M \times M$ matrix inversions, and two reparameterizations of the prior hyperparameters into quantities with direct physical interpretations.

2.1 Linear prediction

Let $\mathbf{x} = (x_1, \dots, x_M)^\top \in \mathbb{R}^M$ be M consecutive speech samples drawn from a short analysis frame. The autoregressive (AR) model of order P asserts that each sample is a linear combination of the P most recent ones plus an excitation residual:

$$x_m = \sum_{p=1}^P a_p x_{m-p} + \varepsilon_m. \quad (15)$$

The vector $\mathbf{a} = (a_1, \dots, a_P)^\top \in \mathbb{R}^P$ collects the *predictor coefficients* and $\boldsymbol{\varepsilon} = (\varepsilon_1, \dots, \varepsilon_M)^\top$ is the *excitation*. In source-filter terms these play distinct roles: $\boldsymbol{\varepsilon}$ represents the glottal excitation and $\mathbf{a}$ encodes the resonance structure of the vocal tract.

Two matrices appear repeatedly. The *AR operator* $\boldsymbol{\Psi}(\mathbf{a}) \in \mathbb{R}^{M \times M}$ is approximately lower-triangular with unit diagonal and sub-diagonal entries given by $-a_1, \dots, -a_P$; it satisfies $\boldsymbol{\varepsilon} = \boldsymbol{\Psi}\mathbf{x}$, or equivalently $\mathbf{x} = \boldsymbol{\Psi}^{-1}\boldsymbol{\varepsilon}$. The *lag matrix* $\mathbf{X} \in \mathbb{R}^{M \times P}$ has m -th row $(x_{m-1}, \dots, x_{m-P})$ and is the design matrix for the normal equations.

The classical probabilistic LP model assumes white Gaussian excitation,

$$\boldsymbol{\varepsilon} \sim \text{Normal}(\mathbf{0}, \nu \mathbf{I}), \quad (16)$$

which yields the likelihood

$$\mathbf{x} \sim \text{Normal}(\mathbf{0}, \nu \boldsymbol{\Psi}^{-1} \boldsymbol{\Psi}^{-\top}). \quad (17)$$

Maximum-likelihood estimation under (17) yields the familiar normal equation $\mathbf{X}^\top \mathbf{X} \mathbf{a} = \mathbf{X}^\top \mathbf{x}$, which is what every standard LPC implementation solves.

The model works as designed for speech with broadband, noise-like excitation. It fails for voiced speech. During voiced phonation the excitation is periodic, not white, and fitting (17) to a pitched signal forces the estimated envelope to track the harmonic partials of the spectrum: the poles of the all-pole filter chase the pitch harmonics rather than the formants. This is the central defect that motivates everything that follows.

2.2 Kernel linear prediction

The fix is to replace the white-noise excitation model with a richer one. We model the excitation as a continuous function $\varepsilon(t)$ over time, approximated by a weighted sum of J basis functions plus a broadband residual:

$$\varepsilon(t) = \varphi(t)^\top \mathbf{w} + \eta(t), \quad (18)$$

where $\varphi(t) = (\varphi_1(t), \dots, \varphi_J(t))^\top$ and $\eta(t)$ is the residual. Sampling at times $t_1, \dots, t_M$ and writing $\Phi \in \mathbb{R}^{M \times J}$ for the design matrix with rows $\varphi(t_m)^\top$ gives $\varepsilon = \Phi \mathbf{w} + \boldsymbol{\eta}$ in vector form.

With independent Gaussian priors

$$\mathbf{w} \sim \text{Normal}(\mathbf{0}, \nu_w \mathbf{I}), \quad \boldsymbol{\eta} \sim \text{Normal}(\mathbf{0}, \nu_e \mathbf{I}), \quad (19)$$

marginalizing over $\mathbf{w}$ yields

$$\varepsilon \sim \text{Normal}(\mathbf{0}, \nu_w \mathbf{K} + \nu_e \mathbf{I}), \quad \mathbf{K} = \Phi \Phi^\top. \quad (20)$$

The matrix $\mathbf{K}$ is a *kernel matrix* with entries $K_{mm'} = \varphi(t_m)^\top \varphi(t_{m'})$, and the key observation of Kameoka et al. (2010) is that any positive semidefinite matrix can serve as $\mathbf{K}$ directly via the kernel trick – $K_{mm'} = k(t_m, t_{m'})$ for some kernel function k – without constructing the basis functions at all.

Substituting (20) into $\mathbf{x} = \Psi^{-1} \varepsilon$ gives the kernelized LP likelihood,

$$\mathbf{x} \sim \text{Normal}(\mathbf{0}, \Psi^{-1}(\nu_w \mathbf{K} + \nu_e \mathbf{I}) \Psi^{-\top}). \quad (21)$$

Classical LP (17) is the special case $\mathbf{K} = \mathbf{I}$ with $\nu = \nu_w + \nu_e$. A periodic kernel with period T concentrates excitation power at harmonics and thereby *decorrelates* the envelope estimate from the pitch structure – exactly the fix we needed.

Excitation semantics It is worth being precise about how the structured and broadband components of ε map onto the GIF picture. We write the excitation as $\varepsilon(t) = s(t) + \eta(t)$, where $s(t) = \varphi(t)^\top \mathbf{w}$ is the structured GP component and $\eta(t)$ is broadband noise, and we distinguish this from $u'(t)$, the true physical DGF. In this thesis $u'(t)$ is the theoretical DGF in the sense of source-filter theory, and inferred $\varepsilon(t)$ is our practical DGF estimate. During voiced regions these are close but not equal, and we always report $\varepsilon(t)$ rather than $s(t)$ alone as our output.

This is deliberate. Even ordinary LPC — the $s = 0$ limit where all excitation comes from $\eta(t)$ — already recovers recognizable DGF estimates from long enough analysis frames, because $\eta(t)$ is free to inject the sharp broadband transients needed to excite all formant resonances simultaneously. Those sharp transients also model the broadband energy injections associated with glottal closure instants in the vocal apparatus. Extending from LPC to IKLP enriches what $s(t)$ can model but leaves $\eta(t)$ active and free to do this useful work — which is precisely what we want.

2.3 Infinite kernel linear prediction

2.3.1 From a fixed kernel to an infinite hypothesis bank

The kernelized model (21) requires designing $\mathbf{K}$ to match the excitation structure, but for voiced speech the fundamental period T is itself unknown and must be inferred. The natural Bayesian response is multiple kernel learning: prepare a bank of I candidate kernels $\mathbf{K}_1, \dots, \mathbf{K}_I$ and define the excitation kernel as their weighted sum,

$$\mathbf{K} = \sum_{i=1}^I \theta_i \mathbf{K}_i, \quad \theta_i \geq 0. \quad (22)$$

Each weight $\theta_i \geq 0$ quantifies the degree to which hypothesis i explains the excitation. In the original formulation of Yoshii & Goto (2013), the bank indexes candidate fundamental periods $T_1, \dots, T_I$ on a grid, so the weights θ_i directly encode a posterior over F0 candidates.

This is an important but too-narrow framing. The index i is a generic *hypothesis axis*: it can index candidate periods, candidate glottal flow morphologies, candidate phonation modalities, or any structured excitation model that can be expressed as a kernel. We will exploit this generality systematically in later chapters, where the bank is populated with GP priors learned from synthetic DGF data. The resulting multiple-kernel likelihood is

$$\mathbf{x} \sim \text{Normal} \left(\mathbf{0}, \Psi^{-1} \left(\nu_w \sum_{i=1}^I \theta_i \mathbf{K}_i + \nu_e \mathbf{I} \right) \Psi^{-\top} \right). \quad (23)$$

MAP estimation of $\boldsymbol{\theta}$ under a regularizing prior does not yield truly sparse weights. Yoshii & Goto (2013) solve this by taking $I \rightarrow \infty$ and placing a *gamma process* (GaP) prior on the infinite weight sequence: for any finite truncation level I ,

$$\theta_i \sim \text{Gamma}(\alpha/I, \alpha), \quad (24)$$

and as $I \rightarrow \infty$ the vector $\boldsymbol{\theta}$ converges to a draw from a GaP with concentration α . It is proven that the effective number of active elements — those with $\theta_i > \varepsilon$ for any $\varepsilon > 0$ — is almost surely finite, and for $I \gg \alpha$ only $\mathcal{O}(\alpha)$ weights are substantially positive. Sparsity is a theorem about the prior, not an approximation.

The full IKLP likelihood is

$$\mathbf{x} \sim \text{Normal}\left(\mathbf{0}, \Psi^{-1} \left(\nu_w \sum_{i=1}^{\infty} \theta_i \mathbf{K}_i + \nu_e \mathbf{I} \right) \Psi^{-\top}\right). \quad (25)$$

This single model simultaneously infers the spectral envelope through $\mathbf{a}$, the fundamental frequency through the dominant θ_i , and voiced/unvoiced status through the ratio $\mathbb{E}[\nu_w]/\mathbb{E}[\nu_w + \nu_e]$ — all at once, in a principled Bayesian framework, with none of the separate heuristic steps that characterize classical DSP approaches to GIF.⁹

2.4 Inference

2.4.1 Priors and variational family

To complete the model we place priors on the positive scale parameters and on $\mathbf{a}$. Following Yoshii & Goto (2013), the scales receive gamma priors,

$$\nu_w \sim \text{Gamma}(a_w, b_w), \quad \nu_e \sim \text{Gamma}(a_e, b_e). \quad (26)$$

For the AR coefficients we generalize the isotropic prior $\mathbf{a} \sim \text{Normal}(\mathbf{0}, \lambda \mathbf{I})$ of the original paper to a full multivariate Gaussian,

$$\mathbf{a} \sim \text{Normal}(\boldsymbol{\mu}_a, \boldsymbol{\Sigma}_a), \quad \mathbf{Q} := \boldsymbol{\Sigma}_a^{-1}, \quad (27)$$

with precision matrix $\mathbf{Q}$. The isotropic case is $\boldsymbol{\mu}_a = \mathbf{0}$, $\mathbf{Q} = \lambda^{-1} \mathbf{I}$. We need the full form in Chapter 3, where informative priors derived from stability theory and spectral constraints replace the isotropic default.

Exact posterior inference is intractable, so we use a mean-field variational family,

$$q(\boldsymbol{\theta}, \mathbf{a}, \nu_w, \nu_e) = q(\mathbf{a}) q(\nu_w) q(\nu_e) \prod_{i=1}^I q(\theta_i), \quad q(\mathbf{a}) = \delta(\mathbf{a} - \mathbf{a}^*), \quad (28)$$

with $\mathbf{a}$ treated as a MAP point estimate. Each remaining factor is updated by coordinate ascent on the ELBO $\mathcal{L}$:

$$\log p(\mathbf{x}) \geq \underbrace{\mathbb{E}[\log p(\mathbf{x} \mid \boldsymbol{\theta}, \mathbf{a}, \nu_w, \nu_e)]}_{\text{expected log-likelihood}} + \sum_z \mathbb{E}[\log p(z)] - \sum_z \mathbb{E}[\log q(z)] =: \mathcal{L}. \quad (29)$$

Compared to MCMC-based GIF approaches — where Auvinen et al. (2014) report runtimes on the order of 2.5 hours of single-CPU time per 25-millisecond analysis frame — variational inference keeps the Bayesian treatment while reducing inference to a sequence of closed-form updates. CAVI is convergence-guaranteed, parallelizable over batches of frames, and compatible with modern autodiff and GPU workflows even when the update steps are computed analytically.

⁹Setting $P = 0$ is also valid: the AR operator becomes $\Psi = \mathbf{I}$ and IKLP reduces to a pure multi-kernel Bayesian inference engine over the excitation process with no envelope component. This is useful as a standalone pitch detector or phonation-type classifier.

2.4.2 Tractable lower bound via matrix inequalities

The expected log-likelihood involves $\mathbb{E}[\log|\mathbf{K}|]$ and $\mathbb{E}[\mathbf{x}^\top \boldsymbol{\Psi}^\top \mathbf{K}^{-1} \boldsymbol{\Psi} \mathbf{x}]$, where $\mathbf{K} = \nu_w \sum_i \theta_i \mathbf{K}_i + \nu_e \mathbf{I}$, and neither expectation is tractable in closed form. Two matrix-variate inequalities provide bounds.

The log-determinant $\log|\mathbf{V}|$ is concave, so a first-order Taylor expansion around any PSD matrix $\boldsymbol{\Omega}$ gives

$$\log|\mathbf{V}| \leq \log|\boldsymbol{\Omega}| + \text{tr}(\boldsymbol{\Omega}^{-1} \mathbf{V}) - M. \quad (30)$$

The quadratic form $\mathbf{z}^\top \mathbf{V}^{-1} \mathbf{z}$ is convex, and an inequality of Sawada et al. (2012) gives

$$\mathbf{z}^\top \left(\sum_i \mathbf{V}_i \right)^{-1} \mathbf{z} \leq \sum_i \mathbf{z}^\top \boldsymbol{\Upsilon}_i^\top \mathbf{V}_i^{-1} \boldsymbol{\Upsilon}_i \mathbf{z}, \quad (31)$$

where $\{\boldsymbol{\Upsilon}_i\}$ are auxiliary matrices summing to the identity.

Applying both inequalities to the expected log-likelihood and optimizing the auxiliary quantities in closed form gives the tractable lower bound $\mathcal{L}' \leq \mathcal{L}$, with optimal values

$$\boldsymbol{\Omega} = \mathbb{E}[\nu_w] \sum_i \mathbb{E}[\theta_i] \mathbf{K}_i + \mathbb{E}[\nu_e] \mathbf{I}, \quad (32)$$

$$\boldsymbol{\Upsilon}_i = \tilde{w}_i \mathbf{K}_i \mathcal{S}^{-1}, \quad \boldsymbol{\Upsilon}_0 = \tilde{\nu}_e \mathcal{S}^{-1}, \quad (33)$$

where the *quadratic operator* $\mathcal{S}$ and harmonic-mean weights are defined by

$$\mathcal{S} = \sum_i \tilde{w}_i \mathbf{K}_i + \tilde{\nu}_e \mathbf{I}, \quad \tilde{\nu}_e = 1/\mathbb{E}[\nu_e^{-1}], \quad \tilde{w}_i = 1/(\mathbb{E}[\nu_w^{-1}] \mathbb{E}[\theta_i^{-1}]). \quad (34)$$

Both $\boldsymbol{\Omega}$ and $\mathcal{S}$ have the form $\nu \mathbf{I} + \sum_i w_i \mathbf{K}_i$; they differ only in which moment or harmonic-mean expectations enter as weights. This shared algebraic structure is what the Woodbury reformulation of Section 2.5 exploits.

2.4.3 GIG posteriors and CAVI updates

The sufficient statistics of a gamma prior are $\log z$ and z ; the bounds above introduce terms in z and $1/z$. The optimal variational posteriors for θ_i, ν_w, ν_e therefore all take the form of a *generalized inverse Gaussian*,

$$\text{GIG}(z \mid \gamma, \rho, \tau) \propto z^{\gamma-1} \exp(-1/2(\rho z + \tau/z)), \quad z > 0, \quad (35)$$

with moments

$$\mathbb{E}[z] = \sqrt{\tau/\rho} K_{\gamma+1}(\sqrt{\rho\tau})/K_{\gamma}(\sqrt{\rho\tau}), \quad \mathbb{E}[z^{-1}] = \sqrt{\rho/\tau} K_{\gamma-1}(\sqrt{\rho\tau})/K_{\gamma}(\sqrt{\rho\tau}), \quad (36)$$

where K_γ is the modified Bessel function of the second kind.

With the LP residual $\mathbf{r} := \Psi(\mathbf{a})\mathbf{x}$, the weighted residual $\mathbf{w} := \mathcal{S}^{-1}\mathbf{r}$, and the per-kernel quadratic statistics

$$\mathbf{u}_i := \Phi_i^\top \mathbf{w}, \quad q_i := \|\mathbf{u}_i\|^2, \quad q_0 := \|\mathbf{w}\|^2, \quad (37)$$

the CAVI parameter updates take the compact form. For the kernel weights:

$$\gamma_{\theta_i} = \alpha/I, \quad \rho_{\theta_i} \leftarrow 2\alpha + \mathbb{E}[\nu_w] \operatorname{tr}(\Omega^{-1}\mathbf{K}_i), \quad \tau_{\theta_i} \leftarrow (1/\mathbb{E}[\nu_w^{-1}]) \left(1/\mathbb{E}[\theta_i^{-1}]^2\right) q_i. \quad (38)$$

For the structured scale:

$$\gamma_w = a_w, \quad \rho_w \leftarrow 2b_w + \sum_i \mathbb{E}[\theta_i] \operatorname{tr}(\Omega^{-1}\mathbf{K}_i), \quad \tau_w \leftarrow \sum_i \left(1/\mathbb{E}[\nu_w^{-1}]^2\right) (1/\mathbb{E}[\theta_i^{-1}]) q_i. \quad (39)$$

For the broadband scale:

$$\gamma_e = a_e, \quad \rho_e \leftarrow 2b_e + \operatorname{tr}(\Omega^{-1}), \quad \tau_e \leftarrow \left(1/\mathbb{E}[\nu_e^{-1}]^2\right) q_0. \quad (40)$$

The AR update is MAP under the generalized prior (27), giving the normal equation

$$(\mathbf{X}^\top \mathcal{S}^{-1} \mathbf{X} + \mathbf{Q})\mathbf{a}^* = \mathbf{X}^\top \mathcal{S}^{-1} \mathbf{x} + \mathbf{Q}\boldsymbol{\mu}_a. \quad (41)$$

All updates are iterated until the ELBO converges or a maximum iteration count is reached.

The dominant costs per iteration are the I trace terms $\operatorname{tr}(\Omega^{-1}\mathbf{K}_i)$ and the operator actions $\mathcal{S}^{-1}\mathbf{r}$ and $\mathcal{S}^{-1}\mathbf{X}$, all requiring solutions to linear systems in matrices of the form $\nu\mathbf{I} + \sum_i w_i \mathbf{K}_i$. The next section reduces all of these to triangular solves on a single small matrix.

2.5 Woodbury-Cholesky reformulation

Both Ω and $\mathcal{S}$ are instances of the abstract operator

$$\mathbf{S}(\nu, \mathbf{w}) := \nu\mathbf{I}_M + \sum_{i=1}^I w_i \mathbf{K}_i \in \mathbb{R}^{M \times M}. \quad (42)$$

The direct approach materializes this $M \times M$ matrix, factors it by Cholesky ($\mathcal{O}(M^3)$), and solves systems against it ($\mathcal{O}(M^2)$ each). For $M = 2048$ samples at $I = 400$ kernels this is expensive; for large frames it is prohibitive.

Suppose each kernel admits a low-rank factorization $\mathbf{K}_i = \Phi_i \Phi_i^\top$ with $\Phi_i \in \mathbb{R}^{M \times R_i}$. In the settings we care about this is exact: the periodic kernels of the original paper are parameterized by Fourier features, and the arc cosine GP priors of later chapters are defined by their feature matrices Φ_i . Form the weighted stacked factor

$$\tilde{\Phi} := [\sqrt{w_1}\Phi_1, \dots, \sqrt{w_I}\Phi_I] \in \mathbb{R}^{M \times L}, \quad L = \sum_i R_i, \quad (43)$$

so that $\mathbf{S}(\nu, \mathbf{w}) = \nu\mathbf{I}_M + \tilde{\Phi}\tilde{\Phi}^\top$. The Woodbury matrix identity gives its inverse as

$$\mathbf{S}^{-1}\mathbf{v} = \nu^{-1}\mathbf{v} - \nu^{-2}\tilde{\Phi}\mathbf{A}^{-1}\tilde{\Phi}^\top\mathbf{v}, \quad \mathbf{A} := \mathbf{I}_L + \nu^{-1}\tilde{\Phi}^\top\tilde{\Phi} \in \mathbb{R}^{L \times L}. \quad (44)$$

The $L \times L$ *core matrix* $\mathbf{A}$ is the only object that needs to be factored. Its Cholesky decomposition $\mathbf{A} = \mathbf{R}^\top\mathbf{R}$ – computed once per operator instance per CAVI iteration – makes the following quantities available through triangular solves alone.

Operator solves To compute $\mathbf{S}^{-1}\mathbf{v}$: form $\mathbf{t} = \tilde{\Phi}^\top\mathbf{v} \in \mathbb{R}^L$, solve the triangular systems $\mathbf{R}^\top\mathbf{y} = \mathbf{t}$ then $\mathbf{R}\mathbf{u} = \mathbf{y}$ to get $\mathbf{u} = \mathbf{A}^{-1}\mathbf{t}$, then return $\nu^{-1}\mathbf{v} - \nu^{-2}\tilde{\Phi}\mathbf{u}$. No explicit inverse is formed.

Log-determinant By the matrix determinant lemma, $|\mathbf{S}| = \nu^M |\mathbf{A}|$, so

$$\log|\mathbf{S}| = M \log \nu + 2 \sum_{\ell=1}^L \log R_{\ell\ell}. \quad (45)$$

Trace of inverse

$$\text{tr}(\mathbf{S}^{-1}) = (M - L)/\nu + \nu^{-1} \text{tr}(\mathbf{A}^{-1}). \quad (46)$$

Per-kernel trace Define $\mathbf{B}_i := \tilde{\Phi}^\top\Phi_i \in \mathbb{R}^{L \times R_i}$. Then

$$\text{tr}(\mathbf{S}^{-1}\mathbf{K}_i) = \nu^{-1} \|\Phi_i\|_F^2 - \nu^{-2} \|\mathbf{R}^{-\top}\mathbf{B}_i\|_F^2, \quad (47)$$

where $\|\mathbf{R}^{-\top}\mathbf{B}_i\|_F^2$ is the squared Frobenius norm of the solution to $\mathbf{R}^\top\mathbf{Z} = \mathbf{B}_i$.

AR normal equation With $\mathcal{S}^{-1}$ available as a triangular-solve operator, the $P \times P$ Gram matrix $\mathbf{G} = \mathbf{X}^\top\mathcal{S}^{-1}\mathbf{X}$ is assembled column-by-column, and (41) is then a single $P \times P$ Cholesky solve of $\mathbf{G} + \mathbf{Q}$.

Every dense operation in the CAVI loop has been reduced to triangular solves on matrices of size at most $\max(L, P) \ll M$.

2.5.1 Complexity

Let M be the frame length, P the AR order, I the number of kernels, R the rank per kernel, $L = IR$ the total factor dimension, and B the number of frames in a batch.

In the dense baseline, building $\mathbf{S}$ costs $\mathcal{O}(IM^2R)$, factoring it costs $\mathcal{O}(M^3)$, each solve costs $\mathcal{O}(M^2)$, and storing it requires $\mathcal{O}(M^2)$ memory.

In the Woodbury-Cholesky regime, the core matrix $\mathbf{A}$ is assembled in $\mathcal{O}(ML^2)$ and factored in $\mathcal{O}(L^3)$. Each operator solve costs $\mathcal{O}(ML + L^2)$, the log-determinant costs $\mathcal{O}(L)$ after factorization, the per-kernel traces cost $\mathcal{O}(L^2R)$ each, and memory drops to $\mathcal{O}(ML + L^2)$: the $M \times M$ covariance is never stored. The crossover is favorable whenever $L \ll M$, precisely the regime of interest.

For B frames the loop vectorizes linearly in B . Since no dense $M \times M$ matrix is ever formed, the formulation is insensitive to its conditioning. Running in 32-bit floating point gives a

substantial throughput gain on accelerators while retaining numerical stability through the Cholesky factorization and a small diagonal jitter on $\mathbf{A}$.

2.6 Prior reparameterizations

The priors (24) and (26) are written in shape-rate form, natural for conjugate analysis but opaque for prior elicitation. We reparameterize both into quantities with direct physical interpretations.

2.6.1 Kernel sparsity: α -calibration

Under (24), define $T := \sum_i \theta_i$ and the normalized weights $p_i = \theta_i/T$. For any finite truncation I , the sum $T \sim \text{Gamma}(\alpha, \alpha)$ has $\mathbb{E}[T] = 1$ and $\text{sd}(T) = \alpha^{-\frac{1}{2}}$, while the proportions follow a symmetric Dirichlet:

$$(p_1, \dots, p_I) \sim \text{Dirichlet}(\alpha/I, \dots, \alpha/I). \quad (48)$$

Small α concentrates mass toward the vertices of the probability simplex – one dominant hypothesis – while large α spreads it uniformly.

A natural summary of effective sparsity is the exponentiated Shannon entropy,

$$I_{\text{eff}} := \exp(\mathbb{E}[H(\mathbf{p})]), \quad H(\mathbf{p}) = - \sum_i p_i \log p_i, \quad (49)$$

which counts the number of equally probable hypotheses that would produce the same entropy. Since the Dirichlet expectation of H is known analytically,

$$\mathbb{E}[H(\mathbf{p})] = \psi(\alpha + 1) - \psi(1 + \alpha/I), \quad (50)$$

we obtain

$$I_{\text{eff}(\alpha; I)} = \exp(\psi(\alpha + 1) - \psi(1 + \alpha/I)). \quad (51)$$

As $I \rightarrow \infty$ with α fixed, I_{eff} saturates at the constant $\exp(\psi(\alpha + 1) + \gamma_{\text{EM}})$ where γ_{EM} is the Euler-Mascheroni constant. At $\alpha = 1$ and large I this gives $I_{\text{eff}} \rightarrow e \approx 2.72$, meaning roughly a handful of simultaneously active hypotheses on average. This matches what Yoshii & Goto (2013) observe empirically at $\alpha = 1$, $I = 400$.

To calibrate α to a desired effective count I_{eff}^0 , one solves (50) for α ; the solution $\alpha^*(I)$ approaches 1 rapidly for $I_{\text{eff}}^0 = e$ as I grows. The local sensitivity of I_{eff} to α near this calibration is described by the log-elasticity

$$b(\alpha; I) := \frac{\partial \log I_{\text{eff}}}{\partial \log \alpha} = \alpha [\psi_1(\alpha + 1) - I^{-1} \psi_1(1 + \alpha/I)], \quad (52)$$

where ψ_1 is the trigamma function. As $I \rightarrow \infty$ this approaches $b \rightarrow \psi_1(2) = \pi^2/6 - 1 \approx 0.645$, giving the local approximation

$$I_{\text{eff}(\alpha; I)} \approx e(\alpha/\alpha^*(I))^{0.645}. \quad (53)$$

So α behaves like a concentration parameter whose effect on sparsity is governed by a fixed exponent near the default calibration.

2.6.2 Excitation power decomposition: the (p, s, κ) -reparameterization

When $b_w = b_e =: b$, the total excitation power $\nu_w + \nu_e \sim \text{Gamma}(a_w + a_e, b)$ and the structured fraction $\pi := \nu_w/(\nu_w + \nu_e) \sim \text{Beta}(a_w, a_e)$ are independent of each other. This motivates the reparameterization

$$a_w = \kappa p, \quad a_e = \kappa(1 - p), \quad b_w = b_e = \kappa/s, \quad (54)$$

under which the total power has prior $\nu_w + \nu_e \sim \text{Gamma}(\kappa, \kappa/s)$ with

$$\mathbb{E}[\nu_w + \nu_e] = s, \quad \text{Var}(\nu_w + \nu_e) = s^2/\kappa, \quad (55)$$

and the structured fraction has prior $\pi \sim \text{Beta}(\kappa p, \kappa(1 - p))$ with

$$\mathbb{E}[\pi] = p, \quad \text{Var}(\pi) = p(1 - p)/(\kappa + 1). \quad (56)$$

The parameter κ controls concentration only — it tightens both distributions simultaneously without moving their means. The parameters (p, s) directly control the prior expected values: s is calibrated to the expected dynamic range of the analysis frame, and p reflects prior beliefs about phonation type. A frame expected to be clearly voiced might receive $p = 0.8$; an unvoiced prior might use $p = 0.1$.

The default of Yoshii & Goto (2013), $(a_w, a_e, b_w, b_e) = (1, 1, 1, 1)$, corresponds to $(p, s, \kappa) = (1/2, 2, 2)$: equal prior probability of structured and broadband excitation, total power expected at 2, with moderate concentration. A unit-power neutral default is $(p, s, \kappa) = (1/2, 1, 1)$.

2.7 Summary

IKLP generalizes classical LP in three directions at once. The white excitation assumption is replaced by a GP prior whose kernel is a sparse mixture learned from the data. The single scalar noise variance is decomposed into structured and broadband components whose ratio tracks voiced/unvoiced status. And the AR prior is generalized to a full multivariate Gaussian, ready to accept the informative spectral priors of Chapter 3.

All of this fits inside a single CAVI loop with closed-form GIG updates and a convergence guarantee. The Woodbury-Cholesky reformulation of Section 2.5 removes the cubic bottleneck, reducing every inner-loop computation to triangular solves on matrices of size $\max(L, P) \ll M$. This makes the method both fast and composable: the kernel bank $\{\Phi_i\}$ is the only interface between the inference engine and the prior on the excitation. Replacing the periodic kernels of the original paper with the arc cosine GP priors of later chapters requires no change to the inference architecture — only the features Φ_i change.

The GIF problem of Chapter 1 requires finding plausible source-filter pairs consistent with a speech observation. IKLP, equipped with the right kernel bank, is exactly a hypothesis-testing engine for that problem.

3 Priors for autoregressive filters

The vocal tract filter is modeled throughout this thesis as an $\text{AR}(P)$ process: a rational all-pole transfer function whose P coefficients $\mathbf{a} = (a_1, \dots, a_P)^\top$ encode the resonant structure of the vocal tract. Estimating $\mathbf{a}$ from speech data is done traditionally by LPC methods, or in our case, by IKLP. Like any Bayesian model, IKLP requires a prior $p(\mathbf{a})$, and throughout their paper Yoshii & Goto (2013) simply assert

$$p(\mathbf{a}) = \text{Normal}(\mathbf{0}, \lambda \mathbf{I}_P), \quad \text{with a fixed } \lambda := 0.1. \quad (57)$$

This choice, which we will refer to as the *isotropic prior*, is clearly something more of a blunt regularizer than an informative prior. It’s there to prevent the normal matrix in (41) from becoming singular, not to encode anything known a priori about the vocal tract.

While it gets the job done, two things about (57) are perhaps unsatisfying. First, the isotropic prior with $\lambda = 0.1$ implies a standard deviation of about 3.16 per coefficient. This is a very loose prior under which most induced AR filters will almost surely be unstable. Second, we actually know a great deal about vocal tract filters: their resonances are spaced roughly 500 Hz apart, their spectral tilt is on the order of -6 dB/octave, and, plain from the physics of speech production, the filter must be at least be passive and causal. Except for causality none of this is encoded in (57).

This chapter then develops two priors that attempt to do better in this regard. The first, the *Monahan prior*, is the closest Gaussian to the uniform distribution over stable AR coefficient vectors; at first sight, this sounds like exactly the thing to do, but we show that it doesn’t work. The second, the *soft spectral shaping prior*, is a new construction that encodes a number of spectral features useful in speech – resonances and spectral rolloff – directly into the prior mean and covariance through a closed-form projection. Both priors remain Gaussian and are therefore compatible with IKLP’s variational inference without modification, and we test their merit in Chapter 7.

3.1 The isotropic prior

3.1.1 Expected output power

To understand the inductive bias that $\lambda = 0.1$ represents, consider what an isotropic Gaussian prior

$$\text{Isotropic}(\lambda) := \text{Normal}(\mathbf{0}, \lambda \mathbf{I}_P) \quad (58)$$

does to the expected output power of an $\text{AR}(P)$ process driven by white noise with variance ν as in equations (15) and (16) from the previous chapter:

$$x_m = \sum_{p=1}^P a_p x_{m-p} + \varepsilon_m, \quad \text{where} \quad \begin{cases} \mathbf{a} \sim \text{Normal}(\mathbf{0}, \lambda \mathbf{I}_P) \\ \varepsilon_m \sim \text{Normal}(0, \nu) \end{cases} \quad (59)$$

Up to scale, then, the output power is determined by the frequency response

$$H(e^{i\omega}) = \frac{1}{A(e^{i\omega})} \quad \text{where} \quad A(z) = 1 - \sum_{p=1}^P a_p z^{-p}. \quad (60)$$

To second order, the Maclaurin expansion of $|H(e^{j\omega})|^2$ in a_p gives

$$|H(e^{j\omega})|^2 \approx 1 + 2 \sum_{p=1}^P a_p \cos(p\omega) - \sum_{p=1}^P \sum_{\ell=1}^P a_p a_\ell \cos((p-\ell)\omega) \quad (61)$$

$$+ 4 \sum_{p=1}^P \sum_{\ell=1}^P a_p a_\ell \cos(p\omega) \cos(\ell\omega) + \dots \quad (62)$$

Taking the expectation over the zero-mean isotropic prior (57), the linear terms and the cross-terms $p \neq \ell$ vanish and the expected spectral power becomes

$$\mathbb{E}_{\mathbf{a}}[|H(e^{j\omega})|^2] \approx 1 - P\lambda + 4\lambda \sum_{p=1}^P \cos^2(p\omega). \quad (63)$$

Integrating over frequency and applying Parseval's identity gives the expected output variance of (59):

$$\mathbb{E}[\text{var}(x_m)] \approx \nu(1 + P\lambda). \quad (64)$$

With $\nu := 1$, $P = 30$ and $\lambda = 0.1$ as in Yoshii & Goto (2013), this gives an expected standard deviation of $\sqrt{1+3} = 2$, which is indeed $O(1)$. Nevertheless, power scales linearly with P , which contradicts the prevalent idea that increasing P just prepares the model for more fine-structured data, not “louder” data. The lesson from (64) is thus: to keep power approximately controlled under an isotropic prior, one needs $\lambda \propto 1/P$. This rule is a known heuristic in signal processing literature (Rabiner et al., 1993, for example; Sørbye & Rue, 2017), here we have derived it from a probabilistic argument.

3.1.2 Stability

Beyond raw power, there is a more fundamental problem: $\text{Isotropic}(\lambda)$ places most of its mass outside the stability region $\mathcal{A}_P \subset \mathbb{R}^P$, which is the set of coefficient vectors for which all roots of $A(z) = 1 - \sum_p a_p z^{-p}$ lie strictly inside the unit disk. Physically, *stability* means the filter's impulse response decays over time: the vocal tract is a passive resonator that dissipates energy, so its transfer function cannot have poles on or outside the unit circle. An unstable AR filter would predict an impulse response that grows without bound, which is physically impossible for a tube of flesh and air (the vocal tract).

For $P = 1$ the stability region is the interval $(-1, 1)$, and a unit Gaussian already places about 32% of its mass outside. For larger P the stability region is a complicated manifold and the fraction of unstable draws tends to grow rapidly – see Figure 7.




Figure 7 | **Probability that a prior draw lies in the stable AR region $\mathcal{A}_P$** for different values of λ with the isotropic prior (58). Stable prior mass collapses quickly under isotropic Gaussian AR priors, unless the $\lambda \propto 1/P$ rule is used with a $O(0.1)$ proportionality constant; the Monahan prior (68) performs underwhelmingly.

This figure also reveals that the proportional constant in the $\lambda \propto 1/P$ rule does seem to matter, because a Gaussian prior for $\mathbf{a}$ entangles expected power output and stability. At one extreme, $\lambda \rightarrow 0$ yields a filter that is trivially stable but will only produce signals with vanishing power. However, turning up expected power output to $O(1)$ by increasing λ then gradually decreases stability. The $\text{Isotropic}(0.1/P)$ prior is seen to strike a reasonable balance here. But can't we encode our preference for stability directly?

3.2 The Monahan prior

3.2.1 The Monahan map

A classical result from time series analysis (Monahan, 1984) is that the partial autocorrelation function (PACF) parameters $\boldsymbol{\varphi} = (\varphi_1, \dots, \varphi_P)^\top$ provide a smooth bijection between the hypercube $(-1, 1)^P$ and the stability region $\mathcal{A}_P$:

$$T : (-1, 1)^P \rightarrow \mathcal{A}_P, \quad \mathbf{a} = T(\boldsymbol{\varphi}). \quad (65)$$

The map T is defined by the so-called Levinson recursion which we illustrate below in Figure 8, but whose details are unimportant for our purpose. Its main value is that drawing $\varphi_p \stackrel{\text{iid}}{\sim} \text{Uniform}(-1, 1)$ produces only stable AR coefficient vectors. The push-forward of this uniform through T defines a density

$$g(\mathbf{a}) = 2^{-P} |\det \nabla_{\mathbf{a}} T^{-1}(\mathbf{a})|, \quad \mathbf{a} \in \mathcal{A}_P, \quad (66)$$

on the stability region, where 2^{-P} is just the density of the uniform base measure on $(-1, 1)^P$. This g is the natural “uniform over stable filters” measure: it assigns equal mass to equal volumes in PACF space, mapped into coefficient space by T .

Since tractable inference in AR models usually requires Gaussian priors, which (66) is not, the best we can do then is to approximate g by a Gaussian $p(\mathbf{a}) = \text{Normal}(\boldsymbol{\mu}_P, \mathbf{S}_P)$ that envelops $\mathcal{A}_P$ in some optimal way.

3.2.2 Deriving the prior

We want the Gaussian to encompass $g(\mathbf{a})$, such that it is neither overconfident about any subregion of $\mathcal{A}_P$, nor wasteful in placing mass far outside it. We thus minimize $D_{\text{KL}}(g \parallel \text{Normal}(\boldsymbol{\mu}_P, \mathbf{S}_P))$ over all Gaussians and solve for $\boldsymbol{\mu}_P$ and $\mathbf{S}_P$. This is a standard convex problem whose unique solution is the moment-matched Gaussian:

$$\boldsymbol{\mu}_P = \mathbb{E}_g[\mathbf{a}] = \mathbf{0}, \quad \mathbf{S}_P = \mathbb{E}_g[\mathbf{a}\mathbf{a}^\top]. \quad (67)$$

Surprisingly, the mean is identically zero by the sign-symmetry of the uniform PACF law and the fact that the Monahan map is odd in each coordinate. The Monahan prior is therefore

$$\text{Monahan}(P) := \text{Normal}(\mathbf{0}, \mathbf{S}_P), \quad (68)$$

where $\mathbf{S}_P$ captures the full geometry of the stability region in P dimensions within the approximation of a Gaussian family.

$\mathbf{S}_P$ has no closed form but is easy to estimate by Monte Carlo: draw N samples $\boldsymbol{\varphi}^{(n)} \stackrel{\text{iid}}{\sim} \text{Uniform}(-1, 1)^P$, compute $\mathbf{a}^{(n)} = T(\boldsymbol{\varphi}^{(n)})$, and average the outer products,

$$\hat{\mathbf{S}}_P = \frac{1}{N} \sum_{n=1}^N \mathbf{a}^{(n)} \mathbf{a}^{(n)\top} \rightarrow \mathbf{S}_P. \quad (69)$$

This converges at standard $\mathcal{O}(N^{-1/2})$ rate and can be precomputed for any desired P . At first sight this all appears promising: the Monahan map guarantees stability by construction, and (68) provides a convenient, tractable prior for IKLP that one might hope captures the shape of $\mathcal{A}_P$ sufficiently well.

3.2.3 Why it does not work

Sampling from $\text{Monahan}(P)$, however, leads to a rapidly vanishing fraction of stable filters as P increases, as shown by the last curve in Figure 7. The issue is geometric: g is supported entirely on $\mathcal{A}_P$, which is a highly non-elliptical subset of $\mathbb{R}^P$, while any nondegenerate Gaussian has full support on $\mathbb{R}^P$. Figure 8 illustrates this in one and two dimensions.

What if we tuck $\text{Monahan}(P)$ inside the support of $g(\mathbf{a})$ and make sure every draw leads to a stable filter? This would be done by minimizing the reverse projection $D_{\text{KL}}(\text{Normal}(\boldsymbol{\mu}_P, \mathbf{S}_P) \parallel g)$. Unfortunately, it is unbounded for any nondegenerate Gaussian since g assigns zero density outside $\mathcal{A}_P$. There is thus no Gaussian approximation that is simultaneously tractable and faithful to the stability constraint: the stable set is simply too irregular to be well-approximated by an ellipsoid.




Figure 8 | **Geometry mismatch.** For $P = 1$ the Monahan map is just $a_1 = \varphi_1$ and $\mathcal{A}_1 = (-1, 1)$. For $P = 2$ it gives $(a_1, a_2) = (\varphi_1(1 - \varphi_2), \varphi_2)$, so uniform PACF draws induce a density $g(\mathbf{a})$ supported on the triangular stable region. The Monahan prior $\text{Monahan}(P)$ matches the first two moments of g , but it cannot reproduce the bounded, non-elliptical shape of that support faithfully, and mass leaks out of the region of stable filters.

3.3 Spectral shaping priors

Setting worries about power and stability aside for a moment, suppose we expect that the AR filter has a resonance near 500 Hz, or that its spectrum rolls off at a particular rate, or that it has a formant structure typical of a particular vowel. All of these are relevant for acoustic phonetics. How can we “fold” such beliefs into the first and second moments of a simple Gaussian prior?

The spectral shaping priors that we introduce in this section tackle this problem in some generality. We present a recipe for encoding “spectral beliefs” into the prior mean of $p(\mathbf{a})$ through a single closed-form projection (82). This works by shifting the burden on the user to encode their desired *spectral features* as *roots of $A(z)$* , which are then imposed on average.

3.3.1 Spectral features as root constraints

Features of an AR filter’s spectral envelope $|H(e^{i\omega})|^2 = 1/|A(e^{i\omega})|^2$ can be traced to the pole structure of $H(z)$, or, equivalently, roots of $A(z)$. That is, controlling where $A(z)$ has its roots is an operational handle on the spectral shape of the filter.

For example, a conjugate pair of roots close to the unit circle at angle ω_0 produces a resonance peak at that frequency; the closer the roots sit to the unit circle, the sharper the peak. A real root near $z = 1$ boosts low frequencies and therefore induces negative spectral tilt, which is exactly the speech-like case of interest. These and other examples of spectral features are collected in Table 1.

One interesting way to encode such spectral features is through *polynomial divisibility*. Saying “we want $A(z)$ to have a root at z_0 ” is the same as saying “the factor $(1 - z_0^{-1}z)$ divides

Spectral feature	$Q(z)$	deg
Resonance at (ρ, ω_0)	$1 - 2\rho \cos(\omega_0)z + \rho^2 z^2$	2
Real pole at ρ	$1 - \rho z$	1
Negative spectral tilt $-6k$ dB/oct	$(1 - \rho z)^k$	k
Seasonal pole of period s	$1 - z^s$	s
Repeated resonance at (ρ, ω_0)	$(1 - 2\rho \cos(\omega_0)z + \rho^2 z^2)^k$	$2k$

Table 1 | **Spectral features and their polynomial encodings.** Each feature corresponds to a polynomial $Q(z)$ whose roots impose the desired pole structure on $|H(z)|^2 = 1/|A(z)|^2$.

$A(z)$.” A collection of desired roots defines a polynomial $Q(z)$, and the spectral constraint becomes

$$Q(z) \mid A(z), \quad (70)$$

which is read as “ Q divides A ”. For instance, a resonance near frequency ω_0 with damping $1 - \rho$ corresponds to the conjugate pair $z = \rho e^{\pm i\omega_0}$, giving the quadratic

$$Q(z) = 1 - 2\rho \cos(\omega_0)z + \rho^2 z^2. \quad (71)$$

Generalizing, assume $K \geq 1$ features are desired simultaneously. The individual $\{Q_k(z)\}_{k=1}^K$ are combined into a single *constraint polynomial*

$$M(z) = \text{lcm}\{Q_1(z), \dots, Q_K(z)\}. \quad (72)$$

Over $\mathbb{C}$ this amounts to collecting the union of desired roots with their maximum multiplicities, and the effective constraint degree is $L_{\text{eff}} = \deg M$. For example, a prior that expects both (1) a resonance and (2) a 18 dB/octave negative rolloff uses $M(z) = Q_1(z) \cdot Q_2(z)$ with $L_{\text{eff}} = 2 + 3 = 5$, where the degrees can be read off the last column of Table 1.

3.3.2 Divisibility is linear

The observation that makes (70) a tractable ansatz for encoding inductive bias towards given spectral features is that polynomial divisibility is a *linear* condition on the coefficients of A . $M(z) \mid A(z)$ if and only if the remainder of the polynomial division $A \bmod M$ vanishes, and this remainder depends linearly on $\mathbf{a}$.

More precisely, for a constraint polynomial $M(z)$ of degree L_{eff} , there exists a matrix $\mathbf{F} \in \mathbb{R}^{L_{\text{eff}} \times P}$ and a vector $\mathbf{c} \in \mathbb{R}^{L_{\text{eff}}}$ such that the remainder

$$\mathbf{f}(\mathbf{a}; M) := \mathbf{F}\mathbf{a} + \mathbf{c} = \mathbf{0} \quad \Leftrightarrow \quad M \mid A. \quad (73)$$

The matrix $\mathbf{F}$ and offset $\mathbf{c}$ are determined by the coefficients of $M(z)$ and the monic leading term $a_0 = 1$ of $A(z)$. The set of $\text{AR}(P)$ coefficient vectors compatible with the constraint is therefore an affine subspace of $\mathbb{R}^P$ of dimension $P - L_{\text{eff}}$.

Worked example To see this concretely, consider an $\text{AR}(4)$ polynomial

$$A(z) = 1 - a_1 z - a_2 z^2 - a_3 z^3 - a_4 z^4 \quad (74)$$

and suppose we want to encode a resonance at (ρ, ω_0) via the quadratic $Q(z) = 1 - cz + dz^2$ where $c = 2\rho \cos(\omega_0)$ and $d = \rho^2$. Divisibility means there exists a quotient $R(z) = 1 - r_1 z - r_2 z^2$ such that $A(z) = Q(z)R(z)$. Expanding and matching the coefficients of z, z^2, z^3, z^4 gives

$$a_1 = c + r_1, \quad a_2 = r_2 - cr_1 - d, \quad (75)$$

$$a_3 = dr_1 - cr_2, \quad a_4 = dr_2. \quad (76)$$

The last two equations pin down the quotient: $r_2 = \frac{a_4}{d}$ and $r_1 = \frac{a_3 + ca_4/d}{d}$. Substituting back into the first two gives two linear equations in $(a_1, \dots, a_4)$ alone,

$$\underbrace{\begin{pmatrix} 1 & 0 & -\frac{1}{d} & -\frac{c}{d^2} \\ 0 & 1 & \frac{c}{d} & \frac{c^2}{d^2} - \frac{1}{d} \end{pmatrix}}_{\mathbf{F}} \begin{pmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{pmatrix} + \underbrace{\begin{pmatrix} -c \\ d \end{pmatrix}}_{\mathbf{c}} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}, \quad (77)$$

from which the matrix $\mathbf{F}$ and vector $\mathbf{c}$ in (73) are read off. The resonance constraint thus carves out a $(P - L_{\text{eff}}) = 2$ -dimensional affine subspace of $\mathbb{R}^4$.

3.3.3 The soft spectral shaping prior

We can now state derive the soft spectral shaping prior for $\mathbf{a}$. Let $p_0(\mathbf{a}) = \text{Normal}(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0)$ represent our initial state of knowledge of the AR coefficients and let $M(z)$ be the constraint polynomial (72) encoding the desired spectral features. For example, the assignments

$$p_0(\mathbf{a}) := \text{Isotropic}(0.1/P), \quad M(z) := Q(z), \quad (78)$$

where $Q(z)$ is the resonance defined in (71), express a state of knowledge where we expect the AR filter to have unit power output, be stable, and have a single resonance.

We now update $p_0(\mathbf{a})$ to $p^*(\mathbf{a})$ by imposing the new piece of information that the statistic $\mathbf{f}(\mathbf{a}; M)$ has zero expectation under the new state of knowledge expressed by p^* – that is,

$$\mathbb{E}_{p^*}[\mathbf{f}(\mathbf{a}; M)] = \mathbf{0}. \quad (79)$$

Knuth & Skilling (2012) have shown that the unique variational potential for updating information is the Kullback-Leibler divergence, so our task is to

$$\text{find } p^* = \underset{p}{\text{argmin}} D_{\text{KL}}(p \parallel p_0) \quad \text{given } \mathbb{E}_p[\mathbf{f}(\mathbf{a}; M)] = \mathbf{0}. \quad (80)$$

Since $p_0(\mathbf{a})$ is Gaussian and $\mathbf{f}(\mathbf{a}; M) = \mathbf{F}\mathbf{a} + \mathbf{c}$ is linear in $\mathbf{a}$, we can solve this in closed form:

$$p^*(\mathbf{a}) = \text{Normal}(\boldsymbol{\mu}_M, \boldsymbol{\Sigma}_M), \quad (81)$$

which is the unique normalized density closest to p_0 under the constraint (79). The solution keeps Σ_0 unchanged, so $\Sigma_M = \Sigma_0$, but it does shift the mean to

$$\mu_M = \mu_0 - \Sigma_0 \mathbf{F}^\top (\mathbf{F} \Sigma_0 \mathbf{F}^\top)^{-1} (\mathbf{F} \mu_0 + \mathbf{c}). \quad (82)$$

Wrapping up, we write the resulting prior formally as

$$\text{SoftSpectral}(\mu_0, \Sigma_0; M) := \text{Normal}(\mu_M, \Sigma_0) \equiv p^*(\mathbf{a}), \quad (83)$$

and refer to it as the *soft spectral shaping prior*.

An example is shown in the left and middle panels of Figure 9, where we encode three resonances that have center frequency and bandwidth typical for the first three vowel formants. The prior spectrum induced by the base prior (left panel) has no discernable resonance structure. The middle panel shows the prior spectrum obtained after updating our state of information per (80). We can see the three resonances clearly imprinted, but what is more interesting is that the prior is noncommittal with respect to their actual locations: the zoom on the F_1 region shows clearly that the actual peak locations “happen around” the mean: this is the maximum entropy principle at work. There is also some interesting tail structure which indicates that the prior is anticipating extra resonances at higher frequencies.




Figure 9 | **Example of imposing spectral features** on an isotropic base prior leading to the soft and hard spectral shaping priors from (83) and (86). We used a constraint polynomial $M(z) = Q_1(z) \cdot Q_2(z) \cdot Q_3(z)$ to encode the spectral features for this example. These are three resonances at $(F_1, F_2, F_3) = (500, 1500, 2500)$ Hz as indicated by the dashed vertical lines. Each panel shows the spectrum mean (grey) together with three samples (colors) of the associated prior. From left to right, these are (1) the base prior $p_0(\mathbf{a}) := \text{Isotropic}(0.1/P)$ with $P = 12$; (2) the soft spectral shaping prior $\text{SoftSpectral}(\mathbf{0}, (0.1/P)\mathbf{I}_P; M)$ with an inset that zooms in on the F_1 region; and (3) the hard spectral shaping prior $\text{HardSpectral}(\mathbf{0}, (0.1/P)\mathbf{I}_P; M)$.

Hard divisibility If one prefers to enforce divisibility $M \mid A$ exactly rather than in expectation – for instance when a particular formant location is known with certainty – this can be derived by solving a related D_{KL} problem. However, since we already know $p^*(\mathbf{a})$ to be Gaussian, we can just condition it on the observation

$$\mathbf{f}(\mathbf{a}; M) = \mathbf{0}. \quad (84)$$

Then the standard conditioning formula gives the hard version of (83): the same mean shift, but the covariance is additionally deflated along the constrained directions,

$$\Sigma_M = \Sigma_0 - \Sigma_0 \mathbf{F}^\top (\mathbf{F} \Sigma_0 \mathbf{F}^\top)^{-1} \mathbf{F} \Sigma_0, \quad (85)$$

and we write the *hard spectral shaping prior* as

$$\text{HardSpectral}(\mu_0, \Sigma_0; M) := \text{Normal}(\mu_M, \Sigma_M). \quad (86)$$

Since we are rarely certain of spectral features in advance, this prior might be overconfident so less useful for applications.

An example is shown in the last panel of Figure 9. Compared to the soft spectral shaping prior spectra (middle panel), it is seen that the hard spectral shaping prior collapses the uncertainty around the first three resonances, which is allowed to increase again towards the tail end of the spectrum.

3.4 Modeling or expanding?

Jaynes (1991) recounts in a memorable passage the confusion of the renowned experimental physicist R. W. Wood, who in 1911 proposed optical experiments to determine whether the sidebands $\omega \pm \nu$ in the amplitude-modulated wave

$$(1 + m \cos \nu t) \cos \omega t = \cos \omega t + \frac{m}{2} \cos(\omega + \nu)t + \frac{m}{2} \cos(\omega - \nu)t \quad (87)$$

are “physically real frequencies actually present,” or whether “there is in reality only one frequency ω , but with varying amplitude.” Today it is clear that the above is a mathematical identity; the two descriptions are two ways of saying the same thing, and “no experiment could have any bearing on it,” as Jaynes wrote.

AR(P) filters raise similar questions. The polynomial $A(z)$ has P roots, and each conjugate pair near the unit circle produces a resonance peak in the spectral envelope $|H(e^{i\omega})|^2$. Now, the vocal tract genuinely has resonances – the formants $F_1, F_2, F_3, \dots$ with their bandwidths $B_1, B_2, B_3, \dots$ – and these are as physically real as anything in acoustics: they correspond to standing wave patterns in the cavities of the vocal tract and show up as broad peaks in the spectrum.

When P is chosen small, say $P \approx 2N_F$ where N_F is the number of formants in the analysis bandwidth, each conjugate pole pair tends to line up with one formant, and the correspon-

dence is fairly direct. In this regime, $A(z)$ is acting as a *physical model* of the vocal tract, and it makes sense to constrain it: stability is a physical requirement (the vocal tract is a passive resonator), spectral tilt reflects the glottal source coupling, and the formant locations themselves are knowable a priori within a vowel class (Peterson & Barney, 1952).

But this is not the only way to read $A(z)$. In signal processing practice, P is routinely set much larger than $2N_F$ – the well-known rule of thumb $P \approx f_s/1000 + 2$ for a sampling rate f_s gives $P = 2$ at 10 kHz, which is generous enough that many poles will have no individual formant counterpart. These “extra” poles shape spectral tilt, model subglottal coupling, smooth over source harmonics, or simply approximate fine spectral structure that no small set of resonances could capture on its own.¹⁰ In this regime, $A(z)$ is no longer a model but just a *rational expansion* of whatever the true transfer function happens to be – a Padé-type approximant whose individual poles are artifacts of the expansion order, not of the physics (Press et al., 1992).

This distinction underlies frequent confusions between poles and formants in acoustic phonetics (Titze et al., 2015). It also matters from the Bayesian point of view. If we are *modeling*, then encoding formant structure into the prior is encoding real physics and we should expect it to help inference. If we are *expanding*, then the same prior could easily turn into a strait-jacket on what is essentially a flexible basis, and we should expect it to hurt inference. These two situations map cleanly on whether we assume informative or noninformative priors.

The same tension applies to the concept of stability itself. In the modeling regime, stability is a genuine physical constraint and enforcing it through the prior is well motivated. But in the expansion regime, a root landing slightly outside the unit circle need not signal physical instability – it may simply mean that the rational approximation of order P is locally imperfect, and penalizing it amounts to penalizing the approximation for being an approximation.

3.4.1 Modeling *and* expanding?

Where does the soft spectral shaping prior (83) stand on all this? In theory at least, rather well, because it operates at what one might call the *suggestion* level.

The constraint polynomial $M(z)$ encodes beliefs about the few poles that *are* likely to have physical counterparts – the formants F_1, F_2, F_3 with their bandwidths, perhaps a spectral tilt factor – and (82) shifts the prior mean to a coefficient vector μ_M that is compatible with these features in expectation. But note that the covariance Σ_0 is left untouched. Indeed, the remaining $(P - L_{\text{eff}})$ degrees of freedom, corresponding to the “expansion” poles that have no individual physical meaning, are entirely free to go wherever the data require. So we are simultaneously modeling *and* expanding.

¹⁰A single broad formant, for instance, can easily be shaped by two or three poles acting together rather than a single conjugate pair (Van Soom & de Boer, 2021). Conversely, a pole may sit far from the unit circle and contribute little more than a nearly flat spectral factor. The correspondence between poles and formants is thus many-to-many, not one-to-one, except approximately in the parsimonious regime.

Nevertheless, even this soft, suggestive encoding of spectral structure may not always help. The $\lambda \propto 1/P$ scaling rule derived in (64) keeps expected power controlled under an isotropic prior, and does so without imposing any spectral shape at all. Whether the soft spectral shaping prior (83) actually improves inference over this simple baseline – or whether the mean shift inadvertently constrains the expansion poles in ways that hurt the fit – is ultimately an empirical question that depends on the data, the phonation type, and the model order P .

3.5 Summary

The IKLP machinery from the previous chapter requires us to choose a prior $p(\mathbf{a})$ for the AR coefficients. In this chapter we discuss three families of such priors, each attempting to encode progressively more structure about the vocal tract filter.

We began with the *isotropic prior* and showed that it entangles two things one would rather control independently: expected output power, which requires $\lambda \propto 1/P$ to stay bounded, and stability, which tends to collapse with P . In practice, the choice $\lambda := 0.1/P$ appears to be a good compromise, however.

The *Monahan prior* attempts to address the stability issue in a more principled way by approximating the uniform distribution over stable AR coefficient vectors, but the geometry of the stability region is too irregular to be faithfully captured by any Gaussian, and the attempt does not survive increasing P .

The *soft* and *hard spectral shaping priors* take a different approach entirely: rather than constraining where the coefficients may live, they encode what the resulting spectrum should roughly look like. These priors answer the question of how to anticipate complex speech-like spectra with something as simple as a Gaussian prior.

The proposed construction rests on the idea that an operational handle on “ $A(z)$ should have certain roots” is given by the concept of polynomial divisibility. That is, we can encode the desired spectral features as roots of $A(z)$, and ask that they divide the latter “on average.” This turns out to be a linear constraint on the coefficients $\mathbf{a}$, and we can solve the associated D_{KL} update in closed form. We present soft and hard flavors of the principle, where the former merely suggests spectral features and the latter insists on them.

This is, as far as we know, a new possibility to have it both ways in general AR modeling: *model* what you know, *expand* where you don’t. The soft and hard spectral shaping priors have broad applicability beyond speech processing and can be used for modeling any time series, whether finance or EEG signals. Because they are derived from a well-posed maximum entropy problem, they form the bridge between work in the literature on *model-based* (informative) priors on the one hand, and *expansion-based* (uninformative) priors on the other (Cuevas et al., 2021; Gibson, 2018; Sørbye & Rue, 2017).

4 From parametric to nonparametric glottal flow models

Glottal flow models (GFMs) describe the source signal $u(t)$ as it drives the vocal tract during voiced speech. These models operate in the time domain because the delicate phase characteristics of the *glottal cycle* are an integral part of vocal communication.¹¹

Decades of empirical work have over time produced a “model zoo” of GFMs: handcrafted waveforms of $u(t)$ with handfuls of carefully chosen parameters. Because these are *parametric* models, they all share the same basic trait: interpretable, but inevitably inflexible (Schleusing, 2012). Their parametric nature limits the range of time domain information they can encode. This motivates the construction of a family of *nonparametric* models, which are more expressive as they are able to grow “parametric capacity” depending on the data at hand.

With the many glottal flow models to follow, it is all too easy to lose connection to the actual physical phenomenon being studied. After a short look at the glottal cycle, we begin with a review of the classic Liljencrants-Fant model. Then we revisit the parametric piecewise polynomial models and generalize them to the nonparametric regime, where they are identified with the *arc cosine kernel* of Cho & Saul (2009). The latter is argued to combine the advantages of both worlds: the interpretability of parametric models with the flexibility of nonparametric ones. This then sets the stage for Chapter 5, which derives the (*quasi*)periodic *arc cosine kernel* as a learnable surrogate model of synthetic glottal flow data.




Figure 10 | **A look at the glottal cycle** via laryngeal stroboscopy of the vocal folds. Each frame here is a picture taken at intervals spanning different cycles to give the illusion of steady-state behavior. The glottis is the opening between the folds. After Chen (2016); Jopling (2009).

¹¹For example, plosives: glottal stops are micro-events at the millisecond level that underlie semantic information two orders of magnitude higher up the scale. The timing of the vocal fold movement is so precisely controlled by our brain that the slightest deviations are studied as biomarkers for neurodegenerative diseases like Parkinson’s (Ma et al., 2020).

The glottal cycle is the periodic opening and closing of the vocal folds, depicted in Figure 10. During one *pitch period*, a single cycle is completed, in this case spanning $T = 6$ ms. GFM's describe the rate of airflow $u(t)$ [typically expressed in cm^3/s] through the opening between the vocal folds called the *glottis*, visible as a black slit in Figure 10. Importantly, it is not at maximum glottis aperture that the acoustic output is greatest; it is at the sudden *glottal closure instant* (GCI), as first shown by Miller (1959) in modern literature. At this point, the moving air column in the vocal tract is abruptly interrupted, and kinetic energy is converted efficiently into acoustic energy which then excites the vocal tract much like plucking a harp or clapping in a resonant room (Chen, 2016, Section 4.6). The sharpness of this transition governs much of the clarity and perceived strength of voiced speech (Fant, 1979). Thus, efficiency demands that the glottal cycle be strongly asymmetric: a slow buildup of flow during the open phase, followed by a rapid, almost impulsive closure to the closed phase.

4.1 The Liljencrants-Fant model

A model that has been very useful in describing how $u(t)$ varies during the glottal cycle is the *Liljencrants-Fant (LF) model* proposed by Fant et al. (1985). It has been the GFM of choice for many joint-inverse filtering approaches¹² but also been studied in its own right.¹³

The model states that $u'(t)$ in a single pitch period of length T consists of three parts, or phases. The *open phase* (O) is modeled as a rising and falling exponential sinusoid part, and transitions to the *closed phase* (C) via an exponential return during the *return phase* (R). The latter is thus part of (C). Mathematically, this is defined piecewise as

$$u'_{\text{LF}}(t) = \begin{cases} -E_e e^{a(t-t_e)} \frac{\sin(\pi t/t_e)}{\sin(\pi t_e/t_p)} & 0 \leq t \leq t_e & \text{(O)} \\ -\frac{E_e}{\varepsilon T_a} (e^{-\varepsilon(t-t_e)} - e^{-\varepsilon(t_c-t_e)}) & t_e < t \leq t_c & \text{(C, R)} \\ 0 & t_c \leq t \leq T & \text{(C)} \end{cases} \quad (88)$$

where $\theta_{\text{LF}} = \{E_e, t_p, t_e, t_c, T_a, T\}$ are the six model parameters explained below in (91) and $\{a, \varepsilon\}$ are determined implicitly from the *closure constraint*

$$\int_0^T u'(t) dt = 0 \quad \text{such that} \quad u(0) = u(T), \quad (89)$$

which holds for any GFM, and the continuity constraint

$$\lim_{t \rightarrow t_e^-} u'(t) = \lim_{t \rightarrow t_e^+} u'(t) \quad (90)$$

which is peculiar to GFM's with smooth return phases rather than hard, discontinuous GCI events (Veldhuis, 1998). Note that like most GFM's, the LF model is defined in $u'(t)$ not $u(t)$

¹²See **Joint source-filter optimization** (page 9).

¹³See Degottex (2010, Section 2.4) for a discussion. Research into the LF model falls mainly into study of its spectral characteristics (Doval et al., 2006) and (re)parametrization, notably Fant (1995); Perrotin et al. (2021).




Figure 11 | **The Liljencrants-Fant model** for $u'_{\text{LF}}(t)$ (blue) and $u_{\text{LF}}(t)$ (red) during a single period of length T . The open phase (O), closed phase (C) and return phase (R) are marked at the top and are determined by the changepoints $\{t_p, t_e, t_c, T\}$. The amplitude at peak excitation t_e is E_e . Not shown here is time constant of the return phase T_a , which during (R) determines the steepness of exponential decay towards zero in $u'(t)$.

due to the radiation effect (7), though $u(t)$ can be recovered easily via integration (8). Figure 11 shows the waveforms for $u'_{\text{LF}}(t)$ and $u_{\text{LF}}(t)$ for a typical setting of θ_{LF} . When varying θ_{LF} the LF waveforms trace a family of broad to narrow asymmetric pulses with smooth to abrupt closure encoding positive ($u(t) \geq 0$) glottal flow.

Parameters of the LF model The LF model achieved broad adoption because it strikes a balance between simplicity and empirical accuracy. After several iterations, Fant et al. (1985) arrived at this parametrization:

$$\theta_{\text{LF}} = \begin{cases} E_e = \min_t u'(t) = u'(t_e) & \text{peak excitation amplitude} \\ t_p = \operatorname{argmax}_t u(t) & \text{instant of peak glottal flow} \\ t_e = \operatorname{argmin}_t u'(t) & \text{instant of peak excitation} \\ t_c & \text{instant of glottal closure (GCI)} \\ T_a & \text{time constant of the return phase} \\ T & \text{fundamental period} \end{cases} \quad (91)$$

which turned out to be effective in capturing the glottal cycle in both modal and more extreme phonations.¹⁴ Note that, apart from E_e , which sets the vertical scale, these are all

¹⁴Actually, Fant et al. present LF as a “four-parameter model”. They ignore E_e and set $T := t_c$ for convenience, thus arriving at a count of four not six as in (91). In their unifying theory of common GFM, Doval et al. (2006) also collapse T to t_c . It is common too in applications (for example, O Cinneide, 2012). Of course, nothing is really lost because a closed phase of any length can be appended to any GFM by piecewise concatenation. We choose to keep t_c distinct from T explicitly because this becomes significant for short return phases and is experimentally observed to be important, especially in clinical scenarios (Kreiman et al., 2007; Rosenberg, 1971).

time domain parameters which determine *changepoints* in the glottal cycle in Figure 10. These tend to covary significantly and can in fact be regressed into a single convenient parameter R_d (Fant, 1995).

From the modeling perspective, the most important changepoints are t_e and t_c : when these lie close together and return is swift (T_a small and vanishing return phase), the glottal cycle is at peak efficiency and $\tilde{u}(t)$ has broadband spectral energy. This results in a clear voice and sharp features in the speech waveform data **d**. Conversely, a smooth return (T_a medium to large) results in a longer return phase with t_c tending to T and the GCI becoming ill-defined. The result here is a more mellow voice and more smooth features in **d**. Both modes are common in voiced speech (Maurer, 2016) and GIF algorithms must be able to handle both cases.

Implementation In practice, the LF model is unwieldy to compute (Gobl, 2017; Veldhuis, 1998). The closure and continuity constraints in (89) and (90) lack an analytical solution, meaning that for each evaluation of $u'(t)$ one must numerically solve for the auxiliary constants $\{a, \varepsilon\}$ by a bisection or root-finding routine. This seems like a high price to pay for an analytical model, and worse, this optimization is brittle. As noted by Perrotin et al. (2021), simplified LF models exist that are perceptually equivalent to LF and easier to work with.

Nevertheless, it remains the de facto standard in parametric source-filter modeling and we therefore use it here. Our code in Appendix C implements the LF equations in JAX (Bradbury et al., 2020), so the bisection routines are compiled to native code for a speedup and the entire model is differentiable and easily batchable. The code also implements four LF parametrizations, including the previously mentioned R_d formulation of Fant (1995).

4.1.1 One of many: the glottal flow “model zoo”

Over the years, a large number of glottal flow models have been proposed in acoustic phonetics. A sample of these is shown in Figure 12, already an impressive lineup by 1986. Despite their visual differences, these models are, essentially, variations on a single theme: they all describe a glottal pulse that tracks the periodic opening and closing of the glottis.




Figure 12 | **A lineup of GF models** back from 1986 (handdrawn?), together with their derivatives (DGFs). These models are visually very similar. From Fujisaki & Ljungqvist.

¹⁵These are LF, KLGLOTT88 (Klatt & Klatt, 1990), Rosenberg C (Rosenberg, 1971), R++ (Veldhuis, 1998).

Doval et al. (2006) demonstrated that the most common GFMs¹⁵ share a common mathematical structure and can be unified under a single framework; the θ_{LF} parameters in (91) can in fact express all of them. These GFMs all model a glottal flow $u(t)$ that is positive or null, continuous, and typically differentiable except at the glottal closure instant (GCI). Within each period, $u(t)$ rises during the opening phase, falls during the closing phase, and returns to zero during the closed phase, possibly through a short return phase that smooths the closure. The derivative $u'(t)$ alternates positive, zero, and negative in the expected order and integrates to zero over one period, satisfying the closure constraint (89). It is continuous in case of a nonempty return phase and discontinuous otherwise (meaning a jump at t_c as in Figure 13), known as hard closure.

Allow partial closure? The strict closure constraint in (89) is of course just a convenient modeling choice and must be expected to be violated to varying degrees in reality, as real folds often leak, especially in breathy or pathological voices. On the one hand, it clearly is a defining feature¹⁶ of the periodic nature of the glottal cycle. On the other, empirical data show the limits of this hard constraint. For example, Kreiman et al. (2007, p. 601) observe that

in many cases in our data, returning flow derivatives to 0 at the end of the cycle conflicted with the need to match the experimental data and conflicted with the requirement for equal areas under positive and negative curves in the flow derivative. [...] These conflicts between model features, constraints, and the empirical data were handled by abandoning the equal area constraint.¹⁷

The approach we take in Chapter 5 is to learn a softened version of the closure constraint (89) from data rather than set it a priori. This means we learn both whether the flow typically returns to zero and, if not, to what extent this constraint is plausibly violated.

Allow negative flow? Most GFMs enforce a positive glottal flow and assert $u(t) \geq 0$. But there are exceptions such as Fujisaki & Ljungqvist (1986). They allow for transient negative flow which could represent a lowering of the vocal folds after GCI (presumably when air is sucked back into the lungs). As with the closure constraint, we will learn this effect from data in Chapter 5 and do not forbid $u(t) < 0$ a priori.

¹⁶The pioneering study on GIF by Miller (1959) had to deconvolve taped speech waveforms with analogue resistor networks and manual optimization. A review article on GIF methods notes the part played by the defining quality of the closure constraint in this search for $u(t)$: “the fine-tuning of the resistors of the inverse network was conducted by searching for such values that yielded zero flow after the instant of glottal closure” (Alku, 2011, p. 626).

¹⁷That “equal area constraint” is just another name for the closure constraint (89).

Degree d	Name	Reference
0	Triangular pulse model	Alku et al. (2002)
1	—	Titze (2000); Verdolini & Titze (1995)
2	KLGLOTT88	Klatt & Klatt (1990) Chang-Sheng Yang & Kasuya (1998)
3	R++	Veldhuis (1998)
3	FL model	Fujisaki & Ljungqvist (1986)
6	—	Childers & Hu (1994)

Table 2 | **Classic piecewise polynomial models** from the literature. Of these, KLGLOTT88 and R++ remain popular today (Doval et al., 2006).

4.2 Classic piecewise polynomial models

Polynomial GFMs occupy an interesting corner of the “model zoo.” They fit into the same piecewise framework we have been using for the LF model, but with each part modeled as a polynomial of degree d . We list several GFMs proposed in the literature in Table 2.

This class of GFMs mainly appeal to simplicity and spectral strength. Piecewise polynomial models admit an analytical solution to the closure constraint while remaining computationally cheap. More importantly, the sharp transitions generated by low-order polynomials have bright spectra with slow decay, which is exactly what it takes to model hard GCIs (Childers & Hu, 1994) and score well on perceptual tests of synthetic speech (Rosenberg, 1971).

To get a feel for these models, we now turn to the simplest case: $d = 0$.




Figure 13 | **The triangular pulse model** for $u'_\Delta(t)$ (blue) and $u_\Delta(t)$ (red) during a single period of length T . The parameters used in this model are identical to Figure 11, except for the lack of a return phase and that the ‘instant’ of maximum excitation t_e now spans the interval $[t_p, t_c]$.

4.2.1 The triangular pulse model

The simplest polynomial model is the *triangular pulse model* proposed by Alku et al. (2002). It takes $u(t)$ piecewise linear and $u'(t)$ piecewise constant, as shown in Figure 13. Mathematically, the triangular pulse model is given by:

$$u'_\Delta(t) = \begin{cases} +E_e \left(\frac{t_c}{t_p} - 1 \right) & 0 < t \leq t_p & \text{(O)} \\ -E_e & t_p < t \leq t_c & \text{(O)} \\ 0 & t_c < t \leq T & \text{(C)} \end{cases} \quad (92)$$

The model has no return phase and only four parameters: $\theta_\Delta = \{E_e, t_p, t_c, T\}$. Note that the amplitude of the first case in (92) is fixed because the closure constraint (89) removes one degree of freedom. Alku et al. (2002) point out that this enables a difficult-to-measure time domain quantity to be expressed as the ratio of two easy-to-measure quantities in the amplitude domain and exploit this fact to measure the open quotient more robustly.

4.3 Parametric piecewise polynomial models

The triangular pulse model (92) contains two jumps in the derivative $u'_\Delta(t)$, so we can write the latter more compactly as a linear combination of two Heaviside functions

$$u'_\Delta(t) = \begin{cases} a_1(t - b_1)_+^0 + a_2(t - b_2)_+^0 & 0 < t \leq t_c & \text{(O)} \\ 0 & t_c < t \leq T & \text{(C)} \end{cases} \quad (93)$$

where $(t - b)_+^0 = \text{Heaviside}(t - b)$ and

$$\mathbf{a} = \begin{pmatrix} a_1 \\ a_2 \end{pmatrix} = E_e \begin{pmatrix} \frac{t_c}{t_p} - 1 \\ -\frac{t_c}{t_p} \end{pmatrix}, \quad \mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} = \begin{pmatrix} 0 \\ t_p \end{pmatrix}. \quad (94)$$

During the open phase, we can interpret the rewritten triangular pulse model as a *parametric piecewise polynomial model* of degree $d = 0$ and order $H = 2$. It has parameters $\theta_2 = \{\mathbf{a}, \mathbf{b}\}$ which describe $H = 2$ jumps with amplitudes $\mathbf{a}$ at locations $\mathbf{b}$. Note that in this formulation the amplitudes $\mathbf{a}$ differ from the piecewise amplitudes in (92) but are still linearly dependent given the changepoints $\mathbf{b}$ and t_c .

Continuing, nothing stops us from generalizing both H and d . The open phase part of (93) becomes:

$$u'_H(t) = \sum_{h=1}^H a_h (t - b_h)_+^d \quad 0 < t \leq t_c \quad \text{(O)} \quad (95)$$

This is the general parametric piecewise polynomial model of degree d and order H . Again, its parameters $\theta_H = \{\mathbf{a}, \mathbf{b}\}$ are the amplitudes $\mathbf{a} = (a_1, \dots, a_H)^\top \in \mathbb{R}^H$ and the changepoints $\mathbf{b} = (b_1, \dots, b_H)^\top \in \mathbb{R}^H$. These scale and shift H piecewise monomials of degree d :




Figure 14 | **The RePU activation functions** — $(t)_+^d$ for $d \in \{0, 1, 2, 3\}$.

$$(t - b)_+^d = (t - b)^d \text{Heaviside}(t - b) = \begin{cases} 0 & t < b \\ (t - b)^d & t \geq b \end{cases} \quad (96)$$

We recognize (96) as the family of activation functions called *rectified power units* (RePUs). These include Heaviside for $d = 0$ and the influential ReLU function for $d = 1$, both shown among higher degree variants in Figure 14. It may sound like a stretch to relate polynomial GFM to neural nets, but this view of things will pay off greatly in Section 4.4 when we take the limit $H \rightarrow \infty$.

4.3.1 Regression view

GFM ultimately play the role of regression models in the GIF problem. From this point of view, the general parametric model in (95) is just a standard *linear model*, and we will now develop it into a Bayesian regression model following MacKay (1998).

Suppose we observe N time samples $\mathbf{t} = \{t_n\}_{n=1}^N$ during one glottal cycle, and fix H basis functions of the form

$$\phi_h(t) = (t - b_h)_+^d, \quad (97)$$

where the changepoints $\mathbf{b}$ and t_c are treated as given. Collecting these into the design matrix $\Phi \in \mathbb{R}^{N \times H}$ with

$$[\Phi]_{nh} = \phi_h(t_n), \quad (98)$$

the general parametric model during the open phase becomes simply

$$\mathbf{u}' = \Phi \mathbf{a}. \quad (99)$$

This is linear in the amplitudes $\mathbf{a}$ given Φ , so the model is rather constrained at this point. Indeed, $\mathbf{a}$ is currently the only source of variability and power in this model because we assume all other parameters fixed. Which prior then for $\mathbf{a}$ should we choose? If we want to maximize model support while adhering to the known additivity and power constraints, the optimal choice is the canonical Gaussian prior over amplitudes (Bretthorst, 1988):

$$p(\mathbf{a} \mid \mathbf{b}) = \text{Normal}(\mathbf{a} \mid \mathbf{0}, \sigma_a^2 \mathbf{I}). \quad (100)$$

Moving on to data space, the induced prior probability of $\mathbf{u}'$ is also Gaussian,

$$p(\mathbf{u}' \mid \mathbf{b}) = \int \delta(\mathbf{u}' - \Phi \mathbf{a}) p(\mathbf{a} \mid \mathbf{b}) d\mathbf{a} = \text{Normal}(\mathbf{u}' \mid \mathbf{0}, \sigma_a^2 \Phi \Phi^\top), \quad (101)$$

due to closure of the Gaussian under linear combinations, and

$$\mathbb{E}_{\mathbf{a} \mid \mathbf{b}}[\mathbf{u}'] = \Phi \mathbb{E}_{\mathbf{a} \mid \mathbf{b}}[\mathbf{a}] = 0, \quad (102)$$

$$\mathbb{E}_{\mathbf{a} \mid \mathbf{b}}[\mathbf{u}' \mathbf{u}'^\top] = \Phi \mathbb{E}_{\mathbf{a} \mid \mathbf{b}}[\mathbf{a} \mathbf{a}^\top] \Phi^\top = \sigma_a^2 \Phi \Phi^\top. \quad (103)$$

Because $\text{rank}(\Phi \Phi^\top) \leq H$, prior samples of $\mathbf{u}'$ are intrinsically confined to an H -dimensional subspace of $\mathbb{R}^N$. Assuming $H < N$, we thus see that linear models like (99) are inherently low-rank. Their expressivity grows with H and saturates when H reaches N .

Closure constraint The closure constraint (89) becomes

$$\int_0^T u'_H(t) dt = \sum_{h=1}^H a_h \int_0^{t_c} \phi_h(t) dt = \sum_{h=1}^H a_h r_h = 0, \quad (104)$$

with $r_h = (t_c - b_h)^{d+1}/(d+1)$. Observe that the prior (100) already satisfies the closure constraint in expectation:

$$\mathbb{E}_{\mathbf{a} \mid \mathbf{b}} \left[\sum_{h=1}^H a_h r_h \right] = \sum_{h=1}^H \mathbb{E}_{\mathbf{a} \mid \mathbf{b}}[a_h] r_h = 0. \quad (105)$$

This motivates setting the prior mean of the $\mathbf{a}$ to zero, otherwise than symmetry of ignorance around $\mathbf{a} = \mathbf{0}$. We can also enforce it:¹⁸

$$\sum_{h=1}^H a_h r_h = 0 \implies p(\mathbf{a} \mid \text{closure}, \mathbf{b}) = \text{Normal}(\mathbf{a} \mid \mathbf{0}, \sigma_a^2 (\mathbf{I} - \mathbf{q} \mathbf{q}^\top)) \quad (106)$$

where $\mathbf{r} = (r_1, \dots, r_H)^\top$ and $\mathbf{q} = \frac{\mathbf{r}}{\|\mathbf{r}\|}$. A convenient way to a sample from this updated prior makes use of the fact that this is a rank-one projection:

$$\mathbf{a} = (\mathbf{q} \mathbf{q}^\top) \mathbf{a} + (\mathbf{I} - \mathbf{q} \mathbf{q}^\top) \mathbf{a} \sim p(\mathbf{a} \mid \mathbf{b}) \quad (107)$$

$$\implies \mathbf{a}_\perp = (\mathbf{I} - \mathbf{q} \mathbf{q}^\top) \mathbf{a} \sim p(\mathbf{a} \mid \text{closure}, \mathbf{b}). \quad (108)$$

How does this look in data space? By the same calculation (102), marginalizing over $\mathbf{a}$ yields

$$p(\mathbf{u}' \mid \text{closure}, \mathbf{b}) = \text{Normal}(\mathbf{u}' \mid \mathbf{0}, \sigma_a^2 \Phi (\mathbf{I} - \mathbf{q} \mathbf{q}^\top) \Phi^\top). \quad (109)$$

¹⁸The general solution is given by $\text{argmin}_f D_{\text{KL}}[f(\mathbf{a}) \parallel p(\mathbf{a} \mid \mathbf{b})]$ under the constraint (104), but the closure under linear conditioning of Gaussian distributions allow us to take a shortcut here.




Table 3 | **Samples of $u(t)$** from the general parametric piecewise polynomial model (95) with or without the closure constraint (104). For each combination of degree d and order H , a total of H changepoints $b_h \sim \text{Uniform}(0, t_c)$ were drawn randomly, and four instances of $\mathbf{u}'$ conditioned on these were sampled either from (101) [**Without closure constraint**] or from (109) [**With closure constraint**]. Then $u(t) = \int_0^t u'(\tau) d\tau$ was obtained by quadrature of $\mathbf{u}'$.

We can then obtain $\mathbf{u}$ from $\mathbf{u}'$ via integration using (8). Compared to samples from the isotropic prior (101), the anisotropic prior (109) induces glottal flows $u(t)$ that tend to look more pulse-like, as shown in Table 3. Indeed, moving to the bottom right corner, it is seen that increasing H and d tends to produce LF-like pulses out of the box!

Regression with (109) ensures that posterior mass vanishes at solutions $\mathbf{u}$ that violate the closure constraint. This illustrates how linear constraints on the amplitudes $\mathbf{a}$ in the linear model (99) can encode GFM properties at the cost of just a single rank-one downdate per constraint. Moving on from linear models to Gaussian processes, these linear constraints become linear functionals known as *interdomain features*, which may be used to impose structure directly in function space, without touching rank, but more challenging mathematically.




Figure 15 | **Neural network view.** (a) The triangular pulse model as a tiny neural network with a single hidden layer, linear readout and t_+^0 activation. (b) The general parametric polynomial model of degree d with order H with hidden weights $\{\mathbf{b}, \mathbf{c}\} \in \mathbb{R}^{2H}$, output weights $\mathbf{a} \in \mathbb{R}^H$ and RePU activation function t_+^d .

Connection to classic polynomial models Observe that the analytical solution (106) to the closure constraint is valid for any degree $d \geq 0$ and order $H \geq 1$, so each classic polynomial model listed in Table 2 can be expressed in our linear model framework given suitable choices of $\mathbf{b}$ and t_c , without specialized derivations.

4.3.2 Neural network equivalence

Figure 15 illustrates that the triangular pulse model of Section 4.2.1 can be seen as a tiny neural network. Likewise, the general parametric piecewise polynomial model (95) can be recast as a single hidden layer of width H with a RePU activation and a linear readout:

$$u'_{\text{NN}}(t) = \sum_{h=1}^H a_h (c_h t - b_h)_+^d. \quad (110)$$

Each unit computes $\phi_h(t) = (\mathbf{w}_h^\top \mathbf{x}_t)_+^d$ with hidden weights $\mathbf{w}_h = (c_h, -b_h)^\top$ and $\mathbf{x}_t = (1, t)^\top$, exactly like a neuron with a bias input. The $3H$ parameters of the neural net are thus $\mathbf{a}$ and $\mathbf{b}$ (as before) and the newly acquired degrees of freedom $\mathbf{c} = (c_1, \dots, c_H)^\top$.

Seen from this point of view, a natural symmetry appears. We argued in Section 4.3.1 that in linear regression context the output weights $\mathbf{a}$ are best assigned an uninformative Gaussian prior, so

$$p(\mathbf{a}) = \text{Normal}(\mathbf{a} \mid \mathbf{0}, \sigma_a^2 \mathbf{I}). \quad (111)$$

But the neural network equivalence makes very clear that the hidden weights $\{\mathbf{b}, \mathbf{c}\}$ actually play similar roles: they enter linearly inside the activation, so the same argument applies, and we are encouraged to assign Gaussians again:

$$p(\mathbf{b}) = \text{Normal}(\mathbf{b} \mid \mathbf{0}, \sigma_b^2 \mathbf{I}), \quad p(\mathbf{c}) = \text{Normal}(\mathbf{c} \mid \mathbf{0}, \sigma_c^2 \mathbf{I}). \quad (112)$$

This step is not obvious from the viewpoint of any of the GFMs we’ve discussed before, where changepoints $\mathbf{b}$ were nonlinear hyperparameters fixed a priori; here they become “active” and “linear” degrees of freedom. Assuming t_c is known, the remaining hyperparameters to describe the open phase are thus the three scales $\boldsymbol{\theta}_{\text{NN}} = \{\sigma_a, \sigma_b, \sigma_c\}$ controlling amplitude, typical changepoint location, and horizontal stretch, respectively.

Turning to data space again and having assigned priors to all parameters, we can now do a full marginalization unlike the conditional in (101). Updating the design matrix Φ to

$$[\Phi]_{nh} = \phi_h(t_n) = (c_h t_n - b_h)_+^d, \quad (113)$$

and assuming the isotropic priors (111) and (112) we have

$$p(\mathbf{u}') = \int \delta(\mathbf{u}' - \Phi \mathbf{a}) p(\mathbf{a}) p(\mathbf{b}) p(\mathbf{c}) d\mathbf{a} d\mathbf{b} d\mathbf{c} \quad (114)$$

$$= \int \text{Normal}(\mathbf{u}' \mid \mathbf{0}, \sigma_a^2 \Phi \Phi^\top) p(\mathbf{b}) p(\mathbf{c}) d\mathbf{b} d\mathbf{c}. \quad (115)$$

This is a Gaussian mixture with varying covariances (not means) as Φ depends on the hidden weights $\{\mathbf{b}, \mathbf{c}\}$. A closed form for this integral exists only for $H \rightarrow \infty$ as $p(\mathbf{u}')$ converges to a Gaussian process, as we will see in Section 4.4.

Closure constraint The closure constraint (89) given $\{\mathbf{b}, \mathbf{c}\}$ remains a linear condition on $\mathbf{a}$, so we can define

$$p(\mathbf{a}, \mathbf{b}, \mathbf{c} \mid \text{closure}) = p(\mathbf{a} \mid \text{closure}, \mathbf{b}, \mathbf{c}) p(\mathbf{b}) p(\mathbf{c}), \quad (116)$$

where $p(\mathbf{a} \mid \text{closure}, \mathbf{b}, \mathbf{c})$ is identical to (106) but with $r_h = (c_h t_c - b_h)_+^{d+1} / c_h (d+1)$. This prior yields a mixture in data space that is similar to (114):

$$p(\mathbf{u}' \mid \text{closure}) = \int \text{Normal}(\mathbf{u}' \mid \mathbf{0}, \sigma_a^2 \Phi (\mathbf{I} - \mathbf{q} \mathbf{q}^\top) \Phi^\top) p(\mathbf{b}) p(\mathbf{c}) d\mathbf{b} d\mathbf{c}. \quad (117)$$

Here both $\mathbf{q} = \frac{\mathbf{r}}{\|\mathbf{r}\|}$ and Φ depend on $\{\mathbf{b}, \mathbf{c}\}$. This integral has no closed form for $H \rightarrow \infty$, but we can still impose the closure constraint on (114) indirectly in the spectral domain via interdomain features (Chapter 5).

Connection to LF and other GFMs The neural network view of the parametric polynomial model also clarifies expressivity. It is well known that a single RePU layer of sufficient width can approximate any continuous waveform on a bounded interval (Hornik, 1993). Therefore, with its analytical closure constraint (116) and the ability to represent sharp changepoint behavior, the parametric piecewise polynomial model in the form (110) effectively unifies the entire family of GFMs encountered in the literature, from triangular pulses to LF-type exponentials and sinusoids, given sufficiently large H .

4.4 Nonparametric piecewise polynomial models

The correspondence of piecewise polynomials GFMs with a single hidden layer shows that their expressivity depends mostly on the order H , as model order is equivalent to layer width in the neural network picture. Indeed, *parametric* linear models have a finite representation capacity determined primarily by rank because the modeled function can only explore the subspace spanned by the basisfunctions.¹⁹ *Nonparametric* models grow capacity as the number of datapoints increases; in the case of Gaussian processes, posterior inference given N datapoints is equivalent to Bayesian linear regression with N linearly independent basisfunctions, hence full-rank (Rasmussen & Williams, 2006).

It thus makes sense to consider the limit $H \rightarrow \infty$ in our search for more expressive GFMs. Continuing where we left, we already saw in (114) that $p(\mathbf{u}')$ is a nontrivial Gaussian covariance mixture weighed by the isotropic priors defined in (112). This can be made explicit:

$$p(\mathbf{u}') = \int \text{Normal}(\mathbf{u}' \mid \mathbf{0}, \sigma_a^2 \Phi \Phi^\top) p(\mathbf{b}) p(\mathbf{c}) d\mathbf{b} d\mathbf{c} \quad (118)$$

$$= \int \text{Normal}(\mathbf{u}' \mid \mathbf{0}, \mathbf{Q}) p(\mathbf{Q}) d\mathbf{Q}, \quad (119)$$

where we defined the push forward

$$p(\mathbf{Q}) = \int \delta(\mathbf{Q} - \sigma_a^2 \Phi \Phi^\top) p(\mathbf{b}) p(\mathbf{c}) d\mathbf{b} d\mathbf{c} \quad (120)$$

which essentially counts the number of $\{\mathbf{b}, \mathbf{c}\}$ configurations that give rise to a given $\mathbf{Q}$. Thus each random sample of $\{\mathbf{b}, \mathbf{c}\} \in \mathbb{R}^{2H}$ drawn from (112) induces in data space a random covariance matrix $\mathbf{Q} \in \mathbb{R}^{N \times N}$,

$$\{\mathbf{b}, \mathbf{c}\} \mapsto \mathbf{Q} = \sigma_a^2 \Phi \Phi^\top, \quad \text{rank}(\mathbf{Q}) \leq H \quad (121)$$

and hence a Gaussian component with weight $p(\mathbf{Q})$ in the mixture (118). Below we will show that by choosing $\sigma_a \propto 1/\sqrt{H}$ and letting $H \rightarrow \infty$, the density of states (120) degenerates to

$$p(\mathbf{Q}) \rightarrow \delta(\mathbf{Q} - \mathbf{K}), \quad \text{rank}(\mathbf{K}) = N \text{ almost surely}, \quad (122)$$

for some $\mathbf{K} \in \mathbb{R}^{N \times N}$ derived below in (158), causing the mixture to collapse to a single Gaussian and thus completing the transition to an indexed Gaussian process with covariance matrix $\mathbf{K}$.

¹⁹More precisely: a higher rank naturally enlarges the subspace of attainable functions, but the prior over amplitudes still constrains where probability mass is placed within that subspace. Geometrically, the model explores an H -dimensional ellipsoid like (101) or (109) rather than the entire hyperplane. When data accumulate, the prior constraints are eventually overwhelmed and the model can move freely on that plane. Random kitchen sinks (Rahimi & Recht, 2008) take these two ideas to the extreme: forgo specialized modeling; just use H random basis functions with uninformative priors and rely on the fact that an entire H -dimensional hyperplane is reachable with high probability.

4.4.1 Mean and covariance description

While the mixture (118) has no closed form for finite H , progress can be made by computing its mean and covariance as before in (102):

$$\mathbb{E}[\mathbf{u}'] = \mathbb{E}_{\mathbf{b},c}[\mathbb{E}_{\mathbf{a}|\mathbf{b},c}[\mathbf{u}']] = \mathbf{0}, \quad (123)$$

$$\mathbb{E}[\mathbf{u}'\mathbf{u}'^\top] = \mathbb{E}_{\mathbf{b},c}[\mathbb{E}_{\mathbf{a}|\mathbf{b},c}[\mathbf{u}'\mathbf{u}'^\top]] = \sigma_a^2 \mathbb{E}_{\mathbf{b},c}[\Phi\Phi^\top], \quad (124)$$

The mean is zero everywhere due to our zero-mean prior $p(\mathbf{a})$. The covariance matrix, however, is nontrivial. To investigate, define

$$\mathbf{C} := \mathbb{E}[\mathbf{u}'\mathbf{u}'^\top] \in \mathbb{R}^{N \times N} \quad (125)$$

with elements

$$[\mathbf{C}]_{nm} = \sigma_a^2 \mathbb{E}_{\mathbf{b},c} \left[\sum_{h=1}^H \phi_h(t_n) \phi_h(t_m) \right], \quad (126)$$

where $\phi_h(t) = (c_h t - b_h)_+^d$ for the neural network model $u'_{\text{NN}}(t)$ was defined in (113). Since the $\phi_h(t)$ are statistically equivalent under i.i.d. priors like the ones we adopted in (112), define a single “mother wavelet” *without index h* as

$$\phi(t; b, c) = (ct - b)_+^d, \quad b \sim \text{Normal}(0, \sigma_b^2), \quad c \sim \text{Normal}(0, \sigma_c^2). \quad (127)$$

Then the elements of $\Phi\Phi^\top$ become sums of H i.i.d. random variables, and their expectation factorizes across h :

$$[\mathbf{C}]_{nm} = \sigma_a^2 \sum_{h=1}^H \mathbb{E}_{\mathbf{b},c}[\phi_h(t_n) \phi_h(t_m)] \quad (128)$$

$$= \sigma_a^2 \sum_{h=1}^H \mathbb{E}_{b,c}[\phi(t_n; b, c) \phi(t_m; b, c)] \quad (129)$$

$$= \sigma_a^2 H \mathbb{E}_{b,c}[\phi(t_n; b, c) \phi(t_m; b, c)] \quad (130)$$

$$= \sigma_a^2 H \text{TACK}(t_n, t_m). \quad (131)$$

Here we anticipate the *temporal arc cosine kernel* of degree d , defined below in (146):

$$\text{TACK}(t, t') := \mathbb{E}_{b,c}[\phi(t; b, c) \phi(t'; b, c)] \quad (132)$$

$$\equiv \int \phi(t; \mathbf{w}) \phi(t'; \mathbf{w}) \text{Normal}(\mathbf{w} \mid \mathbf{0}, \Sigma) d\mathbf{w}, \quad (133)$$

where $\Sigma := \begin{pmatrix} \sigma_b^2 & 0 \\ 0 & \sigma_c^2 \end{pmatrix}$ describes the covariance of the hidden weights $\mathbf{w} = (-b, c)^\top$. For legibility we suppress the dependence of TACK, written $\text{TACK}(t, t')$, on both the degree d and the covariance matrix Σ .

To recapitulate, the first and second central moments of the mixture (118) are known for any H :

$$\mathbb{E}[\mathbf{u}'] = \mathbf{0}, \quad (134)$$

$$\mathbb{E}[\mathbf{u}' \mathbf{u}'^\top] \equiv \mathbf{C} = \sigma_a^2 H \text{TACK}(t, t'). \quad (135)$$

The third, fourth, ... central moments are in general nonzero and difficult to compute. They will vanish, however, when $H \rightarrow \infty$ and we choose $\sigma_a \propto 1/\sqrt{H}$. Nevertheless, even before taking any limit, the first two moments of the prior already mirror something kernel-like. Note that this result hinges critically on the independence assumption for $\{b_h, c_h\}$ in (127). The closure-constrained prior of (117) would couple these parameters nonlinearly, destroying that independence and making the derivation intractable, which is why we temporarily set it aside.

4.4.2 The temporal arc cosine kernel

Of course, I wouldn't have had the courage to attempt this whole calculation had I not already known that a very similar limit has been evaluated successfully in the neural-network-as-a-Gaussian-process literature. The marginalization we performed above is, in fact, a one-dimensional variant of the classic infinite-width limit of feedforward networks studied by Neal (1996); Williams (1998) and later generalized by Cho & Saul (2009).

In these works, the expectation over random weights $\mathbf{w}$ with $\text{Normal}(\mathbf{w} \mid \mathbf{0}, \Sigma)$ priors gives rise to a family of kernels which describe an infinitely wide Bayesian network, and which differ only with respect to the activation function used and the imposed a priori covariance Σ . This same reasoning will now allow us to identify our expected covariance (132) with a time-domain specialization of the *arc cosine kernel family* of degree d .

Table 4 organizes the two kernels we will introduce shortly. From here on we refer to these kernels in prose simply as ACK and TACK, and suppress their dependence on the listed parameters for legibility.

The arc cosine kernel (ACK) of degree d proposed by Cho & Saul (2009) is defined as the expected inner product between infinite-dimensional random RePU features of D -dimensional inputs $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^D$:

Name	Form	Parameters	Eq.
arc cosine kernel	$\text{ACK}(\mathbf{x}, \mathbf{x}') = \frac{1}{\pi} \ \mathbf{x}\ ^d \ \mathbf{x}'\ ^d J_d(\theta)$	d	(137)
temporal arc cosine kernel	$\text{TACK}(t, t') = \frac{1}{2\pi} \ \Sigma^{\frac{1}{2}} \mathbf{x}_t\ ^d \ \Sigma^{\frac{1}{2}} \mathbf{x}_{t'}\ ^d J_d(\theta)$	d, Σ	(146)

Table 4 | **Arc cosine kernels** introduced in this section. The ACK defined on arbitrary input dimension $D \geq 1$ was originally proposed by Cho & Saul (2009). The TACK is a simple application of the ACK for 1D time series.

$$\text{ACK}(\mathbf{x}, \mathbf{x}') = 2 \int (\mathbf{w}^\top \mathbf{x})_+^d (\mathbf{w}^\top \mathbf{x}')_+^d \text{Normal}(\mathbf{w} \mid \mathbf{0}, \mathbf{I}_D) d\mathbf{w}, \quad (136)$$

where the hidden weights $\mathbf{w} \in R^D$, and $\mathbf{I}_D$ is the $D \times D$ identity matrix. This integral representation makes it clear that this is indeed a valid positive definite kernel in any input dimension $D \geq 1$. It admits a beautiful closed form:

$$\text{ACK}(\mathbf{x}, \mathbf{x}') = \frac{1}{\pi} \|\mathbf{x}\|^d \|\mathbf{x}'\|^d J_d(\theta), \quad (137)$$

where

$$\theta = \arccos\left(\frac{\mathbf{x}^\top \mathbf{x}'}{\|\mathbf{x}\| \|\mathbf{x}'\|}\right) \in [0, \pi] \quad (138)$$

is the positive angle between $\mathbf{x}$ and $\mathbf{x}'$, and $J_d(\theta)$ is given by Cho & Saul (2009, Eq. 4) as the generator expression

$$J_d(\theta) = (-1)^d (\sin \theta)^{2d+1} \left(\frac{1}{\sin \theta} \frac{d}{d\theta} \right)^d \left(\frac{\pi - \theta}{\sin \theta} \right). \quad (139)$$

The first few instances of which are

$$J_0(\theta) = \pi - \theta, \quad (140)$$

$$J_1(\theta) = \sin \theta + (\pi - \theta) \cos \theta, \quad (141)$$

$$J_2(\theta) = 3 \sin \theta \cos \theta + (\pi - \theta)(1 + 2 \cos^2 \theta), \quad (142)$$

$$J_3(\theta) = 15 \sin \theta - 11 \sin^3 \theta + (\pi - \theta)(9 \cos \theta + 6 \cos^3 \theta). \quad (143)$$

Compared to the familiar Matérn kernels, the arc-cosine kernel is a rather atypical kernel, and probably only sparingly used in machine learning applications. It is, for one, blatantly nonstationary, and unlike stationary kernels such as the RBF, it can represent asymptotic behavior: its posterior mean need not revert to the prior mean outside the data regime, much like a neural network's response. The closed form (137) shows this structure explicitly: the integral representation factors into a polynomial term $\|\mathbf{x}\|^d \|\mathbf{x}'\|^d$, which controls the overall dynamic range and absorbs asymptotic variance, and an angular term $J_d(\theta)$, which captures the nonlinear thresholding of the RePU activation and thus allows for modeling the kind of changepoint structure we know is essential for GFMs.

The temporal arc cosine kernel (TACK) is the kernel which describes $\mathbf{C} \equiv \mathbb{E}[\mathbf{u}' \mathbf{u}'^\top]$ in (123). We can cast it as an affine shift of the bias augmented ACK on the input dimension $D = 2$. Define the auxiliary vectors

$$\mathbf{x}_t = \begin{pmatrix} 1 \\ t \end{pmatrix}, \quad \mathbf{w} = \begin{pmatrix} -b \\ c \end{pmatrix}. \quad (144)$$

and the covariance matrix for the hidden weights $\{b, c\}$

$$\Sigma = \begin{pmatrix} \sigma_b^2 & 0 \\ 0 & \sigma_c^2 \end{pmatrix}. \quad (145)$$

Then, by substituting these into (137) and integrating with respect to the Gaussian weight prior $\text{Normal}(\mathbf{w} \mid \mathbf{0}, \Sigma)$, we obtain the TACK:

$$\text{TACK}(t, t') = \mathbb{E}_{b,c}[\phi(t; b, c)\phi(t'; b, c)] \quad (146)$$

$$= \int (\mathbf{w}^\top \mathbf{x}_t)_+^d (\mathbf{w}^\top \mathbf{x}_{t'})_+^d \text{Normal}(\mathbf{w} \mid \mathbf{0}, \Sigma) d\mathbf{w} \quad (147)$$

$$= \frac{1}{2} \text{ACK}(\Sigma^{\frac{1}{2}} \mathbf{x}_t, \Sigma^{\frac{1}{2}} \mathbf{x}_{t'}) \quad (148)$$

$$= \frac{1}{2\pi} \|\Sigma^{\frac{1}{2}} \mathbf{x}_t\|^d \|\Sigma^{\frac{1}{2}} \mathbf{x}_{t'}\|^d J_d(\theta), \quad (149)$$

where the factor $1/2$ compensates for the convention used in Cho & Saul (2009). Here,

$$\theta = \arccos\left(\frac{\mathbf{x}_t^\top \Sigma \mathbf{x}_{t'}}{\sqrt{\mathbf{x}_t^\top \Sigma \mathbf{x}_t} \sqrt{\mathbf{x}_{t'}^\top \Sigma \mathbf{x}_{t'}}}\right), \quad (150)$$

is the angle between $\mathbf{x}_t$ and $\mathbf{x}_{t'}$ in the geometry induced by Σ . This anisotropic variant was also used in the context of D -dimensional feature extraction by Pandey & Dukkipati (2014).

Figure 16 makes the role of the degree parameter visible before we move on to the infinite-width limit. Across all three covariance settings, low degrees produce rougher draws with sharper local bends, whereas higher degrees regularize these bends into smoother open-phase trajectories. The covariance parameters then control where this structure is concentrated: smaller σ_b imply a more localized changepoint location and lead to sharp changepoint




Figure 16 | **Prior samples from TACK for different degrees d .** Each column fixes a different covariance setting (σ_b, σ_c) . Within each column the four curves correspond to different values of d . Increasing d progressively softens the hinge-point-like behavior, while decreasing σ_b relative to σ_c makes the prior more localized around the center.

behavior. Smaller values of σ_c just smear out time, which have the effect of longer timescales. Note that $d = 2$ and $d = 3$ are practically indistinguishable.

4.4.3 Taking the Gaussian-process limit

We turn again to the mixture (118). Recall from (121) that each random draw $\{\mathbf{b}, \mathbf{c}\} \in \mathbb{R}^{2H}$ from the prior (112) produces a random covariance matrix $\mathbf{Q} \in \mathbb{R}^{N \times N}$, which elementwise decomposes as a sum of H i.i.d. terms:

$$[\mathbf{Q}]_{nm} = \sigma_a^2 \sum_{h=1}^H \phi_h(t_n) \phi_h(t_m). \quad (151)$$

This sum tends to grow as $\mathcal{O}(H)$, since from (128)

$$\mathbb{E}_{\mathbf{b}, \mathbf{c}}[\mathbf{Q}] = \mathbf{C} \propto H. \quad (152)$$

Such growth of the marginal variance of $p(\mathbf{u}')$ makes the $u'_{\text{NN}}(t)$ model (110) useless as a prior when we want strong inductive bias even for large H , which we do. Therefore, we couple σ_a and H by choosing $\sigma_a \propto 1/\sqrt{H}$, in accordance with the variance-preserving initialization principle (Glorot & Bengio, 2010). Thus, substituting $\sigma_a \rightarrow \sigma_a/\sqrt{H}$ in (151), we get

$$[\mathbf{Q}]_{nm} = \sigma_a^2 \times \left[\frac{1}{H} \sum_{h=1}^H \phi_h(t_n) \phi_h(t_m) \right]. \quad (153)$$

Observe that the quantity between brackets is a Monte Carlo estimate of the TACK (146). Therefore, by the strong law of large numbers we conclude that

$$\left[\frac{1}{H} \sum_{h=1}^H \phi_h(t_n) \phi_h(t_m) \right] \rightarrow \mathbb{E}_{b,c}[\phi(t; b, c) \phi(t'; b, c)] \quad \text{as } H \rightarrow \infty, \quad (154)$$

with deviations vanishing as $\mathcal{O}(1/\sqrt{H})$. Therefore from (132)

$$[\mathbf{Q}]_{nm} \rightarrow \sigma_a^2 \text{TACK}(t_n, t_m) \quad \text{as } H \rightarrow \infty. \quad (155)$$

Equivalently, the induced density $p(\mathbf{Q})$ of covariance matrices collapses to

$$p(\mathbf{Q}) \rightarrow \delta(\mathbf{Q} - \mathbf{K}) \quad \text{as } H \rightarrow \infty, \quad (156)$$

which in turn collapses the mixture (118) to

$$p(\mathbf{u}') \rightarrow \text{Normal}(\mathbf{0}, \mathbf{K}) \quad \text{as } H \rightarrow \infty. \quad (157)$$

Here the *kernel matrix* is the symmetric PSD matrix $\mathbf{K} \in \mathbb{R}^{N \times N} \propto \mathbf{C}$ given by

$$[\mathbf{K}]_{nm} = \sigma_a^2 \text{TACK}(t_n, t_m). \quad (158)$$

In other words, the marginal $p(\mathbf{u}')$ converges entirely to a Gaussian and the first and second moments in (134) become *sufficient statistics* (Jaynes, 2003). All higher central moments vanish

as $\mathcal{O}(1/\sqrt{H})$ because the covariance fluctuations in (154) themselves decay at that rate. Since this argument holds for any finite collection of times $\mathbf{t} = \{t_n\}_{n=1}^N$, the limiting process is indeed officially a Gaussian process:

$$u'_{\text{NN}}(t) \sim \text{GaussianProcess}(0, \sigma_a^2 \text{TACK}(t, t')). \quad (159)$$

Except for H , it has the same hyperparameters as the neural network model of Section 4.3.2: the scales $\boldsymbol{\theta}_{\text{NN}} = \{\sigma_a, \sigma_b, \sigma_c\}$ describing overall power, typical changepoint location and horizontal stretch, respectively; and the two higher-level hyperparameters $\{d, t_c\}$, which determine the polynomial degree and the instant of glottal closure, respectively.

Integrating out changepoints Note that t_c is the only changepoint parameter that survived the Gaussian-process limit. All other changepoints [each changepoint being defined by a pair of amplitude and location parameters $\{a_h, b_h\}$] have been analytically integrated out. That act of marginalization removed the parametric bottleneck carried by finite rank models like (88), (92), (95), (110) completely. This may come as a surprise, because we pointed out before that changepoint parameters as in (91) constitute the main “control points” of any GFM expressed in the time domain: what previously required a discrete arrangement of control points is now averaged into a single continuous covariance function capable of learning the “timing relationships [which] are very important for modelling the glottal flow signal” (Doval et al., 2006, p. 1).

More aspects of (159) are discussed in greater detail in Appendix D.

4.5 Summary

Great care was taken to argue that any *parametric* glottal flow model of the open phase of the glottal cycle can be expressed as a RePU network with a single hidden layer of width H , given H large enough.

Then we showed that in a Bayesian regression context with noninformative priors, the limit $H \rightarrow \infty$ corresponds to a *nonparametric* glottal flow model: a zero-mean GP with the temporal arc cosine kernel. This well-known limit, borrowed from classical GP literature, is applied here to unify the two dominant approaches to GIF. These are [parametric $\Leftrightarrow$ joint source-filter estimation] and [nonparametric $\Leftrightarrow$ inverse filtering], and we attempt here to combine the best of both worlds: parametric interpretability and nonparametric expressivity, respectively.

That is, the proposed probabilistic model has support for any parametric glottal flow model while retaining capacity to “let the data speak for itself” if need be. In the next chapters we will refine this initially tiny support gradually into a strong inductive bias by learning features from synthetic glottal flow simulations.

5 The periodic and quasiperiodic arc cosine kernel

Chapter 4 established parametric and nonparametric glottal flow models for a *single period* of the glottal cycle. In this chapter, our goal is to extend this basic building block to *multiple periods*. Recordings of voiced speech often visually exhibit self-similarity on timescales of $O(10\text{ ms})$ because in normal speech the glottal cycle often reaches steady-state, and a clear pitch is perceived as a consequence.²⁰ An illustrative example for the case of steady-state vowels is given in Figure 17.

To model this steady-state source behavior we proceed in three steps. We begin with an idealized periodic kernel, the *periodic arc cosine kernel* (PACK), which is stationary and perfectly periodic. We then relax this to the *quasiperiodic arc cosine kernel* (QPACK), which retains the same within-cycle carrier while allowing slow amplitude variation and residual baseline structure across cycles. Finally we introduce *gated variants* (g-PACK and g-QPACK) in which the closed phase is incorporated explicitly. This last step breaks stationarity, but it turns out to be a useful one in practice.

The second half of the chapter turns from construction to learning. It is well known from functional data analysis that learning a single shared basis for a large corpus of training data requires aligning the data points, the process of which is called *registration*. (Ramsay & Silverman, 1997). We show how to do this accurately for the EGIFA data by aligning the




Figure 17 | **Examples of steady-state vowels** during three pitch periods. Here (a) is perfectly periodic because it was synthesized with (88); and (b) is quasiperiodic because it is taken from a real recording.²¹ The timescale is indicated by the horizontal bars ———— below the waveforms, each bar having a length of 5 msec.

²⁰Results from clinical trials (Little et al., 2007) indicate that normal (non-pathological) voices are usually something called Type I (Titze, 1995); that is, phonation results in nearly periodic sounds, which supports the notion that uttered vowels in normal speech typically reach a steady-state before “moving on”. Thus when a vowel is perceived with a clear and constant pitch, it is reasonable to assume that the vowel attained steady-state at some point (though perceptual effects forbid a one-to-one correspondence). In practice, steady-state vowels are identified simply by looking for such quasiperiodic strings, which typically consist of about 3 to 5 pitch periods in everyday speech (Rabiner & Schafer, 2007).

²¹Taken from a recorded steady-state instance of the vowel /æ/ at a fundamental frequency of 138 Hz. Source: bdl/arctic_a0017.wav:0.51-0.54sec (Kominék & Black, 2004).

examples to GCI (and optional GOI) events. This also allows us to extract voiced groups from the data, which then serve as high-quality training data to *calibrate* the global kernel hyperparameters. These hyperparameters determine the shared basis seen by the whole corpus, and we learn and evaluate them directly from the registered data.

Three algorithms do the practical work of the chapter. LBGID (Algorithm 1) detects the GCI and GOI events needed for registration; PRISM (Algorithm 2) calibrates a single global basis on the registered corpus; and t-PRISM (Algorithm 3) robustly filters raw period tracks into clean voiced groups before calibration. For the sake of brevity their full details and development have been deferred to the algorithm chapters; here we just accept them as given tools and only motivate their usage.

5.1 The periodic arc cosine kernel

The TACK or temporal arc cosine kernel derived in (146) provides a nonstationary covariance structure for the open phase of a *single* glottal cycle. That covariance reflects the change-point geometry inherited from piecewise polynomial GFMs, and it does so in a principled, nonparametric manner. But a single pitch period cannot exist in isolation: voiced speech is characterized by the imperfect repetition of this basic unit, and so we must now try to extend TACK to periodic signals.

Arguably the most natural route to periodicity is to replace the linear time coordinate in the bias-augmented input of TACK, namely $\mathbf{x}_t = (1, t)^\top$, by a set of *harmonic features* that are periodic simply by construction. This is the classical idea of embedding time on a circle (MacKay, 1998): functions of the embedded coordinates are automatically periodic in the original time variable. The construction proceeds in three steps: define the embedding, apply the arc cosine kernel, and normalize.

5.1.1 Harmonic embedding

Let $\omega_0 = 2\pi/T$ denote the fundamental angular frequency for a period $T > 0$, and let $J \geq 1$ be the *harmonic order*, which controls the richness of the embedding. Define the harmonic feature map

$$\varphi(t) = \begin{pmatrix} \sigma_b \\ \sigma_{c,1} \cos(\omega_0 t) \\ \sigma_{c,1} \sin(\omega_0 t) \\ \vdots \\ \sigma_{c,J} \cos(J\omega_0 t) \\ \sigma_{c,J} \sin(J\omega_0 t) \end{pmatrix} \in \mathbb{R}^{2J+1}, \quad (160)$$

where $\sigma_b > 0$ and $\sigma_{c,j} > 0$ for $j = 1, \dots, J$ are scale parameters playing analogous roles to σ_b and σ_c in Section 4.3.2, but act on the embedded space rather than linear time. These scale parameters modulate the harmonic basis (160) and thereby shape the angular similarity seen by the kernel. The constant coordinate σ_b is not a DC component of the waveform itself;

after normalization it becomes the zeroth harmonic weight and shifts the carrier toward or away from perfect self-alignment. The remaining scales control how strongly each harmonic pair participates in that comparison. In any case, where TACK lived in $\mathbb{R}^2$, we now work in $\mathbb{R}^{2J+1}$ with $(J + 1)$ independent harmonic scales, which gives the kernel much greater spectral flexibility, while still inheriting (some of) the changepoint behavior from its ancestor.

An interesting property of this particular embedding is that its squared norm is constant in t :

$$\|\varphi(t)\|^2 = \sigma_b^2 + \sum_{j=1}^J \sigma_{c,j}^2 (\cos^2(j\omega_0 t) + \sin^2(j\omega_0 t)) = \sigma_b^2 + \sum_{j=1}^J \sigma_{c,j}^2 := R^2. \quad (161)$$

Every point $\varphi(t)$ therefore lies on a sphere of radius R in $\mathbb{R}^{2J+1}$. For the inner product between two embedded points, the cosine addition formula $\cos \alpha \cos \beta + \sin \alpha \sin \beta = \cos(\alpha - \beta)$ collapses the harmonic pairs:

$$\varphi(t)^\top \varphi(t') = \sigma_b^2 + \sum_{j=1}^J \sigma_{c,j}^2 \cos(j\omega_0(t - t')). \quad (162)$$

This depends on t and t' only through their difference, and is manifestly periodic with period T . Define θ as the angle between the two embedded points,

$$\theta(t - t') := \arccos \left(\frac{\varphi(t)^\top \varphi(t')}{R^2} \right), \quad (163)$$

which takes values in $[0, \pi]$ and satisfies $\theta(0) = 0$.

5.1.2 Definition

Applying the ACK (137) to $\varphi(t)$ and $\varphi(t')$ gives

$$\text{ACK}(\varphi(t), \varphi(t')) = \frac{1}{\pi} R^{2d} J_d(\theta(t - t')), \quad (164)$$

which depends on t and t' only through their difference since R is constant. The prefactor R^{2d}/π couples the overall differentiability of sample paths (controlled by d) to the output scale. To circumvent this, we apply standard normalization and define the *periodic arc cosine kernel of degree d* as

$$\text{PACK}(t, t') := \sigma_a^2 \frac{J_d(\theta(t - t'))}{J_d(0)}, \quad (165)$$

where $J_d(\theta)$ was defined in (139) and $J_d(0) = \pi(2d - 1)!!$ is the value of the angular kernel at zero angle. Positive definiteness is preserved because composing a positive definite kernel with a deterministic feature map yields another positive definite kernel. By construction, the variance $\text{PACK}(t, t) = \sigma_a^2$, so σ_a is the only parameter that controls the overall amplitude

Symbol	Name	Role
σ_a	amplitude	marginal standard deviation
σ_b	bias amplitude	DC component of the embedding
$\sigma_{c,1}, \dots, \sigma_{c,J}$	harmonic scales	modulate each harmonic pair
T	period	fundamental period of the kernel
d	degree	smoothness class (RePU order)
J	harmonic order	number of harmonics in embedding

Table 5 | **Hyperparameters of PACK.** Of these, σ_a , σ_b and T are positive scalar, the harmonic scales $\sigma_{c,j}$ are J -dimensional, and d and J are discrete.

scale. Note that for legibility we generally suppress the dependence of $\text{PACK}(t, t')$ on its hyperparameters; we list them in Table 5 for an oversight.

How come PACK is stationary even though ACK and TACK were not? This is simply because we assigned equal weight to each cosine-sine pair in (162). It is a practical choice to ease training the kernel, and does not impair learning nonstationary phenomena. The thresholding behavior of the RePU activations is still present, but it is now replicated periodically.²² This means that PACK can still express very sharp within-cycle structure even though it is now stationary and periodic.

5.1.3 Hyperparameters and the weight simplex

The definition (165) involves $J + 1$ scale parameters (σ_b and $\sigma_{c,1}, \dots, \sigma_{c,J}$) that only enter through the angle θ in (163). Since θ depends on the normalized inner product $\varphi(t)^\top \varphi(t')/R^2$, only the *ratios* of these scales matter. This motivates introducing the *harmonic weights*

$$w_0 = \frac{\sigma_b^2}{R^2}, \quad w_j = \frac{\sigma_{c,j}^2}{R^2} \quad (j = 1, \dots, J), \quad (166)$$

which satisfy the simplex constraint $w_0 + \sum_{j=1}^J w_j = 1$ with $w_j \geq 0$. In terms of these weights, the cosine of the angle (163) becomes

$$\cos \theta(t - t') = w_0 + \sum_{j=1}^J w_j \cos(j\omega_0(t - t')), \quad (167)$$

and the kernel (165) can be evaluated by computing $\theta(t - t')$ from (167) and substituting into J_d . We learn the harmonic weights directly in the simplex space to avoid degeneracies.

²²The neural network interpretation of Section 4.3.2 carries over directly. PACK is the $H \rightarrow \infty$ limit of a single hidden layer with RePU activations applied to the $(2J + 1)$ -dimensional harmonic input (160) rather than the 2-dimensional bias-augmented input.




Figure 18 | **Effect of harmonic order J on PACK.** The top row shows the kernel section $\text{PACK}(t, 0)$ and the bottom row a prior sample. At fixed degree $d = 1$, increasing J decreases the effective within-cycle lengthscale and yields progressively wigglier covariance structure and sample paths. The remaining hyperparameters are fixed at $T = 1$, $\sigma_a^2 = 1$, $w_0 = 0.2$, and uniform harmonic weights $w_1, \dots, w_J$ over the remaining mass.

The zeroth order harmonic weight w_0 plays an interesting role. It descends directly from the bias amplitude σ_b in the TACK construction, but its behavior is altered nontrivially by the harmonic embedding. When $w_0 = 1$ (all harmonic weights zero), the angle θ is identically zero and the kernel degenerates to a constant σ_a^2 , describing a DC signal. Intermediate values of w_0 let the cosines oscillate about zero in (162) so that θ can reach $\pi/2$ and beyond, where $J_d(\theta)$ is able to produce steep, changepoint-like transitions of the type we need to model sharp glottal closures. If $w_0 = 0$, the autocovariance $\text{PACK}(t, 0)$ is allowed to reach zero periodically, which for higher d effectively creates kernels with compact support, although it isn't immediately clear what this implies for modeling speech signals.

At least J has a clear geometric effect. Increasing J inserts higher harmonics directly into the embedding and therefore decreases the effective within-cycle lengthscale of the process. In prior draws this appears as progressively finer oscillatory detail, as seen in Figure 18. Practically, increasing J while keeping all $w_j \sim O(1/J)$ gives rise to smaller effective lengthscales.

5.1.4 Spectral decay and smoothness

Because PACK is both stationary and periodic with period T , its covariance is an even T -periodic function that admits a Fourier cosine expansion

$$\text{PACK}(t, t') = a_0 + \sum_{m=1}^{\infty} a_m \cos(m\omega_0(t - t')), \quad (168)$$

with coefficients

$$a_0 = \frac{1}{T} \int_0^T \text{PACK}(t, 0) dt, \quad a_m = \frac{2}{T} \int_0^T \text{PACK}(t, 0) \cos(m\omega_0 t) dt \quad (m \geq 1). \quad (169)$$

By Bochner's theorem, positive definiteness of PACK is equivalent to $a_m \geq 0$ for all $m \geq 0$, so we can interpret these coefficients as encoding power. The Fourier expansion (168) has a direct interpretation as a *Bochner line spectrum* with spectral density given by

$$\tilde{k}(\xi) = 2\pi a_0 \delta(\xi) + \pi \sum_{m=1}^{\infty} a_m [\delta(\xi - m\omega_0) + \delta(\xi + m\omega_0)]. \quad (170)$$

The decay of the harmonic weights a_m determines both the spectral character of the prior and the mean square differentiability of GP sample paths. For stationary kernels, this is governed by the behavior of the covariance near the origin: the more derivatives $k(t) = k(t, 0)$ has at $t = 0$, the more mean square derivatives the process admits. A familiar example is the Matérn family, whose sample paths are r times mean square differentiable for every integer $r < \nu$. Now, it can be shown that the sample paths of a degree- d PACK kernel are exactly d times mean square differentiable, independent of J . This implies that the harmonic weights decay in power as

$$a_m = \mathcal{O}(m^{-(2d+2)}), \quad (171)$$

corresponding to an approximate rolloff in magnitude of $-6(d+1)$ dB per octave, as illustrated for the $d = 1$ case in Figure 20.

This behavior is also visible in Figure 19: increasing d removes the cusp at the origin and produces progressively smoother kernels with faster spectral decay. Low degrees ($d = 0$ or $d = 1$) therefore yield broad spectra with slowly decaying harmonics, appropriate for modeling sharp glottal closure events.

Classical source models report a spectral tilt of roughly -12 to -18 dB per octave for the glottal flow $u(t)$ itself (Doval et al., 2006; Rosenberg, 1971). Differentiation increases spectral slope by $+6$ dB per octave, so the corresponding glottal flow derivative $u'(t)$ exhibits a tilt of approximately -6 to -12 dB per octave, which corresponds to PACK kernels of degree $d = 0$ and $d = 1$.

5.1.5 Comparison to the standard periodic kernel

The standard periodic kernel in GP literature is the benign periodic squared exponential, which MacKay (1998) obtained by composing the RBF kernel with a circular embedding:

$$k_{\text{per}}(t, t') = \sigma^2 \exp\left(-2 \sin^2\left(\pi \frac{t - t'}{T}\right) / \ell^2\right). \quad (172)$$

Its exact cosine expansion is (Riutort-Mayol et al., 2020, Eq. B.2):




Figure 19 | **Effect of degree d on within-cycle smoothness.** The top row shows the kernel section $\text{PACK}(t, 0)$ and the bottom row a single prior sample. Holding J and the harmonic weights fixed, increasing d softens the changepoint-like geometry and produces smoother periodized waveforms. The remaining hyperparameters are fixed at $J = 1$, $T = 1$, $\sigma_a^2 = 1$, $w_0 = 0.2$, and $w_1 = 0.8$.

$$k_{\text{per}}(t, t') = \sigma^2 e^{-\frac{1}{\ell^2}} \left[I_0(\ell^{-2}) + 2 \sum_{m=1}^{\infty} I_m(\ell^{-2}) \cos(m\omega_0(t - t')) \right]. \quad (173)$$

Thus its entire harmonic envelope is controlled by a single lengthscale ℓ . Changing ℓ only broadens or narrows a smooth low-pass spectrum.

Figure 20 illustrates how PACKs are different in two ways. First, the harmonic weights w_j can redistribute mass freely across harmonics rather than merely broadening a fixed envelope. Second, the coupling between degree d and the asymptotic decay allows the model to learn more realistic, non-exponential spectral rolloff.

5.2 The quasiperiodic arc cosine kernel

Of course, perfect periodicity is too rigid for actual glottal flow; even synthesized glottal flow contains all sorts of drifts and interesting imperfections, which add organic warmth to the synthesized voice. Therefore the next step is to relax PACK into a quasiperiodic variant.

We thus define the *quasiperiodic arc cosine kernel* by

$$\text{QPACK}(t, t') := \text{env}(t, t') \cdot \text{PACK}(t, t') + \text{base}(t, t'). \quad (174)$$

This is a composite kernel modeling amplitude modulation of a periodic carrier, plus a second term that adds a small low pass correction. The full parameterization of (174) is listed below




Figure 20 | **PACK versus the periodic squared exponential in the line spectrum.** The curves show the harmonic weights a_m from (170) in dB. PACK with $d = 1$ exhibits the expected -12 dB per octave tail, while the periodic squared exponential traces out a one-parameter low-pass envelope. The PACK curves use $T = 1$, $\sigma_a^2 = 1$, and $w_0 = 0.2$, with uniform weights over the remaining harmonic mass; the periodic squared exponential uses lengthscale $\ell = 0.20$. The dashed reference line indicates -12 dB per octave.

in Table 6. In what follows, we choose the envelope $\text{env}(t, t')$ and baseline kernel $\text{base}(t, t')$ to be stationary kernels. This makes sure that QPACK is still stationary such that it has a well defined Bochner spectrum. We give an example of the latter in Figure 21 and now turn to each of the three components that make up QPACK.

Carrier kernel The carrier is a PACK of degree d from the previous section. The amplitude modulation introduced by the envelope does not materially alter the high-frequency rolloff, so the degree d still governs closure sharpness and spectral slope, while the weights and harmonic order J determine the detailed within-cycle geometry.

Envelope kernel The envelope kernel, written $\text{env}(t, t')$, modulates this carrier on a much longer timescale ℓ_{env} than the fundamental period T – typically on the order of a dozen periods. We use a unit-variance RBF,

$$\text{env}(t, t') := \exp\left(-\frac{(t - t')^2}{2\ell_{\text{env}}^2}\right), \quad (175)$$

for the envelope because this does not interfere with the differentiability class set by d . In the frequency domain, the product of a slowly varying RBF envelope with the PACK carrier broadens each of the harmonic lines into narrow bands like the ones on display in Figure 21.

Baseline kernel The baseline kernel, written $\text{base}(t, t')$, is a small additive correction. We choose the rational quadratic (RQ) kernel




Figure 21 | **From PACK to QPACK in time and frequency.** The top row shows prior samples over six periods; the bottom row shows the corresponding normalized spectra in dB against harmonic index j . The shared carrier hyperparameters are $d = 1$, $J = 2$, $T = 1$, $\sigma_a^2 = 1$, and $w = (0.2, 0.4, 0.4)$. For QPACK the envelope is a unit-variance RBF with lengthscale $\ell_{\text{env}} = 3$, and the baseline is a rational quadratic kernel with $\sigma_{\text{base}}^2 = 0.02$, $\ell_{\text{base}} = 0.5$, and $\alpha_b = 1$.

$$\text{base}(t, t') := \sigma_{\text{base}}^2 \left(1 + \frac{(t - t')^2}{2\alpha_b \ell_{\text{base}}^2} \right)^{-\alpha_b}, \quad (176)$$

with scale $\sigma_{\text{base}} \ll \sigma_a$, lengthscale $\ell_{\text{base}} \ll \ell_{\text{env}}$, and a general shape parameter $\alpha_b > 0$. Like the RBF, the RQ kernel is infinitely differentiable and will not influence the differentiability of the composite kernel. Its role is to capture weak residual structure which is not phase-locked enough to belong in the carrier but also not unstructured enough to be dismissed as white noise. This includes open-phase ripple caused by glottal-vocal-tract interaction (Ananthapadmanabha & Fant, 1982), as well as violations of the closure condition that occur frequently in more irregular or pathological voices (Kreiman et al., 2007). In spectral terms the baseline adds a broad low-frequency “pedestal” underneath the broadened carrier bands.

5.3 Gated variants

PACK and QPACK are manifestly stationary kernels. After conditioning on data they can produce phase-dependent (nonstationary) posterior means, but *a priori* they assign the same variance to every phase of the cycle, whereas we know that the excitation energy depends strongly on the phase of the glottal cycle.

Symbol	Component	Role
σ_a	carrier	PACK amplitude scale; fixed to 1 in the whitened calibration experiments
$w_0, \dots, w_J$	carrier	harmonic weight simplex inherited from PACK
d	carrier	within-cycle smoothness and spectral slope
J	carrier	harmonic order; effective within-cycle lengthscale control
T	carrier	fundamental period
ℓ_{env}	envelope	timescale of cycle-to-cycle amplitude modulation
σ_{base}	baseline	amplitude of the additive correction
ℓ_{base}	baseline	correlation lengthscale of the additive correction
α_b	baseline	rational quadratic shape parameter

Table 6 | **Hyperparameters of QPACK.** QPACK inherits the full PACK carrier and adds a long-scale envelope and a weak baseline. In the EGIFA calibration experiments the envelope is taken to have unit variance, so only its lengthscale is learned.




Figure 22 | **From PACK to QPACK to g-QPACK.** The top row shows kernel sections $k(t, t')$ at anchor time $t' = 0.75$, which is halfway the first open phase. The bottom row shows prior samples, all with matched carrier hyperparameters. QPACK relaxes strict periodicity by adding a slow envelope; g-QPACK then suppresses covariance in the closed phase of the cycle. The shared carrier settings are $d = 1$, $J = 1$, $T = 1$, and $w = (0.2, 0.8)$. To emphasize the closed phase the baseline correction was suppressed by setting $\sigma_{\text{base}} = 0$. The envelope is a unit-variance RBF with $\ell_{\text{env}} = 4$, and the gate sharpness of g-QPACK is $\beta = 12$.

The next step in taking advantage of prior information is therefore to encode the open/closed asymmetry directly in the kernel itself. This can be done easily by gating the kernel covariance (Alvarado & Stowell, 2016). Define therefore the two *gated variants*

$$\mathbf{g}\text{-PACK}(t, t') := g_\beta(t) g_\beta(t') \text{PACK}(t, t'), \quad (177)$$

$$\mathbf{g}\text{-QPACK}(t, t') := g_\beta(t) g_\beta(t') \text{QPACK}(t, t'). \quad (178)$$

These inherit all properties and hyperparameters of their respective base kernel except the property of stationarity.²³ The periodic gate is defined as

$$g_\beta(t) = \sigma(-\beta \sin(2\pi t/T)), \quad (179)$$

where $\beta > 0$ is a new hyperparameter that controls the sharpness and $\sigma(z) = (1 + e^{-z})^{-1}$ is the standard logistic sigmoid. This is a smooth approximation to a square wave: over one half-cycle $g_\beta(t) \approx 0$, over the other $g_\beta(t) \approx 1$, and the transitions are smooth.

Figure 22 illustrates the behavior of g-QPACK and compares it to its ancestors. As the gate $g_\beta(t)$ is itself periodic, g-PACK remains a proper extension of PACK (and incidentally a rare example of a periodic-but-nonstationary kernel). In practice, however, there is little reason to give up on the extra flexibility of QPACK, so the real model of interest is g-QPACK, not g-PACK, and we will focus on the former from here on.

Note that (179) implies that the first half of the period is closed phase and the second half open phase, with the two subintervals equally long. This means that g-QPACK always generates GF waveforms with fixed open quotient $\text{OQ} = 0.5$ (unless $\beta \rightarrow 0$). We relax this in the next section by warping physical time t into normalized time τ which can stretch and compress the phases and achieve any value of interest for OQ.

5.4 Registration with EGIFA

In the next section, we intend to learn the hyperparameters of the kernels introduced so far from the EGIFA dataset (Dataset 1). EGIFA provides a large number of high-quality glottal flow waveforms $\{u(t)\}$ simulated from a physics-based articulatory speech production model called VocalTractLab (Birkholz, 2013), which is the current open-source state of the art (Gao et al., 2024). We can use the generated $\{u(t)\}$ waveforms as exemplars to calibrate the PACK family on, but this requires aligning them beforehand – a procedure known as *registration* in functional data analysis, hence the title of this section.

In order to maximize learning signal, we therefore

1. extract only clean voiced groups from the raw $\{u(t)\}$ waveforms;
2. calculate their derivatives $\{u'(t)\}$ in a physiologically grounded way;²⁴
3. normalize power across the DGF exemplars;

²³This also means that the gated variants lack well-defined Bochner spectra. They are still in the family of harmonizable kernels, however (Altamirano & Tobar, 2022).

²⁴See Appendix E for details.




Figure 23 | **Registration of one retained EGIFA voiced group.** Top: the raw speech waveform $s(t)$ in physical time, extracted from the indicated .wav file and voiced group. Middle: the corresponding raw glottal flow $u(t)$, again in physical time, with LBGID-detected GCI and GOI landmarks. Bottom: the same group after differentiation, amplitude normalization and registration to normalized time, now shown as $u'(\tau)$. In this representation the GCIs land at integers $\tau = 0, 1, 2, \dots$ and the GOIs at half-integers $\tau = 0.5, 1.5, 2.5, \dots$

4. and finally align them in time such that the GCI and GOI events across different $u'(t)$ waveforms align in *normalized time* τ .

The procedure is illustrated in Figure 23, which follows a single EGIFA group through the registration process. With regard to step 4, time alignment, it is seen that the upper two rows live in physical time t , where period length and OQ vary continually and the GCIs and GOIs appear as irregularly spaced events. The bottom row lives in “subdued” normalized time τ , where the glottal closing and opening instants have been mapped to integer and half integer values of τ , respectively, and $OQ = 0.5$ everywhere.

Rather unfortunately, steps 1 and 4 above are nontrivial data preprocessing steps, and difficult to get right consistently, even for 1D time series. Because data quality is crucial, and with the prospect of reusing these transformations for future large scale benchmarks, we have had to resort to custom algorithms in order to guarantee meaningful alignment. LBGID (Algorithm 1) is a simple heuristic to reliably extract GCIs and GOIs from ground truth glottal flow $u(t)$.




Figure 24 | **Four samples from the registered EGIFA corpus.** From left to right and top to bottom these are waveforms with index $n = 19, 601, 1452, 5892$ out of a total $N = 7697$, with $(M_n, \tau_{nM_n}) = (875, 5), (884, 7), (860, 7), (759, 10)$.

With these in hand, t-PRISM (Algorithm 3) then clusters consistent progressions of such GCIs into voiced groups; that is, it segments the raw $\{u(t)\}$ waveforms into regions where the glottal cycle has reached something resembling steady state.

In this manner, most of the amplitude and phase variations in the raw data $\{u(t)\}$ are normalized to an orderly collection of waveforms with reliable signal to learn from. We leave a detailed discussion of the registration procedure to Dataset 1 and now proceed to the task of *calibration*; that is, learning the hyperparameters of the PACK kernels from the aligned $\{u'(\tau)\}$ exemplars.

5.5 Calibration on EGIFA

In ordinary GP regression we have one data set, learn a distribution over latent functions parametrized by the hyperparameters, and predict in the context of that same data set. Our setting here is quite different: we are given a large collection of N aligned DGF waveforms

$$\mathcal{D} = \{(\tau_n, \mathbf{u}'_n)\}_{n=1}^N, \quad (180)$$

where waveform n has its own registered grid

$$\tau_n = (\tau_{n1}, \dots, \tau_{nM_n}) \quad (181)$$

and standardized DGF samples

$$\mathbf{u}'_n = (u'_{n1}, \dots, u'_{nM_n}). \quad (182)$$

As Figure 24 illustrates, each waveform (182) is a data set in its own right, with its own length M_n , its own irregularly sampled grid (181) and its own number of completed cycles τ_{nM_n} .

Quantile	0%	25%	50%	75%	95%
Cycles covered	5	5	6	8	9

Table 7 | **Quantiles of covered pitch cycles at width $W = 2048$** across the registered EGIFA corpus. At the 95th percentile, the padded width accommodates about 9 cycles per frame.

Given one of the above kernels $k_{\theta}(\tau, \tau')$, the task is then to learn a single set of kernel hyperparameters θ that maximizes the support for $\mathcal{D}$. To distinguish from ordinary GP inference, we call this process *calibration*: learning a common basis from the corpus without yet “finetuning” on individual data points.

This is similar to basis learning in functional data analysis, but here we incorporate GP features like kernel regularization; minibatching; principled handling of irregularly sampled data. Algorithm 2 gives the derivation and further details. The basic idea is to extend the well-known sparse variational inference machinery of Titsias (2009) from a single function to many independent waveforms. We refer to this extension as PRISM, for *process-induced surrogate modeling*, and use it to accomplish the calibration task.

5.5.1 Experimental setup

All calibration experiments use the aligned EGIFA corpus from Section 5.4. After registration, each voiced group is smoothed, differentiated, and whitened within group. For ungated models (PACK, QPACK) we register only GCIs to integers; for gated models (g-PACK, g-QPACK) we additionally register GOIs to half-integers so that the gate transition occurs at a common phase. Each registered waveform is then padded to the common width $W = 2048$ introduced in Table 7 and masked at padding positions.

The total corpus comprises 6500 registered voiced groups (roughly 6.6M samples), split into a training set of 5000 groups and a held-out test set of 1000 groups. Training maximizes the collapsed PRISM objective on the training set. As derived in Algorithm 2, this objective is simply the (Titsias, 2009) collapsed bound summed across a mini-batch of complete waveforms with the standard $I / |\mathcal{B}|$ SVI correction. We optimize with Adam at a learning rate of 10^{-3} for 2500 iterations, using a batch size of 128 waveforms (accumulated over two microbatches of 64 for memory efficiency) and a jitter of 10^{-5} for numerical stability of the Cholesky factor.

We compare carrier degrees $d \in \{0, 1, 2\}$ and harmonic orders $J \in \{16, 32\}$, always with an RBF envelope, a rational quadratic baseline, and $M = 128$ inducing basis functions. For each of the resulting configurations we run 32 random initializations to assess sensitivity to the optimization landscape. Model selection uses not the training bound itself but a held-out reduced-rank log evidence per effective sample on the 500 test groups; this penalizes overfitting the basis to the training corpus. As a sanity check, each trained model is also compared to a matched white-kernel null that retains the same noise scale and inducing structure but removes the learned temporal covariance entirely.




Figure 25 | **Held-out fits from calibrated g-QPACK on EGIFA.** Two retained EGIFA groups are shown together with posterior means and uncertainty bands from the best single held-out g-QPACK run (pack:2, $J = 16$, $M = 128$). The shaded half-cycles $[n, n + 1/2]$ indicate the phase registered as closed. After GOI registration, gating leaves these intervals comparatively de-emphasized while retaining flexible interpolation across the informative open phase.

The practical comparison is between QPACK and g-QPACK. PACK is the conceptual starting point and g-PACK its gated analogue, but once quasiperiodicity and gating are both available it is their combination that matters.




Figure 26 | **Held-out calibration scores on EGIFA.** Each box summarizes 32 independent PRISM initializations for a fixed carrier degree and harmonic order, using 128 inducing basis functions, an RBF envelope, and a rational quadratic baseline. The vertical axis reports mean held-out reduced-rank log evidence per effective sample on 500 test groups. All shown models lie far above their matched white-kernel nulls, omitted for readability. Gating shifts the entire score distribution upward, while harmonic order 16 is consistently more reliable than harmonic order 32 at fixed basis size.

5.5.2 Results

The main empirical conclusion is immediate from Figure 26: gating matters more than the step from strict periodicity to quasiperiodicity alone. Across the entire grid, the gated models score higher than their ungated counterparts, and the gain is not a one-off artifact of a single favorable initialization. This is the empirical reason the chapter’s emphasis shifts from stationary PACK and QPACK to the nonstationary g-QPACK construction.

Within the ungated family, the strongest average configuration is `pack:1` with harmonic order 16, which attains a mean held-out score of about 0.297 plus or minus 0.112 across the 32 runs. Within the gated family, `pack:1` with harmonic order 16 is again best on average at about 0.412 plus or minus 0.089. The best *single* ungated run reaches about 0.429, while the best *single* gated run reaches about 0.569 and is obtained by `pack:2` with harmonic order 16. The degree-0 carrier is clearly too crude, whereas degrees 1 and 2 both capture the open-phase geometry well enough to be competitive.

The role of harmonic order is also visible now from the data rather than only from prior samples. At fixed basis size, moving from 16 to 32 usually hurts rather than helps. This is consistent with the interpretation from Figure 18: a larger harmonic order yields a wigglier carrier and therefore asks the finite PRISM basis to resolve sharper within-cycle structure. That extra flexibility is useful only when the data demand it strongly enough, and on EGIFA the more conservative order 16 turns out to be the better compromise.

Visual inspection agrees with the score summary. The held-out fits in Figure 25 show that a well-calibrated g-QPACK model interpolates complete voiced groups cleanly while leaving the registered closed phase comparatively de-emphasized. This is exactly what the score improvement suggests should happen: the model spends more of its capacity on the part of the cycle where the source is informative.

Why not do PCA?

For exploratory work, PCA is of course a useful baseline. But as a calibration engine it is mismatched to the present problem. The aligned EGIFA groups still have varying length and irregular registered grids, so a PCA pipeline must first commit to a common rectangular representation and then discard uncertainty in the resulting scores. PRISM does the reverse: it learns the shared basis directly on the registered functions and only then projects each waveform to a Gaussian in a common coefficient space.

There is also a conceptual difference. PCA learns whatever orthogonal basis best explains empirical variance in the training matrix. PRISM learns a basis induced by a specific kernel family, so calibration remains tied to the qualitative assumptions built into PACK, QPACK, and g-QPACK. That is the real point of the exercise. We do not merely want a convenient compression of EGIFA; we want a computationally tractable way to tune a theoretically motivated prior before carrying it into the downstream BNGIF model.

5.6 Summary

We began with PACK, the idealized periodic extension of TACK. Its stationarity and line spectrum made it the clean mathematical starting point, but also revealed the limits of perfect periodicity as a model of real glottal flow.

QPACK then relaxed this idealization to realistic quasiperiodic behavior by adding a slow envelope and a weak baseline term. The decisive step, however, was gating: the best-performing models in this chapter are not stationary, because the closed phase is represented directly by the gate rather than only through posterior inference. This is the conceptual reason the chapter’s emphasis shifts from PACK and QPACK to g-QPACK.

To make such kernels learnable on a corpus, alignment via LBGID and t-PRISM was required. That pipeline turned raw EGIFA flows into registered voiced groups in shared cycle time, with GOI-aware half-cycle normalization for the gated models. On that aligned corpus, PRISM could then calibrate one shared surrogate basis. This calibrated basis is the endpoint of the present chapter; the next chapter takes it as input and refines it into richer latent models of the DGF.

6 Surrogate glottal flow models

The story so far has been one of progressive refinement. Chapter 4 derived ACK and TACK from changepoints in classical glottal flow models. Chapter 5 lifted this to PACK, QPACK and g-QPACK and calibrated the resulting kernels on registered EGIFA waveforms. That calibration step already gave a useful global basis, but not yet a satisfactory generative model of DGF.

This chapter takes the next step. We keep the calibrated basis learned by PRISM, but refine the corpus of waveform posteriors themselves into a compact surrogate distribution. The aim is no longer to tune a kernel family, but to learn a genuinely data-driven prior over realistic DGF waveforms while retaining the Gaussian structure needed later by IKLP.

6.1 Refinement

6.1.1 Nonlinear dimensionality reduction

We can learn on the level of hyperparameters, which is a nonlinear optimization of the basis functions in the weight-space GP view.

Via PRISM we learn a basis for the problem.

However, to get better models, it is also necessary to learn on the second level, which is that of amplitudes.

From LF model we know that inherent dimensionality of GF models is 5. Since we normalize out amplitude and OQ we set intrinsic dimensionality of GPLVM to $Q = 3$, which is also convenient for learning. The number of clusters K is arbitrary; same modeling or expanding? question we discussed in Section 3.4

See Algorithm 4

Note: this is effectively a deep GP but linearized, and the time warping adds another GP layer.

Here interesting to see QPACK learns the closed phase: each cluster first specializing to OQ and then shapes

6.1.2 Clustering

xd-gmm to take in the uncertainties

6.1.3 Quantization

quantization means: in the deepest latent space, we quantize into components (basically expand the density), linearize the map, and push back

after we fitted XDGMM we push back

because we linearize, we get just gaussians back, and thus our reduced rank GP model (weight space view (Rasmussen & Williams, 2006)), and thus we can use this immediately as a prior for the DGFs in IKLP

because IKLP is a kernel mixture model, we can thus push each quantized component in the mixture: that's Chapter 7

here in this chapter we just return this as a mixture in direct data space

we can sample from it (with the π weights) and generate more realistic DGF waveforms than parametric models like LF

6.2 Phase warping

Static DGF surrogates are not enough for voiced speech. Even if the shape of one cycle is modeled well, realistic utterances still require the cycles to arrive at plausible times. That is the role of a pitch prior.

We learn that prior on APLAWD in Dataset 2 using t-PRISM Algorithm 3. The training story and model selection are kept there. Here we only use the result: a smooth GP prior over inter-GCI period trajectories.

For one voiced sequence with GCI sample indices $n_{\text{GCI},0}, \dots, n_{\text{GCI},K}$, define the local periods

$$T_k = \frac{n_{\text{GCI},k+1} - n_{\text{GCI},k}}{f_s}, \quad k = 0, \dots, K-1, \quad (183)$$

and the log-period process

$$\ell_k = \log T_k. \quad (184)$$

t-PRISM learns a shared prior over such trajectories. The posterior means already fit retained APLAWD pitch tracks well, and prior draws from the same model generate smooth period sequences on the timescale expected of human speech. Those sampled period tracks are what later drive the phase warps of the surrogate speech model.

6.2.1 Instantaneous phase and pitch

A large portion of the nonstationary variation of voiced speech can be described via a monotone *phase function* alone. Let

- t be physical time (in seconds),
- $\tau(t)$ be the cycle index or phase, dimensionless,
- $f_0(t)$ be the instantaneous pitch (Hz).

By definition,

$$\frac{d\tau}{dt} = f_0(t) \quad (185)$$




Figure 27 | **From learned pitch prior to time warp.** Top: two retained APLAWD period tracks together with their t-PRISM posterior means. Bottom left: prior samples of the instantaneous period function $T(\tau)$ from the learned APLAWD model. Bottom right: the corresponding induced warps $t(\tau) = t_0 + \int_0^\tau T(u) du$. This is the object needed later to turn refined DGF surrogates in cycle time into voiced signals in physical time.

This says that phase advances at a rate given by instantaneous pitch. The key is to work with time as a function of phase, not the other way around. That is the time-warping view.

In the simplest case, that of constant pitch, we have $f_0(t) := 1/T$ such that, integrating (185), we have

$$\tau(t) = \frac{t - t_0}{T} \quad (186)$$

where $t_0 \in \mathbb{R}$ is an arbitrary offset. Thus

$$t(\tau) = t_0 + T\tau. \quad (187)$$

This is simply rescaling and re-offsetting when going from τ space from t , as people would naively do.

The most general solution of (185) however, is:

$$\frac{dt}{d\tau} = \frac{1}{f_0(t(\tau))} := T(\tau), \quad (188)$$

such that




Figure 28 | **Analytic marginal fundamental frequency distribution** implied by the log-period prior. The distribution is lognormal with median 150 Hz, 5-95% range approximately 100-224 Hz (2.5-97.5%: 92-244 Hz).

$$t(\tau) = t_0 + \int_0^\tau T(u) du. \quad (189)$$

This defines physical time t as time-warped by local curvature $T(\tau)$, which is the instantaneous period given by the reciprocal of the instantaneous fundamental frequency $f_0(\tau) = 1/T(\tau)$. These can be indexed either in phase τ or physical time t .

The learned pitch prior therefore induces a prior over monotone warps. Sample a smooth period function $T(\tau)$, integrate it, and a physically plausible phase map follows automatically. This is why continuous-time GP models are useful here: the warp is not an ad hoc interpolation layer but the integral of a probabilistic period process.

6.3 Generative models of voiced speech

We can bring together

- time warping (the phase connection)
- learning formant tracks so PRISM can shine again, + combine with tilt and ARPrior to get better AR coeffs
- q-GPLVM quantized qpack models which have natural OQ variations (rather than learning separately). Thus we use QPACK, not g-QPACK kernels such that we can use our time warps.

Can be used to generate large quantities of voiced speech automatically

This is a surrogate model of voiced speech which can be used for generating test data or speech degradation restoration etc

The idea of data-driven voice-source modeling is not new. For example, (Thomas et al., 2009) build a codebook of prototype inverse-filtered voice-source waveforms from APLAWD. The

present route differs in two ways. First, we stay explicitly probabilistic throughout. Second, the end result is not a deterministic codebook but a mixture of low-rank Gaussian priors that can be used directly by the inference machinery of Chapter 7.

6.3.1 Generating vowels

6.4 Summary

7 Bayesian nonparametric glottal inverse filtering

7.1 Recap of all the ingredients

- IKLP
- AR prior
- Quasiperiodic arc cosine kernel
 - Learnt from synthetic Paule data
- Latent variable models for refining

7.2 Current state-of-the-art GIF algorithms

The distinction matters because our PACK, QPACK, and g-QPACK constructions are *prior* models: they specify in advance what kinds of DGF waveforms are plausible, and in the gated case they make low variance in the registered closed phase an explicit part of the covariance itself. Many modern GIF methods operate differently. Rather than building the open/closed asymmetry into the prior, they weight the data-fit term in the likelihood so that residuals in the closed phase count more heavily; the estimate is then pushed toward smaller closed-phase error, which tends to produce cleaner closed phases as a consequence of inference. That is an ingenious mechanism, but it is conceptually different from gating: it acts through observation weighting, whereas g-QPACK encodes the phase asymmetry directly in the prior.

7.2.1 Weighted LP variants

Weighted variants of LP were first introduced by (Ma et al., 1993) which are motivated to decrease coupling between source harmonics and filter resonances during GIF estimation.

CP SLP and WLP QCP

(Airaksinen et al., 2014)

In recent papers this together with IAIF is considered currently state of the art.

Based on WLP (weighted LP as the principle, not the WLP from (Chien et al., 2017)) with a specific window/weight function called AME: attenuated main excitation, is how the weight function is called, which was first proposed in the context of formant estimation of high-pitched vowels (Alku et al., 2013), and is a type of parametrized window.

Uses multiple pitch periods; how many? Uses 20-50 ms.

This is a pitch-synchronous algorithm

7.2.2 IAIF

7.2.3 CCD

7.3 Learning priors from synthetic data

7.3.1 No ground truth

The central difficulty in GIF evaluation is that the true glottal flow is not available in natural speech. As emphasized in the review of (Kadiri et al., 2021), noninvasive recording of the ground-truth glottal volume velocity is not possible in ordinary speech production, so every evaluation protocol must work around a missing reference. The standard workaround is to use synthetic data, where the generating flow is known by construction. Traditionally this has meant analytic source models such as LF; here we instead use synthetic waveforms generated through VocalTractLab and PAULE so that the learned priors are calibrated on a richer space of physiologically plausible variation.

7.3.2 PAULE and CommonVoice

PAULE is a neural model that inverts speech waveforms to the control parameter (CP) trajectories that drive VocalTractLab (VTL). In other words it attempts to reconstruct what the articulators did to arrive at loosely similar sounding speech (like children acquiring and mimicking speech). While (as its authors acknowledge) listening to the resynthesized recordings show that reconstruction quality is quite poor, our aim here is not the latter but rather exploration of human possible glottal flow, operationalized via VTL.

To do this, we forked VTL and patched it to allow extracting glottal flow and used that on the PAULE training set (ie, inverted CP trajectories which are then inputted to VTL to get the GF waveforms). This got us some 21k examples, of which most were “duds”, ie, synthesized GFs and speech that are nondescript murmured babblings.

The next step is to use LBGID on this to extract reasonable quality GCI and GOI estimates, then detect voiced regions with our pitch track model (t-PRISM), then take the DGF, killing the very large (unphysically breathiness) DC in the signal.²⁵

A direct indicator of synth quality is just synth length (ie the length of synthesized speech and GF waveforms). We found that short lengths generally sounded very similar like the same muffled grunt, and we filter these out.

We then can make accurate measurements of T and OQ. Our goal here is to cover enough variation of OQ to learn the closed phase and spectral rolloff (since latter is heavily correlated with the former). Due to synth quality we also artificially make the closed phase more perfectly closed, since this is physiologically observed and apparently non-reached (non-modeled) part of GF waveforms by PAUL.

²⁵ Although we could have also modeled it with PRISM-VFF directly by adding a RBF kernel to soak up low varying baselines. But this would then need to be unlearned by the model during BNGIF inference time.

7.3.3 Finding high quality voiced groups

We use t-PRISM for that, and our LBGID algorithm. The main conceptual idea here is that we use *an explicit pitch track model* to do the extraction rather than ad hoc heuristics.

Tricks we use:

- Normalize data to (0,1): this means we freeze mean to 0 **and** kernel variance to 1.
- Lengthscale in ballpark order: 10. This is simply to cut number of iterations, the model does find lengthscales 15 if initialized from 1, but takes much longer.
- Inverse-ECDF initialization of inducing points. This is important because trajectory lengths are long tailed, and we need to place predictive power where it matters. In 1D we can just sample from the empirical distribution, with interpolation.
- Set ν based on event statistics. Optimizing ν is possible, but more stable optimization. Results depend only very slightly on exact value of ν as long as < 10 .

For the rest standard Adam optimizer with standard settings.

Heuristic for setting ν We set ν using a simple tail-probability heuristic derived from the expected frequency of spike events. After normalization, the GP prior has unit variance and the learned observation noise standard deviation is approximately $\sigma \approx 0.15$. Typical spikes have amplitude around 5, corresponding to residuals of size $r \approx \frac{5}{0.15} \approx 30 - 35$ noise standard deviations. We interpret spikes as events that should lie in the heavy tails of the Student-t likelihood rather than being explained by the Gaussian core. Therefore we choose ν such that the tail probability of a Student-t distribution satisfies $P(|X| > r) \approx p_{\text{spike}}$, where p_{spike} is the empirical spike rate (here about 1%). Using the large- r asymptotic tail behaviour $P(|X| > r) \approx C * r^{-\nu}$ and ignoring constants gives the practical rule $\nu \approx -\frac{\log(p_{\text{spike}})}{\log(r)}$. For $r \approx 35$ and $p_{\text{spike}} \approx 0.01$, this yields $\nu \approx 1 - 2$, which corresponds to a Cauchy-like heavy-tailed likelihood. This choice ensures that extreme residuals are downweighted rather than forcing the GP to explain them through excessive curvature.

7.4 EGIFA

The dataset itself is described in Dataset 1. Here we only need its role as an evaluation benchmark. Relative to Chapter 5, the question has changed: the issue is no longer how to register EGIFA for learning, but how to compare inferred source signals against its reference waveforms and against the evaluation suite of (Chien et al., 2017). The gauge and delay issues that matter for this comparison are taken up below and in Appendix G.

7.5 Evaluation on EGIFA

Low F1 vowels cannot be resolved with GIF?

From (Chien et al., 2017): > It has been observed by some researchers that some vowels with a low first formant frequency cannot be adequately analyzed by an inverse filtering algorithm, whereas the vowel /ä/ has a first-formant frequency that is sufficiently high to avoid interference with the primarily low-frequency energy distribution of glottal source

Is this the case for us? Check if close rounded vowels /o/ and /u/ are also harder. This is case for all traditional algos

F0 dependence

Our decoupling and larger window size should also remedy F0 performance. Is that the case? Just for our best model

GCI estimates

Rather than using error injected ground truth estimates, we just use ground truth (LBGID) and actual GCI detectors, which provide for a more realistic error injection

We compare performance per GCI backend

We can also comparee GCI backend quality by testing on APLAWD

Window sizes

shorter windows do generally help for speech

for vowel relatively ok

failure modes at 0.00 quantile:

1. genuine different explanation found due to nonidentifiability => perhaps VI restarts could help here
2. F1 f0 => AR is sculpting excitation

For learning PRISM wise from EGIFA we set WIDTH=2048 This has number of cycles covered at 0.25, 0.5, 0.75, 0.99: 5, 6, 8, 9; which corresponds very well to marginal F0 distribution, so perfect

Reason of setting WIDTH low: more dense Z coverage at 18/cycle for M=128 inducing points => hoped to help inference with modeling sharp peaks

About CCD

CCD comes out really badly in (Chien et al., 2017), but relatively well in our paper Anecdotal evidence(ie, looking at actual inferred DGFs) suggests that CCD often has really good performance relative to other trad algos

The reason for that is probably that CCD is very sensitive to GCI estimation, and (Chien et al., 2017) area_gci estimator failed heavily there.

7.6 Evaluation on VTRFormants

BNGIF is not just a GIF; it's a full grey-box generative model of both voiced and nonvoiced speech, though GIF is only defined reasonably accurately for the voiced case, as the algorithms need and feed upon the correlations in the quasiperiodic glottal to resolve the inherent ambiguity in GIF (ie blind inversion where both filter and source must be inferred).

7.7 Summary

We optimize by projecting to a common basis. Memory becomes $\mathcal{O}(BMr + Ir^2)$ and compute $\mathcal{O}(Mr^2 + Ir^2 + r^3) \approx \mathcal{O}(Mr^2 + Ir^2)$ as $M \gg r$. With these optimizations in place, we can get to BNGIF speed: linear complexity in frame size. We work at native EGIFA fs of 20 kHz. Raw runtimes: in x32 runs at 5x realtime, in x64 runs at 0.5x realtime. So in x32 1 sec of audio is processed in 0.2 sec, and in x64 in 2 sec. The entire EGIFA corpus can be processed in about 3 minutes in x32 (this does not include jax tracing and loading the corpus and postprocessing results, which about double the time).

Working in x32 can sometimes result in nans due to the GIG distribution-associated computations. This happens in 1% in case of EGIFA/vowel and 0.1% in case EGIFA/speech. In the former case nan-occurrence can be reduced to below 0.001% by doing random restarts without any effect on performance. Future work could wrap the numerically sensitive parts in x64 in JAX so both speed and stability are retained.

8 Conclusion

Here is a test of mono fonts. EGIFA_raw_speech_file.

A person speaks 16,000 words on average each day: glottal folds are used plenty, so disorders arise quickly when used inappropriately. (Andrade-Miranda et al., 2024)

Recent article in Nature Parkinson's disease "Speech impairments, on the other hand, are particularly prevalent, affecting up to 90% of people with PD" "Growing evidence indicates that speech and language alterations often precede the defining motor signs and PD diagnosis **by as much as a decade**" (Cao et al., 2025) "Recent research highlights the value of **objective acoustic speech markers** in detecting PD in the initial or even prodromal (e.g., RBD) stages^{23,25,26}. This implicates a potential window for timely interventions, such as monoamine oxidase B inhibitors, anticholinergics and future novel therapies that may confer improved efficacy in early PD stages" (Cao et al., 2025)

> The recommended and widely employed tests are 1) sustained phonation of the vowel /a/ in a single breath, which helps assess breath control, vocal fold function and vocal quality (Cao et al., 2025) > 2) rapid repetition of syllables such as /pa/, /ta/ and /ka/ to evaluate consonant and vowel articulation, articulation rate and regularity, coordination and speech timing;

Important features: vocal fold oscillation irregularity, breathiness and noise, **and vocal tract resonance fluctuations**

Detection rate in clinical lab conditions is already very high But algorithms could be improved And perhaps we could have apps on mobile phone to cut costs even more; this would mean nonideal conditions so better algos needed

For example (Sasou & Chen, 2023) propose a GIF method for extracting features from Parkinsonian speech, but conclude that, while their method is effective, it is "calculation-intensive and time-consuming". BNGIF is ofcourse much slower than direct LP based methods, but it can still run at 5x realtime speed on an old RTX 2080, so with modern hardware it can be made essentially as fast as one pleases. Another advantage of using JAX is that the code was written in Python, has the advantage of the scripting vibe and mature ecosystem, while the code can be exported to many devices, and for each device optimized. It can be embedded e.g. as a C++ library with specific optimization for the processor without any work done, via the MLIR system. Although all this is experimental, it shows the appeal of turning signal-processing-based more-heuristic approaches into modern ML. Another advantage is end-to-end differentiation, which enables automated large scale turning and learning.

Online learning Priors for the expansion coefficients $\mathbf{a}$ can be updated from the posterior from previous frame. Same for the AR prior $\mathbf{a}$ (nameclash here). Likewise, values of $\{T, t_c\}$ are correlated across frames and this can be exploited too.

8.1 Future work

We put a probabilistic model handle on the GIF problem, so our assumptions are made explicit and can be criticized and improved (MacKay, 2005).

On the one hand, grey box approach means: no true large scale data collection needed, and we have tried to show how to make most of theoretical knowledge, each nat of constraint built in requiring another $O(\dots)$ less data. On the other, collecting a high quality dataset from synthesized glottal flow does not hurt either, and can be used both for more realistic learning and benchmarking; and is high priority. Ideally, of course, flow can be measured directly or via proxy by high speed area imaging or by neck contact sensors.

Holistic optimization IKLP is fully differentiable end-to-end (though will require x64 for that) if CAVI updates are made locally. So future work should just optimize parameters directly on it, rather than out multi-stage pipeline.

Learning the time warp from data It would be better to ditch LBGID and learn it from data. But scaling would be an issue. (Kazlauskaitė et al., 2018).

IKLP extensions To KARMA and mixture of AR priors

q-GPLVM q-GPLVM as used here is a very stable two-step procedure, but it would an interesting research direction to include the GMM objective directly into the Bayesian GPLVM ELBO. This would enable “activation-aware” quantization.

Spectral priors for ARMA Extend spectral priors for MA and ARMA models. The linear divisibility machinery itself is agnostic: it works just as well for numerator polynomials, hence for antiresonances, notches, and radiation factors in ARMA or pole-zero models. What fails here is only the **all-pole interpretation** of roots of A as dips in $|H|^2$.

Generate testing databases Use of AR prior and surrogate GF can generate large quantities of test data for other algorithms. Currently synthetic datasets such as (Alku et al., 2019) are typically generated from LF model, but this isn’t too realistic. For example (Airaksinen et al., 2014) determined AME weight function parameters from a corpus of synthetic vowels generated from LF model by varying hyperparameters on a grid: (LF parameters $\times$ F0 $\times$ poles in a 8 pole AR) All of these can be improved by models developed in this paper: (surrogate DGF source $\times$ F0 from pitch tracks $\times$ poles from higher order AR with just specifying the formants with our AR prior) This allows huge quantities of more realistic data to be generated nonparametrically while knowing ground truth prefectly. Indeed when using VocalTractLab ground truth like F0 or GCIs must be derived from the signal (as we did with LBGID), and generating large quantities of data is computationally intensive – eg 60k samples used in (Airaksinen et al., 2014) – which our method can do extremely efficient due to its surrogate nature: just sample DGF as a Gaussian-weighted comibnaton of basis functions, sample an AR prior, and filter.

Other Possible improvements/future directions:

- $p(a)$ not as a delta function in IKLP
- work in spectral domain completely
- model $h(t)$ as a GP conditional on $e(t)$ given: alternate optimization
- interframe correlations as in (Mehta et al., 2012)
- learn from more realistic glottal flow simulations like (Avhad et al., 2022; Schoder et al., 2024; Zhang et al., 2020); perhaps learn first a mapping from glottal flow area to glottal flow from physics and then use real life measured glottal flow area databases like (Andrade-Miranda et al., 2024)

Datasets

Dataset 1. EGIFA

EGIFA is the synthetic benchmark introduced by (Chien et al., 2017). It combines radiated speech waveforms with matched glottal-flow waveforms generated by a high-fidelity articulatory speech production model. In this thesis it serves two distinct purposes. In Chapter 5 it provides the corpus on which the PACK family is calibrated. Later, in Chapter 7 it serves as the controlled benchmark on which inferred source signals can be compared against reference glottal flow.

This makes EGIFA unusually valuable. Unlike natural-speech corpora, it gives access not only to the observable speech waveform but also to the source waveform we would actually like to learn. At the same time, that does not make the data ready-made for learning. The raw recordings still vary in duration, phase, and phonatory stability, and the gated kernels of Chapter 5 need a common phase convention before their prior can mean the same thing across the corpus.

1.1. Registration for learning

For learning, we turn the raw EGIFA recordings into a corpus of clean voiced groups in shared cycle time. The four steps are the same ones only summarized in Chapter 5:

1. detect GCI and GOI landmarks from the glottal flow waveform;
2. fit a robust pitch track and retain only the steady voiced stretches;
3. warp each surviving group from physical time t to normalized time τ ;
4. smooth, differentiate, and standardize the result to obtain DGF waveforms $u'(\tau)$.

The first two steps are handled by LBGID and t-PRISM, whose algorithmic details are given in Algorithm 1 and Algorithm 3. Here the main point is only their division of labor. LBGID finds candidate glottal landmarks from the flow signal itself. t-PRISM then decides which stretches of the resulting period sequence are consistent enough to keep. The output is not a cycle-by-cycle annotation but a set of voiced groups whose pitch evolution is stable enough for learning.

The time warp itself is intentionally simple. For ungated models the GCIs are mapped to the integers,

$$\tau(t) = \text{interp}(t; n_{\text{GCI},0} \mapsto 0, n_{\text{GCI},1} \mapsto 1, n_{\text{GCI},2} \mapsto 2, \dots), \quad (190)$$

so that one unit of τ corresponds to one glottal cycle. For gated models, the GOIs are additionally mapped to half-integers,

$$n_{\text{GOI},k} \mapsto k + 1/2, \quad (191)$$




Figure 29 | **Registration of one retained EGIFA voiced group.** Top: the raw speech waveform $s(t)$ in physical time, extracted from the indicated .wav file and voiced group. Middle: the corresponding raw glottal flow $u(t)$, again in physical time, with LBGID-detected GCI and GOI landmarks. Bottom: the same group after differentiation, amplitude normalization and registration to normalized time, now shown as $u'(\tau)$. In this representation the GCIs land at integers $\tau = 0, 1, 2, \dots$ and the GOIs at half-integers $\tau = 0.5, 1.5, 2.5, \dots$

so that the open/closed boundary always occurs at the same registered phase. This is not a claim that the physical open quotient is $1/2$. It is only a phase convention that lets the gated kernels compare like with like across the corpus.

After registration, the flow is smoothed on a physiologically meaningful timescale of 0.1 ms, differentiated, and standardized within group. At the native EGIFA sampling rate of 20 kHz, the associated resampling step lands essentially back at the original rate, so the operation acts mainly as regularization rather than downsampling.

EGIFA does not provide ground-truth GCIs or GOIs. Still, the detected landmarks imply cycle lengths and open quotients, namely

$$T_{0,k} = \frac{n_{\text{GCI},k+1} - n_{\text{GCI},k}}{f_s}, \quad F_{0,k} = \frac{1}{T_{0,k}}, \quad \text{OQ}_k = \frac{n_{\text{GOI},k} - n_{\text{GCI},k}}{n_{\text{GCI},k+1} - n_{\text{GCI},k}}. \quad (192)$$




Figure 30 | **The EGIFA registration pipeline on one recording.** (a) A short raw-flow window with LBGID-detected GCI and GOI events. (b) The corresponding period sequence over the full recording, together with the t-PRISM posterior mean and uncertainty band; periods with weight above one are retained as inliers and contiguous retained runs become voiced groups. (c) One retained voiced group after GOI-aware registration, smoothing, differentiation, and per-group standardization, shown as a waveform in shared cycle time. The lightly shaded half-cycles indicate the phase later suppressed by the gated kernels.

Aggregated over the extracted EGIFA frames, these induced values look sensible. In particular, the F_0 modes line up well with the nominal pitch settings of the corpus, and the implied open quotients stay in a plausible range.




Figure 31 | **Implied F_0 and OQ distributions from LBGID landmarks on EGIFA.** Left: the frame-level F_0 distribution implied by the detected GCIs, with dashed lines indicating EGIFA's nominal pitch settings. Right: the corresponding frame-level OQ distribution implied by the detected GCI/GOI pairs.

On EGIFA this yields 5612 registered voiced groups in total. The calibration experiments of Chapter 5 use 5000 for training and 500 for held-out evaluation. This is the corpus on which PRISM learns one shared low-rank basis.

1.2. Benchmark Caveats

EGIFA also reappears later as an evaluation benchmark, but there the emphasis changes. The question is no longer how to register the data for learning, but how to compare inferred source signals against the reference waveforms in a way that respects the underlying ambiguities of GIF. The original evaluation suite of (Chien et al., 2017) is useful and still worth reproducing, but its time-domain scores are highly sensitive to delay and affine alignment. That is why the thesis keeps the learning use of EGIFA separate from the later evaluation use. The former needs a clean registered corpus; the latter needs careful treatment of gauge and delay, as discussed in Appendix G and Chapter 7.

Dataset 2. APLAWD

APLAWD is a speech corpus with synchronized speech, electroglottograph (EGG), differentiated EGG (DEGG), and hand-corrected GCI reference markings (Naylor et al., 2007). It has been used both for data-driven voice-source modeling (Thomas et al., 2009) and for large-scale GCI detection evaluation (Serwy, 2017). Serwy reports a hand-verified markings database over 10,944 waveforms, which is exactly what makes the corpus useful here: APLAWD supplies trustworthy landmark sequences from which a prior over smooth pitch-period trajectories can be learned.

In this thesis APLAWD is not a glottal-flow learning corpus. It enters only through the geometry of pitch. The pitch-track prior of Algorithm 3 is trained on APLAWD and then reused later for voiced-group selection on EGIFA and CommonVoice. That role is narrower than EGIFA’s role in Chapter 5, but it is also cleaner. For learning a pitch prior, accurate GCI landmarks matter more than matched source waveforms.

The manual markings are important because the quantity of interest is not the speech waveform itself but the inter-GCI geometry it induces. For recording n with GCI sample indices $n_{\text{GCI},n,0}, \dots, n_{\text{GCI},n,M_n}$ at sampling rate f_s , define the period sequence

$$T_{nm} = 1000 \frac{n_{\text{GCI},n,m+1} - n_{\text{GCI},n,m}}{f_s} \quad (m = 0, \dots, M_n - 1), \quad (193)$$

measured in milliseconds. We then work with the log-periods




Figure 32 | **One APLAWD excerpt with DEGG landmarks.** Top: the acoustic waveform. Bottom: the synchronized DEGG trace together with the hand-corrected GCI references and QuickGCI estimates. The database provides this kind of alignment throughout the corpus, which is why it is the right place to learn a pitch prior.




Figure 33 | **Two APLAWD diagnostics.** Left: QuickGCI timing errors against the hand-corrected GCI references. The mode is centered close to the expected larynx-to-microphone delay. Right: local absolute jitter measured over retained voiced groups. Most of the mass lies well below the usual pathology threshold.

$$\ell_{nm} = \log_{10} T_{nm}, \quad (194)$$

which stabilize the scale and turn multiplicative pitch variation into additive variation. The training data for Algorithm 3 is simply the collection of these sequences over the corpus, after a global whitening transform.

2.1. Landmark quality

Because APLAWD contains both speech and hand-corrected GCI references, it also provides a natural place to sanity-check the geometry we later ask the pitch model to explain. Serwy’s Hilbert-phase detectors already showed that QuickGCI is close to the stronger GADFLI detector on this corpus (Serwy, 2017). For the thesis the main point is simpler: the detector errors are tightly localized, and the local jitter measured over voiced groups remains in a healthy range.

This does not prove that every landmark is perfect, and it need not. The thesis only needs a corpus where the dominant pitch geometry is trustworthy enough to learn a smooth prior from it. APLAWD meets that standard.

2.2. Learning the pitch prior

The goal is to learn a single GP prior that captures how pitch typically rises and falls over a voiced stretch of speech. This is not a source model. It is a model of period trajectories. The target is therefore the collection of log-period series (194), not the speech waveform.

We train the model with t-PRISM from Algorithm 3. The Student-t likelihood is useful here because even in APLAWD some period sequences contain long silent gaps, detector slips, or locally erratic closures. The robust weights absorb such observations instead of forcing the prior to explain them. Model selection is done by held-out log likelihood per effective test point.

Kernel	M	Held-out log likelihood	Lengthscale	Obs. std.
rbf	64	$2.117 + -0.012$	15.25	0.137
matern:52	64	$2.070 + -0.009$	17.35	0.144
rbf	32	$2.051 + -0.011$	15.28	0.147
matern:32	64	$2.029 + -0.008$	18.65	0.148
matern:52	32	$2.008 + -0.010$	17.23	0.153

Table 8 | **Held-out model selection on APLAWD.** Values are mean test log likelihood per effective test point with the uncertainty estimate from `experiments/svi/aplawd/stats.R`.

The winning model is an RBF prior with $M = 64$ inducing points.

The result is clean. An RBF prior with $M = 64$ inducing points wins decisively enough to become the default pitch-track prior used throughout the thesis. The learned lengthscale is about 15 cycles, which is the quantity of real interest here: it sets the timescale on which the model expects pitch to vary smoothly. Neither the rational-quadratic alternative nor the learned- ν Student-t extension improves on this winner in held-out score.

2.3. Voiced-group segmentation

Once learned, the pitch prior can be turned around and used as a robust filter on new GCI-annotated period tracks. Given a period sequence $\{T_{nm}\}_{m=0}^{M_n-1}$, t-PRISM returns local weights

$$w_{nm} = \mathbb{E}[\lambda_{nm}], \quad (195)$$

and in this thesis we use the conservative rule

$$w_{nm} > 1 \quad (196)$$

to declare period m an inlier. Contiguous inlier runs are then taken as voiced groups. This is the segmentation rule later reused on EGIFA and CommonVoice.

This is where the APLAWD chapter meets the rest of the thesis. APLAWD provides the corpus on which the pitch prior is learned and selected. Later chapters reuse that prior as a robust notion of normal voiced pitch, but they do not need to repeat its training story.




Figure 34 | **Voiced-group selection on one APLAWD pitch track.** Top: the observed inter-GCI periods together with the t-PRISM posterior mean and uncertainty band. Bottom: the inferred weights $w_k = \mathbb{E}[\lambda_k]$. Weight collapses carve the period track into retained voiced runs and rejected gaps.

Dataset 3. CommonVoice

The CommonVoice corpus used here is a derived paired-waveform dataset: each recording is stored both as speech and as a matched glottal-flow waveform. Relative to EGIFA, the material is longer, less orderly, and closer to the large-corpus setting for which the registration pipeline was ultimately designed.

In the present thesis CommonVoice is mainly a scaling target. The LBGID + t-PRISM + registration pipeline is first developed and stress-tested on EGIFA, where the synthetic control is tighter. Later, the same machinery is reused on CommonVoice to construct a larger registered corpus for source modeling beyond the benchmark setting.

(Gao et al., 2024; Sun et al., 2022)

(Birkholz, 2013; Krug et al., 2021)

Dataset 4. VTRFormants

VTRFormants (Deng et al., 2006) is the TIMIT-derived corpus of hand-corrected formant and bandwidth tracks, distributed as frame-wise FB files aligned with the original TIMIT waveforms and phone labels (Garofolo et al., 1993).

Its role is to provide reference resonance trajectories for later evaluation. EGIFA asks whether a model can recover glottal source structure when reference flow is available. VTRFormants asks whether the same probabilistic machinery yields useful formant estimates on real speech.

4.1. Spectral tilt and TIMIT

The same TIMIT vowel material used to construct the reference formant tracks also gives a useful estimate of spectral rolloff.²⁶ What we want is not just the tilt of voiced speech itself, but an estimate of the tilt contributed by the *vocal tract filter*. I have not been able to find direct empirical estimates of that quantity in the literature.

Now, broadband slope is not uniquely attributable in source-filter decompositions: depending on the assumed glottal model and smoothing procedure, some of the slope can be assigned to the source and some to the filter (Fant, 1960). Nevertheless, physical models of the glottal cycle and lip radiation do imply a certain range of plausible tilt for these stages of speech production. We know that:

1. the glottal flow $u(t)$ itself has a rolloff of about -12 dB/oct for abrupt closure, and can be steeper, up to about -18 dB/oct, for smoother return phases (Doval et al., 2006);
2. lip radiation contributes about $+6$ dB/oct (Schroeder, 1999);
3. what remains to be estimated is the filter-side term; a quantity that should be interpreted as an effective residual tilt rather than a directly observed property of the vocal tract, but still not entirely unidentifiable;
4. these add up to the spectral tilt of the observed speech signal.

So by measuring 4 and subtracting 1 and 2, we can estimate 3.

Estimating the spectral tilt of voiced speech can be done with the marked vowel segments on which the VTRFormants reference tracks are based. So the data are essentially quasi-periodic vowels.

One could try to estimate rolloff from two frequencies an octave apart, but that is too noisy in practice because the FFT power spectrum is strongly modulated by harmonic structure. Instead, we smooth by regression. For each voiced vowel segment, we estimate the FFT magnitude spectrum of the speech signal in dB and fit the linear trend

$$P(f) = a \log_2 f + b. \quad (197)$$

²⁶Note that ‘rolloff’, ‘spectral tilt’, and ‘spectral slope’ all loosely mean the same thing, although some authors reserve rolloff for asymptotic high-frequency decay and tilt for an overall spectral slope.

Then

$$P(2f) - P(f) = a. \quad (198)$$

so the coefficient a itself is the estimated rolloff in dB per octave.

The fit is only performed above a cutoff of about 200 Hz. The reason is that the additive rolloff picture only really makes sense beyond the glottal-formant region; below that, the low-frequency source structure dominates. Both Doval et al. (2006) and Fulop & Disner (2011) point to about 200 Hz as a reasonable cutoff for this regime.

Figure 35 shows four representative fits. Averaging over the full set of usable TIMIT vowel segments, the resulting voiced-speech tilt is about $-8(2)$ dB per octave. Combining this with the source-side heuristics above gives an effective residual filter tilt of about $-2(2)$ dB per octave.




Figure 35 | **Example of spectral tilt fits on TIMIT vowel segments.** Each panel shows the empirical power spectrum together with the fitted trend $a \log_2 f + b$. The dashed vertical line marks the 200 Hz cutoff above which the fit is performed. The four examples span a representative range from steeper to flatter observed rolloff.

This is consistent with mild tilt as predicted by lossy acoustic tubes (de Boer, 2008) and with Chen (2013)'s theoretical estimate of the decay rate of formants (Eq. 24) which is -3 dB/oct, as opposed to Fant's suggestion of -6 dB/oct (Eq. 26).

Algorithms

Algorithm 1. LBGID

LBGID stands for *level-based glottal instant detection*.

Glottal closure and opening instants (GCIs and GOIs) define the pseudoperiodicity of speech and are important speech features with many practical applications (Drugman et al., 2012). Within the thesis, GCIs and GOIs are used as temporal anchors for the kernels modeling the glottal cycle. Learning these kernels from example $\{u(t)\}$ glottal flow waveforms requires these instants to be known. The task LBGID attempts to solve is to extract consistent (GCI, GOI) pairs from a given $u(t)$, which can then be used downstream for aligning a set of $\{u(t)\}$ waveforms.

In a physics-driven high-quality synthesizer like VocalTractLab (Birkholz, 2013; Krug et al., 2021), glottal closure is not given as a symbolic event and often does not appear as a sharp isolated spike. Instead, it appears as an emergent soft low-flow regime embedded in a continuous waveform whose baseline, amplitude, and degree of breathiness can all vary over time. LBGID extracts the glottal instants based on the following assumption: closure happens where the flow $u(t)$ is below a certain threshold, and the glottal instants are the strongest *opposing changes* in $u(t)$ that flank it.

1.1. Adaptive envelope normalization

A difficulty with glottal airflow signals is that absolute amplitude is rarely interpretable on its own. Slow baseline drift and physical effects like variation in subglottal pressure all alter amplitude to various degrees. To measure effective flow rather than simulator artifacts, LBGID tracks adaptive upper and lower envelopes, denoted by $\text{roof}(t)$ and $\text{floor}(t)$, by detecting peaks and troughs and interpolating between them with a standard peak picking algorithm. These define a smoothly varying local amplitude range,

$$A(t) = \text{roof}(t) - \text{floor}(t) \geq 0, \quad (199)$$

which sets the local scale for the threshold. Note that we assume the polarity of $u(t)$ to be positive; that is, $u(t) \geq 0$ everywhere.

1.2. Level-based closed phase detection

With the local amplitude frame in hand, candidate closed phases are identified via a simple relative threshold,

$$\text{level}(t) = \text{floor}(t) + \alpha(\text{roof}(t) - \text{floor}(t)), \quad (200)$$

for a small constant α , and points satisfying $u(t) \leq \text{level}(t)$ are treated as belonging to a closed phase. This yields contiguous low-flow segments that naturally partition the signal into bounded candidate intervals. In the thesis we use $\alpha = 0.01$ on EGIFA and $\alpha = 0.10$ on CommonVoice. The latter is needed because the inferred glottal flow there tends to retain a




Figure 36 | **LBGID on EGIFA and CommonVoice**. Each panel shows $u(t)$ together with $\text{roof}(t)$, $\text{floor}(t)$, the relative threshold $\text{level}(t)$, and the detected (GCI, GOI) pairs. On EGIFA we use $\alpha = 0.01$; on CommonVoice we use $\alpha = 0.10$. Note how the level function rejects the non-voiced part (roughly before 300 ms) in the bottom panel.

more persistent breathy floor. The three functions in (200) are illustrated in Figure 36 and Figure 37.

1.3. Dynamic pairing inside the closed phase

Within each low-flow segment, LBGID searches for a pair of maximally opposed changes in $u(t)$. First, the derivative $u'(t)$ is computed by finite differences. Then algorithm evaluates all pairs of sample indices (i, j) , where $i < j$, within the segment by the score

$$E(i, j) = -u'(i) u'(j). \quad (201)$$

Among these admissible ordered pairs, the one maximizing $E(i, j)$ is selected. The sample indices of the optimal pair are then output as the glottal closure instant (GCI) and glottal opening instant (GOI), respectively. Equation (201) is thus just a cheap proxy for quantifying the “abruptness” of a candidate glottal instant pair.

Figure 36 shows the result on an excerpt from EGIFA (Dataset 1) and one from CommonVoice (Dataset 3). EGIFA has better defined closed phases than CommonVoice due to synthesis




Figure 37 | **Zoom-in of a single closed phase.** Top: $u(t)$ together with $\text{roof}(t)$, $\text{floor}(t)$ and $\text{level}(t)$, with the selected low-flow region shaded. Bottom: $u'(t)$ on the same window, with the chosen GCI and GOI marked.

quality (even boasting handcrafted control trajectories for the synthesizer) and the choice of α is relatively inconsequential. By contrast, CommonVoice is much more challenging, and the choice of α must balance accuracy (which requires α small) and false positives like splitting up closed phases when a bump happens in the middle (which happens when α is small).

Figure 37 shows a more detailed view of the algorithm for a given located closed phase. The algorithm first finds the candidate closed phase by amplitude, and then locates the strongest negative-to-positive dynamic pair inside it.

1.4. Evaluation

There is no direct ground truth for the GCI and GOI estimates themselves, except for manually checking every output (Serwy, 2017). On EGIFA, a proxy for LBGID’s performance is simply to look at the implied quantities

$$F_{0,k} = \frac{f_s}{n_{\text{GCI},k+1} - n_{\text{GCI},k}}, \quad \text{OQ}_k = \frac{n_{\text{GOI},k} - n_{\text{GCI},k}}{n_{\text{GCI},k+1} - n_{\text{GCI},k}}, \quad (202)$$

whose distributions over frames in EGIFA are shown in Figure 31. These look reasonable and show a healthy amount of variation.

1.5. Summary

LBGID estimates (GCI, GOI) pairs from given $u(t)$ waveforms via level-based glottal instant detection.

It does this with two heuristics: closed phases correspond to intervals where flow is relatively small, and the GCI and GOI events happen with “maximally opposite abruptness.”

These simple rules appear to be sufficient for the registration of the EGIFA and CommonVoice datasets.

Algorithm 2. PRISM

PRISM stands for *process-induced surrogate modeling*.

PRISM is an adapter from variable-length datasets to Gaussians in a common amplitude space. The idea is not to fit one GP per waveform and stop there, but to learn a single shared kernel-induced basis and express each dataset in that finite-dimensional basis with calibrated uncertainty. Each dataset is modeled by its own latent function corrupted independently by globally scaled noise.

The name signifies the following. “Process-induced” means that each observation is indexed by an arbitrary input grid and may have arbitrary length, so no ad-hoc resampling or interpolation is required. “Surrogate modeling” refers to the fact that the learned basis defines a cheaper probabilistic surrogate for the original collection of functions, to be optionally refined or compressed by the user.

In Chapter 5 this shared-basis view is used for *calibration*: one global kernel is fit to a large registered EGIFA corpus of over 7k waveforms and 6.6M samples. Later chapters use the same representation for downstream surrogate modeling. PRISM is therefore best thought of as the bridge from datasets of functions to tractable Gaussian coordinates in an amplitude space of fixed dimension R .

The mechanism that makes this tractable is an extension of the classic collapsed variational bound proposed by Titsias (2009). For Gaussian likelihoods, the local inducing-variable posteriors can be optimized analytically and substituted back into the objective, so only global quantities need to be trained. Projection to per-dataset Gaussians then becomes a cheap post-training step. This chapter develops that construction explicitly.

2.1. Data and model

We have N independent datasets. Dataset n consists of an input grid

$$\mathbf{t}_n = (t_{n1}, \dots, t_{nM_n}) \quad (203)$$

and observations

$$\mathbf{y}_n = (y_{n1}, \dots, y_{nM_n}), \quad (204)$$

where M_n may vary freely with n . For simplicity the inputs are written as scalar times, but nothing below depends on that; the same construction works for arbitrary input dimension. Indeed, any model compatible with Titsias (2009) can be handled by PRISM. An important restriction is the assumption of a Gaussian likelihood, which is what makes the local states collapse analytically.

Each dataset is modeled by an independent latent function

$$f_n(\cdot) \sim \text{GaussianProcess}(0, k_\theta(\cdot, \cdot)), \quad (205)$$

sharing the same kernel hyperparameters θ across all n but with independent draws of f_n . The kernel k_θ in (205) ties the draws together and acts as a global regularizer in function space.²⁷ The likelihood is given by

$$\mathbf{y}_n \mid f_n \sim \text{Normal}(f_n(\mathbf{t}_n), \sigma^2 \mathbf{I}_{M_n}) \quad (206)$$

which expresses the assumption that observations are conditionally independent given f_n , and imposes additional regularization through the global noise scale σ .

2.2. From kernel to shared basis

Choose R inducing locations

$$\mathbf{z} = (z_1, \dots, z_R) \quad (207)$$

in the input domain and treat them as global learnable parameters. More general inducing *features* are also possible (Wilk et al., 2020), but vanilla inducing point evaluations are convenient for more general nonstationary kernels such as g-QPACK (177). For dataset n , define the inducing variables

$$\mathbf{u}_n = f_n(\mathbf{z}) \in \mathbb{R}^R. \quad (208)$$

Since f_n is a GP,

$$\mathbf{u}_n \sim \text{Normal}(\mathbf{0}, \mathbf{K}_{zz}), \quad (209)$$

where $\mathbf{K}_{zz} \in \mathbb{R}^{R \times R}$ with $(\mathbf{K}_{zz})_{rj} = k_\theta(z_r, z_j)$ is the same for all n . The GP conditional on $\mathbf{u}_n$ is

$$f_n(\mathbf{t}_n) \mid \mathbf{u}_n \sim \text{Normal}(\mathbf{K}_{nz} \mathbf{K}_{zz}^{-1} \mathbf{u}_n, \mathbf{K}_{nn} - \mathbf{Q}_{nn}), \quad (210)$$

where $\mathbf{K}_{nz} \in \mathbb{R}^{M_n \times R}$ has (m, r) entry $k_\theta(t_{nm}, z_r)$ and

$$\mathbf{Q}_{nn} = \mathbf{K}_{nz} \mathbf{K}_{zz}^{-1} \mathbf{K}_{zn} \quad (211)$$

is the Nyström approximation to $\mathbf{K}_{nn}$.

We emphasize that the column space of $\mathbf{K}_{nz}$ is shared across all datasets because the inducing locations $\mathbf{z}$ are shared. That is the common basis PRISM learns. The local variables $\mathbf{u}_n$ only say how dataset n “sits” in that basis.

2.3. Collapsed variational bound

The cost of keeping local parameters A direct sparse variational treatment would introduce one Gaussian variational factor per dataset,

$$q_n(\mathbf{u}_n) = \text{Normal}(\boldsymbol{\mu}_n, \boldsymbol{\Sigma}_n), \quad (212)$$

²⁷Since functions are infinite dimensional, some prior notion of smoothness is typically needed to make the problem well-posed (Sivia & Skilling, 2006).

and maximize a separate ELBO for every n . That is workable for a single dataset, but for a large corpus it means carrying a synchronized local Gaussian state for every dataset during training. PRISM avoids this by keeping the local variational problem only long enough to solve it analytically, and then substituting the solution back into the objective.

Collapsed bound Following (Titsias, 2009), the per-dataset variational bound can be written as

$$\begin{aligned} \mathcal{L}_n(q_n) = \mathbb{E}_{q_n} \left[\log \text{Normal}(\mathbf{y}_n \mid \mathbf{K}_{nz} \mathbf{K}_{zz}^{-1} \mathbf{u}_n, \sigma^2 \mathbf{I}_{M_n}) \right] \\ - D_{\text{KL}}(q_n(\mathbf{u}_n) \parallel p(\mathbf{u}_n)) - \frac{1}{2\sigma^2} \text{tr}(\mathbf{K}_{nn} - \mathbf{Q}_{nn}). \end{aligned} \quad (213)$$

For Gaussian likelihoods the optimal factor is available in closed form:

$$q_n(\mathbf{u}_n) = \text{Normal}(\boldsymbol{\mu}_n, \boldsymbol{\Sigma}_n), \quad (214)$$

with

$$\mathbf{C}_n = \mathbf{K}_{zz} + \sigma^{-2} \mathbf{K}_{zn} \mathbf{K}_{nz}, \quad (215)$$

$$\boldsymbol{\mu}_n = \sigma^{-2} \mathbf{K}_{zz} \mathbf{C}_n^{-1} \mathbf{K}_{zn} \mathbf{y}_n, \quad \boldsymbol{\Sigma}_n = \mathbf{K}_{zz} \mathbf{C}_n^{-1} \mathbf{K}_{zz}. \quad (216)$$

Substituting this optimum back into (213) gives the collapsed Titsias bound

$$\mathcal{L}_n(\mathbf{z}, \theta, \sigma^2) = \log \text{Normal}(\mathbf{y}_n \mid \mathbf{0}, \mathbf{Q}_{nn} + \sigma^2 \mathbf{I}_{M_n}) - \frac{1}{2\sigma^2} \text{tr}(\mathbf{K}_{nn} - \mathbf{Q}_{nn}). \quad (217)$$

The first term is the low-rank marginal likelihood and the second penalizes the discarded covariance.²⁸

The PRISM ELBO is then just the sum over independent datasets,

$$\mathcal{L}(\mathbf{z}, \theta, \sigma^2) = \sum_{n=1}^N \mathcal{L}_n(\mathbf{z}, \theta, \sigma^2), \quad (218)$$

which is seen to depend only on the global parameters $(\mathbf{z}, \theta, \sigma^2)$.

Evaluating the bound In practice we whiten the basis so the local prior becomes isotropic. Let $\mathbf{K}_{zz} = \mathbf{L}\mathbf{L}^\top$ and define the amplitude coordinates

$$\mathbf{a}_n = \mathbf{L}^{-1} \mathbf{u}_n, \quad (219)$$

so that $\mathbf{a}_n \sim \text{Normal}(\mathbf{0}, \mathbf{I}_R)$. The corresponding feature map is

$$\boldsymbol{\varphi}(t) = \mathbf{L}^{-1} \mathbf{k}_z(t), \quad \mathbf{k}_z(t) = (k_\theta(z_1, t), \dots, k_\theta(z_R, t))^\top, \quad (220)$$

²⁸The newer bounds of (Titsias, 2025) can also be used here. For BNGIF they did not yield better models in practice and were less efficient to batch, so we stayed with the original bound of (Titsias, 2009).

and the design matrix for dataset n is

$$\Phi_n = \begin{pmatrix} \varphi(t_{n1})^\top \\ \vdots \\ \varphi(t_{nM_n})^\top \end{pmatrix} \in \mathbb{R}^{M_n \times R}. \quad (221)$$

Then the sparse GP mean becomes a BLR mean,

$$\mathbf{y}_n = \Phi_n \mathbf{a}_n + \varepsilon_n, \quad \varepsilon_n \sim \text{Normal}(\mathbf{0}, \sigma^2 \mathbf{I}_{M_n}). \quad (222)$$

Define

$$\mathbf{A}_n = \frac{\Phi_n^\top}{\sigma} \in \mathbb{R}^{R \times M_n}, \quad \mathbf{B}_n = \mathbf{I}_R + \mathbf{A}_n \mathbf{A}_n^\top, \quad \mathbf{v}_n = \mathbf{A}_n \mathbf{y}_n \in \mathbb{R}^R. \quad (223)$$

Both terms in (217) are then computed efficiently from $\mathbf{A}_n$, $\mathbf{B}_n$, and $\mathbf{v}_n$, so the cost per dataset scales as $\mathcal{O}(R^2 M_n)$ and the $M_n \times M_n$ matrix $\mathbf{Q}_{nn} + \sigma^2 \mathbf{I}$ is never formed.

For the log marginal term, write $\mathbf{Q}_{nn} + \sigma^2 \mathbf{I} = \sigma^2 (\mathbf{I}_{M_n} + \mathbf{A}_n^\top \mathbf{A}_n)$. By the matrix determinant lemma,

$$\log \det(\mathbf{Q}_{nn} + \sigma^2 \mathbf{I}) = M_n \log \sigma^2 + \log \det(\mathbf{B}_n). \quad (224)$$

The Woodbury identity gives

$$(\mathbf{Q}_{nn} + \sigma^2 \mathbf{I})^{-1} = \sigma^{-2} (\mathbf{I}_{M_n} - \mathbf{A}_n^\top \mathbf{B}_n^{-1} \mathbf{A}_n), \quad (225)$$

so the quadratic form is

$$\mathbf{y}_n^\top (\mathbf{Q}_{nn} + \sigma^2 \mathbf{I})^{-1} \mathbf{y}_n = \sigma^{-2} (\|\mathbf{y}_n\|^2 - \mathbf{v}_n^\top \mathbf{B}_n^{-1} \mathbf{v}_n). \quad (226)$$

For the trace correction, $\text{tr}(\mathbf{Q}_{nn}) = \|\Phi_n\|_F^2$, so

$$-\frac{1}{2\sigma^2} \text{tr}(\mathbf{K}_{nn} - \mathbf{Q}_{nn}) = -\frac{1}{2\sigma^2} (\text{tr}(\mathbf{K}_{nn}) - \|\Phi_n\|_F^2). \quad (227)$$

Every quantity in (217) is thus computable from the global parameters and the current dataset alone.

2.4. Stochastic optimization

Because the datasets are independent, an unbiased estimate of (218) is available from any minibatch $\mathcal{B} \subset \{1, \dots, N\}$,

$$\hat{\mathcal{L}} = \frac{N}{|\mathcal{B}|} \sum_{n \in \mathcal{B}} \mathcal{L}_n(\mathbf{z}, \theta, \sigma^2). \quad (228)$$

The minibatch unit is a *complete dataset*, not an individual observation. This is what keeps irregular grids and unequal lengths harmless: each dataset contributes one collapsed term, regardless of its internal geometry. In practice we optimize (228) with Adam following (Orvieto & Gower, 2025) at a learning rate of 10^{-3} . We found this stable and amenable to float32 optimization, which gives a substantial speed gain over float64.

The contrast with the ELBO of stochastic variational inference (Hensman et al., 2013, Eq. 3) is interesting here. There, stochastic optimization over very large datasets is obtained by introducing and maintaining an explicit variational state of size $\mathcal{O}(R^2)$ in order to have the objective factorize over datapoints (and allow more general likelihoods beyond (206)). We avoid that move here because the factorization is already over complete datasets. The local inducing posteriors are collapsed analytically and only the global parameters $(\mathbf{z}, \theta, \sigma^2)$ are updated.

2.5. Projection to the PRISM representation

Once $(\mathbf{z}, \theta, \sigma^2)$ are trained, PRISM maps any dataset of arbitrary length M_n to a Gaussian in the fixed amplitude space $\mathbb{R}^R$. At the fitted global parameters, the local inducing posterior is simply the optimal factor (216) evaluated for the new dataset:

$$\mathbf{u}_n \mid \mathbf{y}_n \sim \text{Normal}(\boldsymbol{\mu}_n, \boldsymbol{\Sigma}_n). \quad (229)$$

Passing to amplitude coordinates via $\mathbf{a}_n = \mathbf{L}^{-1}\mathbf{u}_n$ gives

$$\mathbf{a}_n \mid \mathbf{y}_n \sim \text{Normal}(\mathbf{m}_n, \mathbf{S}_n), \quad (230)$$

$$\mathbf{S}_n = \mathbf{B}_n^{-1}, \quad \mathbf{m}_n = \mathbf{B}_n^{-1}\mathbf{v}_n. \quad (231)$$

This reuses $\mathbf{B}_n$ and $\mathbf{v}_n$ from the collapsed bound, so projection is essentially for free.

To summarize, the learned PRISM representation consists of the shared basis (220) together with the projected Gaussian coordinates

$$\mathcal{R}_{\text{PRISM}} = \left(\boldsymbol{\varphi}, \{(\mathbf{m}_n, \mathbf{S}_n)\}_{n=1}^N \right). \quad (232)$$

This is a collection of Gaussian blobs in a common $\mathbb{R}^R$ together with the basis that generated them, and therefore defines an implicit density in amplitude space suitable for downstream processing.

2.6. Surrogate modeling

Importantly, PRISM alone does not yet produce a good *global* surrogate model of the corpus. What it learns is the regularized basis (220) together with the Gaussian coordinates (232) of the observed data in that basis. This is already enough to define a set of *local* Gaussian surrogate posteriors. For dataset n , one can sample amplitudes

$$\mathbf{a}_n^* \sim \text{Normal}(\mathbf{m}_n, \mathbf{S}_n) \quad (233)$$

cheaply, push them through the synthesis map $\Phi_n \mathbf{a}_n^*$, and thereby resample from the local surrogate posterior induced by the shared basis. The posterior covariances are what make this cheap resampling possible.

To obtain a better corpus-level surrogate model, extra structure must be learned on top of the Gaussian collection in (232). One may ignore the covariances and apply ordinary tools such as PCA, clustering, or Gaussian mixtures to the posterior means. Or one may keep the covariances and use errors-in-variables (uncertain-input) variants of those same ideas. For very large collections one generally also wants to quantize or compress the number of Gaussian blobs, for example by hierarchical clustering or by the latent-variable model of the next chapter. In this thesis we take the latter route and hand the PRISM representation off to q-GPLVM (Algorithm 4), which learns a more compact generative model in the same shared amplitude space.

2.7. Summary

PRISM learns a shared, global kernel-induced basis from a possibly large collection of independent variable-length time series without storing or updating per-dataset variational parameters. In terms of GP approximation, PRISM is a reduced-rank projection of the GP prior k_θ to a rank- R BLR. The infinite-dimensional function space is projected onto the span of the learned basis, and the remaining degrees of freedom are absorbed by the observation noise σ^2 .

After training, each dataset is projected to a Gaussian amplitude posterior (230) in a common, R -dimensional amplitude space. Together with the shared basis (220), this defines an implicit density in that space that can be used as-is for cheap surrogate sampling or passed to more structured downstream models.

Algorithm 3. t-PRISM

The Gaussian likelihood in PRISM (Algorithm 2) is analytically convenient, but inappropriate when data contain outliers. For learning a pitch-track model from APLAWD (Dataset 2), for example, the observations are inter-GCI period measurements, where a single detector failure, silent gap, or irregular closure can create an isolated gross outlier. More generally, any functional series may contain observations that should be downweighted rather than fitted exactly.

To address this, we introduce t-PRISM here as the “robust” version of PRISM. It replaces the Gaussian likelihood by a Student-t while retaining the same collapsed basis-learning structure. Tractability is maintained by representing the Student-t as a Gaussian scale mixture with one local precision variable per observation. Conditional on those precisions the likelihood is Gaussian again, so the collapsed bound from Algorithm 2 applies with the single modification that the noise covariance becomes diagonally weighted. The result is again a minibatch algorithm over complete datasets, and like PRISM it outputs Gaussian amplitude posteriors in a shared space.

Figure 38 illustrates the principle. Under a Gaussian likelihood, the posterior mean is pulled toward the locally irregular periods because the model must explain every residual at full strength. t-PRISM instead preserves the smooth trend and incorporates outliers by reducing their corresponding local weights $w_m \ll 1$.




Figure 38 | **t-PRISM on a local APLAWD period track.** The top panel compares PRISM and t-PRISM on one excerpt of inter-GCI periods from the APLAWD pitch-track corpus, using the learned pitch track model of Dataset 2. The bottom panel shows the inferred precision weights $w_m = \mathbb{E}[\lambda_m]$. Weights $w_m \geq 1$ mark observations consistent with the local trend; smaller weights $w_m < 1$ signal outliers. Note that t-PRISM ignores the outlier at $m = 37.5$.

3.1. Student-t as a scale mixture

Recall the PRISM setup from Algorithm 2. Given our kernel of choice k_θ , we are presented with N independent datasets $\{(\mathbf{t}_n, \mathbf{y}_n)\}_{n=1}^N$, where dataset n has length M_n , latent function

$$f_n(\cdot) \sim \text{GaussianProcess}(0, k_\theta(\cdot, \cdot)), \quad (234)$$

and inducing variables $\mathbf{u}_n = f_n(\mathbf{z})$ at shared inducing locations $\mathbf{z}$. In PRISM this prior is paired with the Gaussian observation model

$$\mathbf{y}_n \mid f_n \sim \text{Normal}(f_n(\mathbf{t}_n), \sigma^2 \mathbf{I}_{M_n}), \quad (235)$$

as in (206). t-PRISM leaves this prior side untouched. The only change is that the homoscedastic Gaussian likelihood is replaced by a locally weighted one: we introduce one latent precision variable per observation and integrate it out to obtain a Student-t likelihood.

The Student-t likelihood with degrees of freedom $\nu > 0$ and scale σ^2 is recovered from the standard Normal-Gamma augmentation

$$\lambda_{nm} \sim \text{Gamma}(\nu/2, \nu/2), \quad y_{nm} \mid f_{nm}, \lambda_{nm} \sim \text{Normal}(f_{nm}, \sigma^2/\lambda_{nm}). \quad (236)$$

Marginalizing over λ_{nm} recovers $\text{StudentT}(y_{nm} \mid f_{nm}, \nu, \sigma^2)$ exactly. For dataset n , stack the per-sample precisions into

$$\mathbf{\Lambda}_n = \text{diag}(\lambda_{n1}, \dots, \lambda_{nM_n}). \quad (237)$$

The augmented likelihood is then Gaussian with diagonal but heteroscedastic noise covariance:

$$\mathbf{y}_n \mid f_n, \mathbf{\Lambda}_n \sim \text{Normal}(f_n(\mathbf{t}_n), \sigma^2 \mathbf{\Lambda}_n^{-1}). \quad (238)$$

Large λ_{nm} means that sample (n, m) is trusted; small λ_{nm} means that its residual is down-weighted. The prior mean of every λ_{nm} is one, so if the data contain no clear outliers the posterior expected precisions remain close to one. Then $\mathbf{W}_n \approx \mathbf{I}$ below, and the weighted formulas reduce continuously back to the PRISM expressions.

Figure 39 shows the intended fallback behavior. When there is no clear evidence for contamination, robustness costs essentially nothing: t-PRISM simply behaves like PRISM.

3.2. Variational family and objective

We introduce a mean-field variational family (Blei et al., 2017) over the latent precisions,

$$q_n(\mathbf{\Lambda}_n) = \prod_{m=1}^{M_n} q_{nm}(\lambda_{nm}), \quad q_{nm}(\lambda_{nm}) = \text{Gamma}(\alpha_{nm}, \beta_{nm}), \quad (239)$$

while the GP part is kept exactly as in the PRISM derivation. That is, for fixed precisions we still introduce a Gaussian factor $q_n(\mathbf{u}_n)$ over the inducing variables, use the exact GP




Figure 39 | **t-PRISM reducing to PRISM in the absence of outliers.** On this local APLAWD period track the observations are already consistent with a smooth trend. The PRISM and t-PRISM fits are nearly indistinguishable, and all inferred weights $w_m \geq 1$ and $O(1)$.

conditional $p(f_n \mid \mathbf{u}_n)$, and collapse the local Gaussian state analytically as in (213) and (217). The per-dataset ELBO before collapse is

$$\begin{aligned} \mathcal{L}_n^t = & \mathbb{E}_{q_n}[\log p(\mathbf{y}_n \mid f_n, \mathbf{\Lambda}_n)] - D_{\text{KL}}(q_n(\mathbf{u}_n) \parallel p(\mathbf{u}_n)) \\ & - \sum_{m=1}^{M_n} D_{\text{KL}}(q_{nm}(\lambda_{nm}) \parallel p(\lambda_{nm})). \end{aligned} \quad (240)$$

Define the expected precision and expected log-precision under q_{nm} :

$$\begin{aligned} w_{nm} &= \mathbb{E}[\lambda_{nm}] = \alpha_{nm} / \beta_{nm} \\ \ell_{nm} &= \mathbb{E}[\log \lambda_{nm}] = \psi(\alpha_{nm}) - \log \beta_{nm}, \end{aligned} \quad (241)$$

where ψ is the digamma function (MacKay & Bauman Peto, 1994).

3.3. Weighted collapsed variational bound

Given fixed expected precisions, write

$$\mathbf{W}_n = \text{diag}(w_{n1}, \dots, w_{nM_n}). \quad (242)$$

Conditional on $\mathbf{W}_n$, the likelihood is Gaussian with noise covariance $\sigma^2 \mathbf{W}_n^{-1}$. This is exactly the PRISM setting with weighted observation noise, so the collapsed bound applies again.

Define the weighted matrices

$$\begin{aligned}
\mathbf{A}_n^W &= \frac{1}{\sigma} \boldsymbol{\Phi}_n^\top \mathbf{W}_n^{1/2} \in \mathbb{R}^{R \times M_n} \\
\mathbf{B}_n^W &= \mathbf{I}_R + \mathbf{A}_n^W (\mathbf{A}_n^W)^\top \in \mathbb{R}^{R \times R} \\
\mathbf{v}_n^W &= \mathbf{A}_n^W \mathbf{W}_n^{1/2} \mathbf{y}_n \in \mathbb{R}^R.
\end{aligned} \tag{243}$$

In the unweighted case $\mathbf{W}_n = \mathbf{I}$ these reduce exactly to $\mathbf{A}_n$, $\mathbf{B}_n$, and $\mathbf{v}_n$ from (223).

The collapsed Gaussian contribution becomes

$$\mathcal{L}_n^{\text{coll}}(\mathbf{W}_n) = \log \text{Normal}(\mathbf{y}_n \mid \mathbf{0}, \mathbf{Q}_{nn} + \sigma^2 \mathbf{W}_n^{-1}) - \frac{1}{2\sigma^2} \text{Tr}(\mathbf{W}_n(\mathbf{K}_{nn} - \mathbf{Q}_{nn})) \tag{244}$$

The weighted predictive moments needed below are

$$\begin{aligned}
m_{nm} &= \boldsymbol{\varphi}(t_{nm})^\top (\mathbf{B}_n^W)^{-1} \mathbf{v}_n^W, \\
v_{nm} &= k(t_{nm}, t_{nm}) - \boldsymbol{\varphi}(t_{nm})^\top \left(\mathbf{I} - (\mathbf{B}_n^W)^{-1} \right) \boldsymbol{\varphi}(t_{nm}).
\end{aligned} \tag{245}$$

The only new part is the expected log-likelihood under the random precisions. At sample (n, m) , conditioning on λ_{nm} gives a Gaussian variance σ^2/λ_{nm} . Its normalization therefore contributes

$$-\frac{1}{2} \log \left(\frac{\sigma^2}{\lambda_{nm}} \right) = -\frac{1}{2} \log \sigma^2 + \frac{1}{2} \log \lambda_{nm}, \tag{246}$$

so after expectation the log-precision enters through $\ell_{nm} = \mathbb{E}[\log \lambda_{nm}]$. The quadratic penalty is weighted by $w_{nm} = \mathbb{E}[\lambda_{nm}]$. Substituting the moments from (245) gives, up to constants,

$$\mathbb{E}[\log p(\mathbf{y}_n \mid f_n, \boldsymbol{\Lambda}_n)] = -\frac{1}{2} \sum_{m=1}^{M_n} \left(\log \sigma^2 - \ell_{nm} + \frac{w_{nm}}{\sigma^2} ((y_{nm} - m_{nm})^2 + v_{nm}) \right) \tag{247}$$

and the full per-dataset t-PRISM ELBO is

$$\begin{aligned}
\mathcal{L}_n^t(\mathbf{W}_n, \ell_n) &= \mathcal{L}_n^{\text{coll}}(\mathbf{W}_n) + \frac{1}{2} \sum_{m=1}^{M_n} \ell_{nm} \\
&\quad - \sum_{m=1}^{M_n} D_{\text{KL}}(q_{nm}(\lambda_{nm}) \parallel \text{Gamma}(\nu/2, \nu/2)).
\end{aligned} \tag{248}$$

Local CAVI updates for the precision variables Differentiating (248) gives the optimal values for tightening the ELBO:

$$\alpha_{nm}^* = (\nu + 1)/2, \quad \beta_{nm}^* = \frac{1}{2} \left(\nu + ((y_{nm} - m_{nm})^2 + v_{nm})/\sigma^2 \right). \tag{249}$$

The shape α is fixed once ν is chosen and does not depend on the data. In the deployed APLAWD model of this thesis we use the conservative choice $\nu = 1$, giving deliberately heavy tails for occasional gross detector failures or silent gaps. As noted in Dataset 2, a learned- ν extension was also tried, but did not improve held-out score over this simpler fixed choice. Only the rate β changes, increasing when the current residual $(y_{nm} - m_{nm})^2 + v_{nm}$ is large relative to σ^2 . Large residuals therefore drive w_{nm} down in the next weighted collapse. That is mechanism driving “robustness.”

3.4. Training

For a minibatch $\mathcal{B}$ of complete datasets, the stochastic objective is

$$\widehat{\mathcal{L}}^t = \frac{N}{|\mathcal{B}|} \sum_{n \in \mathcal{B}} \mathcal{L}_n^t. \quad (250)$$

For each dataset n in the minibatch:

1. initialize $w_{nm} = 1$ for all m ,
2. repeat for a small fixed number of steps:
 - run the weighted collapsed GP computation with current $\mathbf{W}_n$, obtaining (m_{nm}, v_{nm}) via (245),
 - update $(\alpha_{nm}, \beta_{nm})$ via (249) and refresh (w_{nm}, ℓ_{nm}) via (241),
3. evaluate $\mathcal{L}_n^t$ via (248).

Then take a gradient step in the global parameters $(\mathbf{z}, \theta, \sigma^2, \nu)$. The Gamma parameters are local to the current minibatch pass and discarded afterwards. As in PRISM, no large per-dataset state persists across iterations. In the experiments of this thesis we use a fixed three inner CAVI steps, do not differentiate through those local updates, and optimize the global parameters with the same Adam setup and float32 arithmetic as in PRISM.

3.5. Projection to the PRISM representation

After training, each dataset is projected to a Gaussian amplitude posterior exactly as in Algorithm 2.5., now using the final converged weights from the inner loop. The posterior formulas are weighted versions of (230) and (231):

$$\begin{aligned} \mathbf{a}_n \mid \mathbf{y}_n &\sim \text{Normal}(\mathbf{m}_n, \mathbf{S}_n) \\ \mathbf{S}_n &= (\mathbf{B}_n^W)^{-1} \\ \mathbf{m}_n &= (\mathbf{B}_n^W)^{-1} \mathbf{v}_n^W. \end{aligned} \quad (251)$$

This again reuses quantities already computed during the inner loop. The resulting robust representation is therefore

$$\mathcal{R}_{\text{t-PRISM}} = \left(\boldsymbol{\varphi}, \{(\mathbf{m}_n, \mathbf{S}_n)\}_{n=1}^N, \{\mathbf{w}_n\}_{n=1}^N \right), \quad (252)$$

the same kind of fixed-dimensional Gaussian amplitude collection as (232). The weight vectors $\mathbf{w}_n$ can be retained as auxiliary diagnostics when one wants local inlier scores in the original observation domain.

3.6. Summary

t-PRISM relaxes the Gaussian noise assumption of PRISM, which allows it to “robustly” learn a shared basis for a set of exemplars contaminated with outliers. This is achieved by replacing the Gaussian likelihood with a Student-t likelihood written as a Gaussian scale mixture.

Conditional on the local precision variables, the model is Gaussian again, so the same collapsed bound can be reused with weighted noise covariance $\sigma^2 \mathbf{W}_n^{-1}$. The local CAVI updates (249) then shrink the weights of observations whose residuals are too large for the current smooth fit. Training therefore remains a minibatch procedure over complete datasets, and the final output remains a Gaussian amplitude posterior per dataset.

In this thesis the fitted weight vectors $\mathbf{w}_n$ are used both indirectly, through the robust amplitude posteriors, and directly, as local inlier scores on period trajectories. The motivation for developing t-PRISM came from needing to solve the following problems:

- learning a pitch-track model from the pitch trajectory exemplars of APLAWD in Dataset 2;
- and segmenting pitch trajectories into voiced groups for the EGIFA and CommonVoice datasets (Dataset 1 and Dataset 3, respectively).

Algorithm 4. q-GPLVM

As we established in Algorithm 2, PRISM projects a collection of variable-length waveforms into Gaussians in a shared amplitude space $\mathbb{R}^R$. More precisely, the PRISM representation (232) is

$$\mathcal{R}_{\text{PRISM}} = (\varphi, \{(\mathbf{m}_n, \mathbf{S}_n)\}_{n=1}^N), \quad (253)$$

where the shared feature map $\varphi(t)$ projects waveform n from data space to amplitude space through the posterior

$$\mathbf{a}_n \mid \mathbf{y}_n \sim \text{Normal}(\mathbf{m}_n, \mathbf{S}_n). \quad (254)$$

This defines the implicit density in amplitude space

$$p_{\text{PRISM}(\mathbf{a})} = \frac{1}{N} \sum_{n=1}^N \text{Normal}(\mathbf{a} \mid \mathbf{m}_n, \mathbf{S}_n), \quad (255)$$

which is not yet a compact global generative model. The density (255) remains tied to the training set, and scoring a new waveform against it still requires keeping all N components around.

q-GPLVM compresses (255) into a compact generative prior. It first learns a Bayesian GPLVM in a lower-dimensional manifold space, then quantizes the aggregated manifold posterior into a small Gaussian mixture, and finally pushes that mixture forward through the learned GP map back to amplitude space. The result is a mixture of K Gaussians in $\mathbb{R}^R$ that summarizes the training distribution without reference to individual training waveforms.

q-GPLVM stands for *quantized Gaussian process latent variable model*. The name refers to the quantization step: the GPLVM posterior, which lives on a continuous manifold, is discretized into a finite mixture. It is a variation on the Bayesian GPLVM of Titsias & Lawrence (2010) with two additions that matter in this thesis: first, the model accepts diagonal amplitude uncertainty inherited from PRISM after one global whitening transform; second, Gaussian components in manifold space are pushed through the learned GP by moment matching, using the same Gaussian kernel integrals that underlie the variational bound (Damianou et al., 2014).

4.1. Data, amplitude, and manifold space

The q-GPLVM pipeline has three levels, and keeping their inputs and outputs distinct avoids most notation drift.

Level 0: data space This is the space of observed waveforms. The input to PRISM is the collection

$$\mathcal{D} = \{(bm(t)_n, \mathbf{y}_n)\}_{n=1}^N, \quad (256)$$

with waveform n sampled at its own grid $\mathbf{t}_n$ and values $\mathbf{y}_n$. After all compression steps, q-GPLVM returns a prior that again lives in this same space.

Level 1: amplitude space PRISM maps each waveform to a Gaussian in amplitude space $\mathbb{R}^R$. Its shared feature map $\varphi(t)$ ((220) in Algorithm 2) defines the linear synthesis model

$$y(t) \approx \varphi(t)^\top \mathbf{a}, \quad \mathbf{a} \sim \text{Normal}(\mathbf{0}, \mathbf{I}), \quad (257)$$

so Level 1 is summarized exactly by (255). This is a *linear* dimensionality reduction: variable-length waveforms become Gaussian amplitude coordinates in fixed dimension R .

Level 2: manifold space The Bayesian GPLVM introduces manifold coordinates $\mathbf{x}_n \in \mathbb{R}^Q$ with $Q \ll R$ and learns the posterior

$$q(\mathbf{X}) = \prod_{n=1}^N \text{Normal}(\mathbf{x}_n \mid \boldsymbol{\mu}_n^x, \mathbf{S}_n^x). \quad (258)$$

This is a *nonlinear* dimensionality reduction from amplitude space to manifold space. Quantization then replaces the aggregated posterior by a small mixture

$$q_{K(\mathbf{x})} = \sum_{k=1}^K \pi_k \text{Normal}(\mathbf{x} \mid \mathbf{c}_k, \mathbf{C}_k), \quad (259)$$

which is later pushed back to Level 1 and then composed with PRISM to return to Level 0.

4.2. The Bayesian GPLVM

The theoretical foundation is the Bayesian GP latent variable model of Titsias & Lawrence (2010), which we now review.

Observations and notation PRISM returns full amplitude covariances $\mathbf{S}_n$, whereas the modified Bayesian GPLVM used here accepts known *diagonal* observation noise in amplitude space. We therefore apply one global affine whitening transform

$$\tilde{\mathbf{a}} = \mathbf{W}(\mathbf{a} - \mathbf{b}), \quad (260)$$

chosen so that the transformed PRISM covariance family is close to diagonal on average. If $(\tilde{\mathbf{m}}_n, \tilde{\mathbf{S}}_n)$ denotes the whitened posterior of waveform n , we keep only the diagonal part

$$\mathbf{d}_n = \text{diag}(\tilde{\mathbf{S}}_n) \in \mathbb{R}^R, \quad (261)$$

and, to avoid a forest of tildes, continue to write $\mathbf{a}_n$ for the whitened posterior mean.

Collect these means into the amplitude matrix

$$\mathbf{A} \in \mathbb{R}^{N \times R}, \quad (262)$$

whose n -th row is $\mathbf{a}_n^\top$. Thus waveform n enters the GPLVM not as a point estimate alone, but as the pair $(\mathbf{a}_n, \mathbf{d}_n)$ in amplitude space.

Model

For each amplitude coordinate $r = 1, \dots, R$, introduce an independent latent function $g_r(\cdot)$ with shared GP prior

$$g_r(\cdot) \sim \text{GaussianProcess}(0, k(\cdot, \cdot)), \quad (263)$$

and an observation model

$$a_{nr} \mid g_r, \mathbf{x}_n \sim \text{Normal}(g_r(\mathbf{x}_n), \beta^{-1} + d_{nr}), \quad (264)$$

where $\mathbf{x}_n \in \mathbb{R}^Q$ is the manifold coordinate of waveform n , β^{-1} is the learned isotropic GPLVM noise level, and d_{nr} is the known diagonal amplitude uncertainty inherited from PRISM. The original model of (Titsias & Lawrence, 2010) is recovered by setting $\mathbf{d}_n = \mathbf{0}$ for all n . The latent prior is isotropic:

$$p(\mathbf{x}_n) = \text{Normal}(\mathbf{0}, \mathbf{I}_Q). \quad (265)$$

The kernel is ARD squared exponential,

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp\left(-\frac{1}{2} \sum_{q=1}^Q \alpha_q (x_q - x'_q)^2\right), \quad (266)$$

with per-dimension inverse lengthscales α_q . When the model drives $\alpha_q \rightarrow 0$ during training, dimension q is effectively switched off, providing automatic selection of the intrinsic latent dimensionality (Titsias & Lawrence, 2010). In the EGIFA q-GPLVM fit used throughout the thesis, all three manifold dimensions remain active, though to different degrees as seen in Figure 40.

Variational inference

Integrating out the manifold coordinates from the marginal likelihood is intractable because they enter the kernel matrix nonlinearly. Following (Titsias & Lawrence, 2010), introduce P inducing inputs $\mathbf{Z} \in \mathbb{R}^{P \times Q}$ in manifold space and optimize the factorized Gaussian posterior (258), with diagonal latent covariance

$$\mathbf{S}_n^x = \text{diag}(S_{n1}^x, \dots, S_{nQ}^x). \quad (267)$$

For point observations ($\mathbf{d}_n = \mathbf{0}$), the collapsed variational bound has the standard form

$$\mathcal{F}(q) = \tilde{\mathcal{F}}(q) - D_{\text{KL}(q(\mathbf{X}) \parallel p(\mathbf{X}))}, \quad (268)$$

where $\tilde{\mathcal{F}}(q)$ is the sparse-GP expected log-likelihood and the KL term is analytic. The inducing variables are again collapsed analytically, exactly as in Algorithm 2.




Figure 40 | **Representative training behavior and ARD selection on EGIFA.** Left: ELBO trace for one typical q-GPLVM run. Right: the learned inverse lengthscales in manifold space. The optimization stabilizes cleanly, and the ARD kernel keeps all three manifold dimensions active, though not equally: the third carries the sharpest variation and the second is the smoothest.

Our diagonal-uncertainty extension leaves this collapsed structure intact. Write $\mathbf{D}_n = \text{diag}(\mathbf{d}_n)$ and let w_n denote the n -th diagonal entry of the collapsed precision matrix in data space. Then the modified bound is

$$\mathcal{F}_{\text{diag}(q)} = \mathcal{F}(q) - \frac{1}{2} \sum_{n=1}^N w_n \text{tr}(\mathbf{D}_n). \quad (269)$$

This is exactly the additional term obtained by taking the expectation of the point-observation bound under the known Gaussian amplitude uncertainty. It is not part of the original model of (Titsias & Lawrence, 2010); it is the thesis extension that lets q-GPLVM consume uncertain PRISM amplitudes directly. No new variational factors are introduced, and the kernel expectations remain those of (Titsias & Lawrence, 2010). The variational parameters $\{\boldsymbol{\mu}_n^x, \mathbf{S}_n^x\}_{n=1}^N$, the inducing inputs $\mathbf{Z}$, and the kernel and noise parameters $\{\sigma_f^2, \alpha_1, \dots, \alpha_Q, \beta\}$ are therefore still optimized jointly by gradient ascent on the same collapsed objective, now augmented by (269).

The Ψ statistics

All tractable computation in the Bayesian GPLVM reduces to these expectations of kernel quantities under $q(\mathbf{x}_n)$:

$$\begin{aligned}
\psi_0 &= \sum_{n=1}^N \mathbb{E}_{q(\mathbf{x}_n)} [k(\mathbf{x}_n, \mathbf{x}_n)], \\
[\Psi_1]_{np} &= \mathbb{E}_{q(\mathbf{x}_n)} [k(\mathbf{x}_n, \mathbf{z}_p)], \\
[\Psi_2]_{pp'} &= \sum_{n=1}^N \mathbb{E}_{q(\mathbf{x}_n)} [k(\mathbf{x}_n, \mathbf{z}_p) k(\mathbf{x}_n, \mathbf{z}_{p'})],
\end{aligned} \tag{270}$$

with $\Psi_1 \in \mathbb{R}^{N \times P}$ and $\Psi_2 \in \mathbb{R}^{P \times P}$. Note that Ψ_2 is *not* $\Psi_1^\top \Psi_1$: the integration couples the two kernel evaluations through the shared uncertain input $\mathbf{x}_n$.

For the ARD RBF kernel (266) with $q(\mathbf{x}_n) = \text{Normal}(\boldsymbol{\mu}_n, \mathbf{S}_n^x)$, all three statistics admit closed forms. The first statistic is trivial: $\psi_0 = N\sigma_f^2$ since $k(\mathbf{x}, \mathbf{x}) = \sigma_f^2$ for any $\mathbf{x}$. For Ψ_1 , the Gaussian integral yields

$$[\Psi_1]_{np} = \sigma_f^2 \prod_{q=1}^Q (\alpha_q S_{nq}^x + 1)^{-1/2} \exp\left(-\frac{\alpha_q (\mu_{nq} - z_{pq})^2}{2(\alpha_q S_{nq}^x + 1)}\right), \tag{271}$$

which is a kernel evaluation at $\boldsymbol{\mu}_n$ with inflated lengthscales $\alpha_q^{-1} + S_{nq}^x$. When $\mathbf{S}_n^x \rightarrow \mathbf{0}$ (no input uncertainty), this recovers the standard kernel evaluation $k(\boldsymbol{\mu}_n, \mathbf{z}_p)$. Ψ_2 has a similar closed form involving the midpoint and distance between inducing pairs $(\mathbf{z}_p, \mathbf{z}_{p'})$ (Titsias & Lawrence, 2010). During training this diagonal-Gaussian case is the only one needed. Later, in the push-forward step, the quantized manifold components carry full covariances $\mathbf{C}_k$, and then we use the corresponding full-Gaussian version of the same integrals (Damianou et al., 2014).

Two properties that matter downstream: first, all three statistics decompose as sums over n , so stochastic minibatching is straightforward (Lalchand et al., 2022); second, the same Gaussian-integral mechanism extends to any Gaussian input covariance, not only the diagonal $\mathbf{S}_n^x$ used in training. This is what makes the push-forward tractable.

4.3. Quantization

After training, the Bayesian GPLVM provides a variational posterior $q(\mathbf{x}_n) = \text{Normal}(\boldsymbol{\mu}_n, \mathbf{S}_n^x)$ for each training waveform. Its aggregated posterior

$$q(\mathbf{x}) = 1/N \sum_{n=1}^N q(\mathbf{x}_n) \tag{272}$$

is already a generative model in manifold space, but with N components it still encodes the entire training set. Quantization replaces it by the compact K -component approximation ($K \ll N$)

$$q(\mathbf{x}) \approx \sum_{k=1}^K \pi_k \text{Normal}(\mathbf{x} \mid \mathbf{c}_k, \mathbf{C}_k), \quad \sum_{k=1}^K \pi_k = 1, \tag{273}$$

where $\mathbf{c}_k \in \mathbb{R}^Q$ and $\mathbf{C}_k \in \mathbb{R}^{Q \times Q}$ are the mean and full covariance of component k . Full covariances are retained because local curvature and anisotropy in manifold space are precisely the structure we want to preserve.

Extreme deconvolution

The variational posteriors $q(\mathbf{x}_n)$ are not point estimates: each carries a covariance $\mathbf{S}_n^x$ that quantifies genuine uncertainty about the latent position. A standard GMM fit to the posterior means $\boldsymbol{\mu}_n$ alone would ignore this uncertainty and overfit to noise in the mean estimates. *Extreme deconvolution* (XDGMM) (Bovy et al., 2011) is a variant of EM for Gaussian mixtures that accounts for known, heterogeneous measurement uncertainties on the inputs. It fits (273) to the collection $\{(\boldsymbol{\mu}_n, \mathbf{S}_n^x)\}_{n=1}^N$ by treating each $\boldsymbol{\mu}_n$ as a noisy observation of the true manifold coordinate with known covariance $\mathbf{S}_n^x$.

The background component

Training data inevitably contain outliers: waveforms that PRISM projected poorly, or that lie outside the main phonation modes. To prevent these from distorting the learned mixture, we initialize component $k = 0$ as a broad background Gaussian whose mean and covariance are set to the empirical mean and covariance of all $\boldsymbol{\mu}_n$. The remaining $K - 1$ components are initialized from k -means on the posterior means. During EM, the background component absorbs waveforms that do not fit any well-defined cluster, and its weight π_0 is generally observed to become negligible.

The output of the quantization step is a small parametric object $\{\pi_k, \mathbf{c}_k, \mathbf{C}_k\}_{k=1}^K$ that stands in for the entire training collection. On EGIFA, the learned manifold is already visibly organized by open quotient, and the fitted ellipsoids then provide a compact cover of that structure, as shown in Figure 41.

4.4. Push-forward

We now have a GMM in manifold space and a trained GP map from $\mathbb{R}^Q$ to amplitude space $\mathbb{R}^R$. Each manifold component must therefore be pushed through the GPLVM. This can be viewed in two ways: either one samples from the quantized manifold mixture and propagates those samples through the nonlinear map directly, or one first constructs the moment-matched Gaussian prior in amplitude space and only then synthesizes waveforms from that approximation.

Recovered inducing parameters

During training, the inducing variables $\mathbf{u}_r$ were eliminated analytically from the bound. For the push-forward we need them back. Let $\mathbf{a}^r \in \mathbb{R}^N$ denote the r -th column of the amplitude matrix $\mathbf{A}$. From the optimal variational posterior (Titsias, 2009) applied within the GPLVM,

$$\begin{aligned}\hat{\mathbf{S}}_u &= \mathbf{K}_{ZZ}(\mathbf{K}_{ZZ} + \beta \boldsymbol{\Psi}_2^{\text{tr}})^{-1} \mathbf{K}_{ZZ}, \\ \hat{\mathbf{m}}_r &= \beta \mathbf{K}_{ZZ}(\mathbf{K}_{ZZ} + \beta \boldsymbol{\Psi}_2^{\text{tr}})^{-1} [\boldsymbol{\Psi}_1^{\text{tr}}]^\top \mathbf{a}^r,\end{aligned}\tag{274}$$




Figure 41 | **Manifold structure and quantization on EGIFA.** Left: the three-dimensional manifold space learned by q-GPLVM, with each waveform colored by open quotient. Right: the $K = 5$ fitted Gaussian ellipsoids used for quantization in that same manifold space. The left panel shows that open quotient varies smoothly over the manifold rather than appearing as disconnected clusters, while the right panel shows how XDGMM compresses that continuous posterior into a small mixture.

where Ψ_1^{tr} and Ψ_2^{tr} are the training-time Ψ statistics. The posterior covariance $\hat{\mathbf{S}}_u \in \mathbb{R}^{P \times P}$ is shared across all output dimensions; the posterior means are collected into $\hat{\mathbf{M}} = [\hat{\mathbf{m}}_1, \dots, \hat{\mathbf{m}}_R] \in \mathbb{R}^{P \times R}$.

Moment matching

For one manifold component $\mathbf{x} \sim \text{Normal}(\mathbf{c}, \mathbf{C})$, the induced distribution of the amplitude vector

$$\mathbf{a} = (g_1(\mathbf{x}), \dots, g_R(\mathbf{x}))^\top \quad (275)$$

under the GP is non-Gaussian due to the nonlinearity of the map. The standard approach is moment matching: choose the Gaussian $\text{Normal}(\boldsymbol{\mu}_a, \mathbf{S}_a)$ whose first two moments equal those of the true push-forward.

The required expectations are the same Ψ statistics of (270), now evaluated at the GMM component $(\mathbf{c}, \mathbf{C})$ rather than at a variational posterior $(\boldsymbol{\mu}_n, \mathbf{S}_n^x)$. Since the ARD RBF formula (271) remains valid for any Gaussian input, including one with full covariance $\mathbf{C}$, no new derivations are needed.²⁹ Write

²⁹With a full $\mathbf{C}$, the product over dimensions in (271) generalizes to the matrix expression $|\boldsymbol{\Lambda}^{-1}\mathbf{C} + \mathbf{I}|^{-1/2} \exp\left(-1/2(\mathbf{c} - \mathbf{z}_p)^\top (\boldsymbol{\Lambda} + \mathbf{C})^{-1}(\mathbf{c} - \mathbf{z}_p)\right)$, where $\boldsymbol{\Lambda} = \text{diag}(\alpha_1^{-1}, \dots, \alpha_Q^{-1})$.

$$\begin{aligned}
\psi_0 &= \mathbb{E}_{\mathbf{x}}[k(\mathbf{x}, \mathbf{x})], \\
\psi_1 &\in \mathbb{R}^P \text{ with } [\psi_1]_p = \mathbb{E}_{\mathbf{x}}[k(\mathbf{x}, \mathbf{z}_p)], \\
\mathbf{\Psi}_2 &\in \mathbb{R}^{P \times P} \text{ with } [\mathbf{\Psi}_2]_{pp'} = \mathbb{E}_{\mathbf{x}}[k(\mathbf{x}, \mathbf{z}_p)k(\mathbf{x}, \mathbf{z}_{p'})],
\end{aligned} \tag{276}$$

for the per-component statistics, suppressing the dependence on $(\mathbf{c}, \mathbf{C})$.

4.4.2.1. Mean

The predictive mean for amplitude coordinate r at a known point $\mathbf{x}$ is $\hat{\mathbf{m}}_r^\top \mathbf{K}_{ZZ}^{-1} \mathbf{k}_Z(\mathbf{x})$, where $\mathbf{k}_Z(\mathbf{x}) \in \mathbb{R}^P$ is the cross-covariance vector to the inducing inputs. Taking the expectation over $\mathbf{x} \sim \text{Normal}(\mathbf{c}, \mathbf{C})$ replaces $\mathbf{k}_Z(\mathbf{x})$ by its mean ψ_1 :

$$\boldsymbol{\mu}_a = \widehat{\mathbf{M}}^\top \mathbf{K}_{ZZ}^{-1} \psi_1 \in \mathbb{R}^R. \tag{277}$$

4.4.2.2. Covariance

The law of total variance gives

$$\mathbf{S}_a = \underbrace{\mathbb{E}_{\mathbf{x}}[\text{Cov}(\mathbf{a} \mid \mathbf{x})]}_{\text{averaged predictive variance}} + \underbrace{\text{Cov}_{\mathbf{x}}(\mathbb{E}[\mathbf{a} \mid \mathbf{x}])}_{\text{input uncertainty}}. \tag{278}$$

The first term averages the GP predictive variance (which is the same for every amplitude coordinate r since all outputs share the kernel) under the input distribution. Since outputs are conditionally independent given $\mathbf{x}$, this term is isotropic:

$$\mathbb{E}_{\mathbf{x}}[\text{Cov}(\mathbf{a} \mid \mathbf{x})] = \gamma \mathbf{I}_R, \tag{279}$$

where

$$\gamma = \psi_0 + \beta^{-1} - \text{Tr}(\mathbf{K}_{ZZ}^{-1} \mathbf{\Psi}_2) + \text{Tr}(\mathbf{K}_{ZZ}^{-1} \hat{\mathbf{S}}_u \mathbf{K}_{ZZ}^{-1} \mathbf{\Psi}_2). \tag{280}$$

The four contributions are respectively: the prior variance at $\mathbf{x}$, the observation noise, minus the variance explained by the inducing inputs, plus the posterior uncertainty on those inducing inputs.

The second term in (278) captures how the predictive mean varies as $\mathbf{x}$ moves over the component. It introduces cross-output correlations via

$$\text{Cov}_{\mathbf{x}}(\mathbb{E}[\mathbf{a} \mid \mathbf{x}]) = \widehat{\mathbf{M}}^\top \mathbf{K}_{ZZ}^{-1} (\mathbf{\Psi}_2 - \psi_1 \psi_1^\top) \mathbf{K}_{ZZ}^{-1} \widehat{\mathbf{M}}, \tag{281}$$

which has rank at most P .

Combining, the moment-matched covariance is

$$\mathbf{S}_a = \gamma \mathbf{I}_R + \widehat{\mathbf{M}}^\top \mathbf{K}_{ZZ}^{-1} (\mathbf{\Psi}_2 - \psi_1 \psi_1^\top) \mathbf{K}_{ZZ}^{-1} \widehat{\mathbf{M}}. \tag{282}$$




Figure 42 | **Direct samples from the quantized manifold mixture in data space.** Here one first samples a manifold component and manifold coordinate, then pushes that sample through the learned GPLVM map and the PRISM synthesis map all the way to Level 0. This shows the nonlinear generative mechanism itself, before moment matching replaces it by a Gaussian approximation in amplitude space.

The message is that the push-forward covariance decomposes into an isotropic term (from averaged predictive noise) and a low-rank term (from input uncertainty). Only the latter carries information about the manifold geometry.

The mixture weights

The weights π_k of the manifold-space GMM (273) are a prior on manifold space. Under the push-forward, they remain the mixture weights in amplitude space:

$$p(\mathbf{a}) \approx \sum_{k=1}^K \pi_k \text{Normal}(\mathbf{a} \mid \boldsymbol{\mu}_k^a, \mathbf{S}_k^a), \quad (283)$$

where $(\boldsymbol{\mu}_k^a, \mathbf{S}_k^a)$ is the moment-matched push-forward of the k -th component via (277) and (282). The weights transform only under Bayesian conditioning on observed data, which happens downstream in IKLP, not here.

Composing with PRISM

The mixture (283) lives in the PRISM amplitude space $\mathbb{R}^R$. To obtain a prior over waveforms $\mathbf{g} \in \mathbb{R}^W$ evaluated on a grid of W time points, recall that PRISM synthesizes via the linear map

$$\mathbf{g} = \boldsymbol{\Phi} \mathbf{a} + \boldsymbol{\varepsilon}, \quad \boldsymbol{\varepsilon} \sim \text{Normal}(\mathbf{0}, \sigma^2 \mathbf{I}_W), \quad (284)$$

where $\boldsymbol{\Phi} \in \mathbb{R}^{W \times R}$ is the design matrix from (221) in Algorithm 2. Because this step is linear, pushing a Gaussian in amplitude space through it preserves Gaussianity:

$$\mathbf{g} \mid k \sim \text{Normal}(\boldsymbol{\Phi} \boldsymbol{\mu}_k^a, \boldsymbol{\Phi} \mathbf{S}_k^a \boldsymbol{\Phi}^\top + \sigma^2 \mathbf{I}_W), \quad (285)$$




Figure 43 | **Samples from the moment-matched prior in data space.** The prior panel draws from the full mixture (286), while the other panels condition on individual mixture components. Compared with Figure 42, these draws are visibly rougher because the nonlinear push-forward has been replaced by a Gaussian approximation in amplitude space, whose covariance includes the extra isotropic uncertainty term of (280).

and the full prior over waveforms is

$$p(\mathbf{g}) \approx \sum_{k=1}^K \pi_k \text{Normal}(\mathbf{g} \mid \Phi \boldsymbol{\mu}_k^a, \Phi \mathbf{S}_k^a \Phi^\top + \sigma^2 \mathbf{I}_W). \quad (286)$$

Low-rank structure

The covariance $\Phi \mathbf{S}_k^a \Phi^\top$ has rank at most $\min(R, \text{rank}(\mathbf{S}_k^a))$. From (282), $\mathbf{S}_k^a$ is a rank- P correction of a scalar multiple of the identity, so the effective rank of the nontrivial part is at most P , which is typically comparable to Q . Each component of (286) is therefore a low-rank Gaussian prior on waveforms plus a diagonal nugget from the PRISM and GPLVM observation noise terms. The cascaded rank reduction is worth emphasizing: PRISM projects from the infinite-dimensional function space to rank R ; the GPLVM further reduces to an effective rank of at most Q (selected automatically by the ARD mechanism); quantization preserves this by working within the Q -dimensional manifold space; and the push-forward inherits the rank bound.³⁰

For DGF waveforms, the ARD mechanism typically selects Q in the range 3 to 8. Each mixture component then defines a waveform prior supported on a subspace of dimension at most Q out of a potentially much larger W . This is genuine dimensionality reduction, and the entire computation required to produce it is amortized at training time.

³⁰More specialized analyses of pushing Gaussians through nonlinear maps, including higher-order corrections beyond moment matching, are given by (Zhang & Shin, 2021) and (Zhang & Shin, 2022).

4.5. Summary

q-GPLVM is a three-stage compression pipeline that converts a collection of PRISM amplitude posteriors into a compact mixture prior over waveforms.

1. A Bayesian GPLVM (Titsias & Lawrence, 2010) discovers the low-dimensional manifold structure of the PRISM amplitude vectors, producing a variational posterior $q(\mathbf{x}_n) = \text{Normal}(\boldsymbol{\mu}_n, \mathbf{S}_n^x)$ for each training waveform in a Q -dimensional manifold space with Q selected automatically by ARD.
2. Extreme deconvolution (Bovy et al., 2011) quantizes the aggregated posterior into a K -component GMM in this manifold space, respecting the per-waveform uncertainties via an errors-in-variables formulation. A broad background component absorbs outliers.
3. Moment matching pushes each GMM component through the trained GP map using the same Ψ statistics that appear in the GPLVM bound, giving a Gaussian in amplitude space ((277), (282)). Because PRISM synthesis is linear, the result extends to a mixture of low-rank Gaussians over waveforms (286).

The entire computation is amortized: once trained, the pipeline outputs a fixed set of K component parameters $\{\pi_k, \boldsymbol{\mu}_k^a, \mathbf{S}_k^a\}_{k=1}^K$ in amplitude space. These are the Gaussian clusters consumed by Chapter 6, where the means are folded into second-moment matrices and a stationary-to-nonstationary interpolation produces the kernel bank for IKLP.

Bibliography

- Airaksinen, M., Raitio, T., Story, B., & Alku, P. (2014). Quasi Closed Phase Glottal Inverse Filtering Analysis With Weighted Linear Prediction. *IEEE/ACM Transactions on Audio, Speech, And Language Processing*, 22(3), 596–607. <https://doi.org/10.1109/TASLP.2013.2294585>
- Alku, P. (1992). *Glottal Wave Analysis with Pitch Synchronous Iterative Adaptive Inverse Filtering*. 11.
- Alku, P. (2011). Glottal Inverse Filtering Analysis of Human Voice Production — A Review of Estimation and Parameterization Methods of the Glottal Excitation and Their Applications. *Sadhana*, 36(5), 623–650. <https://doi.org/10.1007/s12046-011-0041-5>
- Alku, P., Bäckström, T., & Vilkman, E. (2002). Normalized Amplitude Quotient for Parameterization of the Glottal Flow. *The Journal of the Acoustical Society of America*, 112(2), 701–710. <https://doi.org/10.1121/1.1490365>
- Alku, P., Murtola, T., Malinen, J., Kuortti, J., Story, B., Airaksinen, M., Salmi, M., Vilkman, E., & Geneid, A. (2019). OPENGLOT – An Open Environment for the Evaluation of Glottal Inverse Filtering. *Speech Communication*, 107, 38–47. <https://doi.org/10.1016/j.specom.2019.01.005>
- Alku, P., Pohjalainen, J., Vainio, M., Laukkanen, A.-M., & Story, B. H. (2013). Formant Frequency Estimation of High-Pitched Vowels Using Weighted Linear Prediction. *The Journal of the Acoustical Society of America*, 134(2), 1295–1313. <https://doi.org/10.1121/1.4812756>
- Altamirano, M., & Tobar, F. (2022). Nonstationary Multi-Output Gaussian Processes via Harmonizable Spectral Mixtures. *Proceedings of The 25th International Conference on Artificial Intelligence and Statistics*, 3204–3218.
- Alvarado, P. A., & Stowell, D. (2016). Gaussian Processes for Music Audio Modelling and Content Analysis. *Arxiv:1606.01039 [Cs, Stat]*.
- Alzamendi, G. A., & Schlotthauer, G. (2017). Modeling and Joint Estimation of Glottal Source and Vocal Tract Filter by State-Space Methods. *Biomedical Signal Processing and Control*, 37, 5–15. <https://doi.org/10.1016/j.bspc.2016.12.022>
- Ananthapadmanabha, T. V., & Fant, G. (1982). Calculation of True Glottal Flow and Its Components. *Speech Communication*, 1(3), 167–184. [https://doi.org/10.1016/0167-6393\(82\)90015-2](https://doi.org/10.1016/0167-6393(82)90015-2)
- Andrade-Miranda, G., Chatzipapas, K., Arias-Londoño, J. D., & Godino-Llorente, J. I. (2024, December). *GIRAFE: Glottal Imaging Dataset for Advanced Segmentation, Analysis, and Facilitative Playbacks Evaluation* (Issue arXiv:2412.15054). arXiv. <https://doi.org/10.48550/arXiv.2412.15054>

- Antsaklis, P. J., & Michel, A. N. (2006). *Linear Systems*. Springer Science & Business Media.
- Atal, B. S., & Hanauer, S. L. (1971). Speech Analysis and Synthesis by Linear Prediction of the Speech Wave. *The Journal of the Acoustical Society of America*, 50(2B), 637–655.
- Auvinen, H., Raitio, T., Airaksinen, M., Siltanen, S., Story, B. H., & Alku, P. (2014). Automatic Glottal Inverse Filtering with the Markov Chain Monte Carlo Method. *Computer Speech & Language*, 28(5), 1139–1155. <https://doi.org/10.1016/j.csl.2013.09.004>
- Avhad, A., Li, Z., Wilson, A., Sayce, L., Chang, S., Rousseau, B., & Luo, H. (2022). Subject-Specific Computational Fluid-Structure Interaction Modeling of Rabbit Vocal Fold Vibration. *Fluids*, 7(3), 97. <https://doi.org/10.3390/fluids7030097>
- Becker, T., Jessen, M., & Grigoras, C. (2008,). Forensic Speaker Verification Using Formant Features and Gaussian Mixture Models. *Ninth Annual Conference of the International Speech Communication Association*.
- Birkholz, P. (2013). Modeling Consonant-Vowel Coarticulation for Articulatory Speech Synthesis. *Plos One*, 8(4), e60603.
- Blei, D. M., Kucukelbir, A., & McAuliffe, J. D. (2017). Variational Inference: A Review for Statisticians. *Journal of the American Statistical Association*, 112(518), 859–877. <https://doi.org/10.1080/01621459.2017.1285773>
- Bleyer, I. R., Lybeck, L., Auvinen, H., Airaksinen, M., Alku, P., & Siltanen, S. (2017). Alternating Minimisation for Glottal Inverse Filtering. *Inverse Problems*, 33(6), 65005. <https://doi.org/10.1088/1361-6420/aa6eb8>
- Boersma, P. (2001). *Praat, a System for Doing Phonetics by Computer*. 5(9), 341–345.
- Bonastre, J.-F., Kahn, J., Rossato, S., & Ajili, M. (2015). Forensic Speaker Recognition: Mirages and Reality. *S. Fuchs/d*, 255.
- Bovy, J., Hogg, D. W., & Roweis, S. T. (2011). Extreme Deconvolution: Inferring Complete Distribution Functions from Noisy, Heterogeneous and Incomplete Observations. *The Annals of Applied Statistics*, 5(2B). <https://doi.org/10.1214/10-AOAS439>
- Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., & Wanderman-Milne, S. (2020). JAX: Composable Transformations of Python+ NumPy Programs, 2018. URL [Http://github.com/google/jax](http://github.com/google/jax), 4, 16.
- Bretthorst, G. L. (1988). *Bayesian Spectrum Analysis and Parameter Estimation*.
- Calvetti, D., & Somersalo, E. (2018). Inverse Problems: From Regularization to Bayesian Inference. *Wires Computational Statistics*, 10(3), e1427. <https://doi.org/10.1002/wics.1427>
- Cao, F., Vogel, A. P., Gharahkhani, P., & Renteria, M. E. (2025). Speech and Language Biomarkers for Parkinson's Disease Prediction, Early Diagnosis and Progression. *Npj Parkinson's Disease*, 11(1), 57. <https://doi.org/10.1038/s41531-025-00913-4>

- Chang-Sheng Yang, & Kasuya, H. (1998). Automatic Estimation of Formant and Voice Source Parameters Using a Subspace Based Algorithm. *Proceedings of the 1998 IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP '98 (Cat. No.98CH36181)*, 2, 941–944. <https://doi.org/10.1109/ICASSP.1998.675421>
- Chen, C. J. (2013). *Acoustics of Human Voice Production*. 5.
- Chen, C. J. (2016). *Elements of Human Voice*. WORLD SCIENTIFIC. <https://doi.org/10.1142/9891>
- Chien, Y.-R., Mehta, D. D., Guenason, J., Zanartu, M., & Quatieri, T. F. (2017). Evaluation of Glottal Inverse Filtering Algorithms Using a Physiologically Based Articulatory Speech Synthesizer. *IEEE/ACM Transactions on Audio, Speech, And Language Processing*, 25(8), 1718–1730. <https://doi.org/10.1109/TASLP.2017.2714839>
- Childers, D. G., & Hu, H. T. (1994). Speech Synthesis by Glottal Excited Linear Prediction. *The Journal of the Acoustical Society of America*, 96(4), 2026–2036. <https://doi.org/10.1121/1.411319>
- Cho, Y., & Saul, L. (2009). Kernel Methods for Deep Learning. *Advances in Neural Information Processing Systems*, 22.
- Cox, R. T. (1946). Probability, Frequency and Reasonable Expectation. *Physics*.
- Cuevas, A., López, S., Mandic, D., & Tobar, F. (2021). Bayesian Autoregressive Spectral Estimation. *2021 IEEE Latin American Conference on Computational Intelligence (LA-CCI)*, 1–7. <https://doi.org/10.1109/LA-CCI48322.2021.9769834>
- Damianou, A. C., Titsias, M. K., & Lawrence, N. D. (2014, September). *Variational Inference for Uncertainty on the Inputs of Gaussian Process Models* (Issue arXiv:1409.2287). arXiv.
- de Boer, B. (2008). Acoustic Tubes with Maximal and Minimal Resonance Frequencies. *The Journal of the Acoustical Society of America*, 123(5), 3732. <https://doi.org/10.1121/1.2935231>
- Degottex, G. (2010). *Glottal Source and Vocal-Tract Separation* [Doctoral dissertation].
- Degottex, G., Kane, J., Drugman, T., Raitio, T., & Scherer, S. (2014). COVAREP — A Collaborative Voice Analysis Repository for Speech Technologies. *2014 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 960–964. <https://doi.org/10.1109/ICASSP.2014.6853739>
- Deller, J. R. (1983). On the Time Domain Properties of the Two-Pole Model of the Glottal Waveform and Implications for LPC. *Speech Communication*, 2(1), 57–63.
- Deng, L., Cui, X., Pruvencok, R., Huang, J., Momen, S., Chen, Y., & Alwan, A. (2006). A Database of Vocal Tract Resonance Trajectories for Research in Speech Processing. *2006 IEEE International Conference on Acoustics Speech and Signal Processing Proceedings*, 1, I–I.
- Diaconis, P., & Freedman, D. (1999). Iterated Random Functions. *SIAM Review*, 41(1), 45–76.

- Doval, B., D'Alessandro, C., & Henrich, N. (2006). The Spectrum of Glottal Flow Models. *Acta Acustica United with Acustica*, 92(6), 1026–1046.
- Drugman, T., Alku, P., Alwan, A., & Yegnanarayana, B. (2019a, December). *Glottal Source Processing: From Analysis to Applications* (Issue arXiv:1912.12604). arXiv.
- Drugman, T., Bozkurt, B., & Dutoit, T. (2011). Causal–Anticausal Decomposition of Speech Using Complex Cepstrum for Glottal Source Estimation. *Speech Communication*, 53(6), 855–866. <https://doi.org/10.1016/j.specom.2011.02.004>
- Drugman, T., Bozkurt, B., & Dutoit, T. (2019b). A Comparative Study of Glottal Source Estimation Techniques. *Arxiv:2001.00840 [Cs, Eess]*.
- Drugman, T., Thomas, M., Gudnason, J., Naylor, P., & Dutoit, T. (2012). Detection of Glottal Closure Instants From Speech Signals: A Quantitative Review. *IEEE Transactions on Audio, Speech, And Language Processing*, 20(3), 994–1006. <https://doi.org/10.1109/TASL.2011.2170835>
- Ekström, A. G., Moran, S., Sundberg, J., & Lameira, A. (2023). *PREQUEL: Supervised Phonetic Approaches to Analyses of Great Ape Quasi-Vowels*.
- Fant, G. (1960). *Acoustic Theory of Speech Production*. Mouton.
- Fant, G. (1979). Glottal Source and Excitation Analysis. *STL-QPSR*, 1(1979), 85–107.
- Fant, G. (1995). The LF-model Revisited. Transformations and Frequency Domain Analysis. *Speech Trans. Lab. Q. Rep., Royal Inst. Of Tech. Stockholm*, 2(3), 40.
- Fant, G., Liljencrants, J., & Lin, Q.-g. (1985). A Four-Parameter Model of Glottal Flow. *STL-QPSR*, 4(1985), 1–13.
- Flanagan, J. L. (1965). *Speech Analysis Synthesis and Perception*. Springer Berlin Heidelberg. <https://doi.org/10.1007/978-3-662-00849-2>
- Frigola, R. (2015). *Bayesian Time Series Learning with Gaussian Processes* [Doctoral dissertation].
- Fu, Q., & Murphy, P. (2006). Robust Glottal Source Estimation Based on Joint Source-Filter Model Optimization. *IEEE Transactions on Audio, Speech and Language Processing*, 14(2), 492–501. <https://doi.org/10.1109/TSA.2005.857807>
- Fujisaki, H., & Ljungqvist, M. (1986). Proposal and Evaluation of Models for the Glottal Source Waveform. *ICASSP '86. IEEE International Conference on Acoustics, Speech, And Signal Processing*, 11, 1605–1608. <https://doi.org/10.1109/ICASSP.1986.1169239>
- Fulop, S. A. (2011). *Speech Spectrum Analysis*. Springer.
- Fulop, S. A., & Disner, S. F. (2011). Examining the Voice Bar. *Proceedings of Meetings on Acoustics 162ASA*, 14(1), 60002.

- Gao, Y., Birkholz, P., & Li, Y. (2024). Articulatory Copy Synthesis Based on the Speech Synthesizer VocalTractLab and Convolutional Recurrent Neural Networks. *IEEE/ACM Transactions on Audio, Speech, And Language Processing*, 32, 1845–1858. <https://doi.org/10.1109/TASLP.2024.3372874>
- Garofolo, J. S., Lamel, L. F., Fisher, W. M., Pallett, D. S., Dahlgren, N. L., Zue, V., & Fiscus, J. G. (1993). TIMIT Acoustic-Phonetic Continuous Speech Corpus. (*No Title*).
- Ghosal, S., & van der Vaart, A. (2017). *Fundamentals of Nonparametric Bayesian Inference* (1st ed.). Cambridge University Press. <https://doi.org/10.1017/9781139029834>
- Gibson, J. (2018). Entropy Power, Autoregressive Models, and Mutual Information. *Entropy*, 20(10), 750. <https://doi.org/10.3390/e20100750>
- Glorot, X., & Bengio, Y. (2010). Understanding the Difficulty of Training Deep Feedforward Neural Networks. *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics*, 249–256.
- Gobl, C. (2017). Reshaping the Transformed LF Model: Generating the Glottal Source from the Waveshape Parameter Rd. *Interspeech 2017*, 3008–3012. <https://doi.org/10.21437/Interspeech.2017-1140>
- Harrison, P. (2013). *Making Accurate Formant Measurements: An Empirical Investigation of the Influence of the Measurement Tool, Analysis Settings and Speaker on Formant Measurements* [Doctoral dissertation].
- Hensman, J., Fusi, N., & Lawrence, N. D. (2013, September). *Gaussian Processes for Big Data* (Issue arXiv:1309.6835). arXiv. <https://doi.org/10.48550/arXiv.1309.6835>
- Herbst, C. T., Lohscheller, J., Švec, J. G., Henrich, N., Weissengruber, G., & Fitch, W. T. (2014). Glottal Opening and Closing Events Investigated by Electroglottography and Super-High-Speed Video Recordings. *Journal of Experimental Biology*, 217(6), 955–963. <https://doi.org/10.1242/jeb.093203>
- Hornik, K. (1993). Some New Results on Neural Network Approximation. *Neural Networks*, 6(8), 1069–1072.
- Jaynes, E. T. (1984). Prior Information and Ambiguity in Inverse Problems. *Inverse Problems. An International Journal on the Theory and Practice of Inverse Problems, Inverse Methods and Computerized Inversion of Data*, 14, 151–166.
- Jaynes, E. T. (1991). Notes on Present Status and Future Prospects. In W. T. Grandy & L. H. Schick (Eds.), *Maximum Entropy and Bayesian Methods: Maximum Entropy and Bayesian Methods* (pp. 1–13). Springer Netherlands. https://doi.org/10.1007/978-94-011-3460-6_1
- Jaynes, E. T. (2003). *Probability Theory: The Logic of Science*. Cambridge University Press.
- Jopling, R. L. (2009). *Voice Initiation and Voice Offset Patterns in Normal Females: Investigated by High Speed Digital Imaging* [Doctoral dissertation].

- Kadiri, S. R., Alku, P., & Yegnanarayana, B. (2021). Extraction and Utilization of Excitation Information of Speech: A Review. *Proceedings of the IEEE*, 109(12), 1920–1941. <https://doi.org/10.1109/JPROC.2021.3126493>
- Kameoka, H., Ono, N., & Sagayama, S. (2010). Speech Spectrum Modeling for Joint Estimation of Spectral Envelope and Fundamental Frequency. *IEEE Transactions on Audio, Speech, And Language Processing*, 18(6), 1507–1516. <https://doi.org/10.1109/TASL.2009.2036287>
- Kazlauskaitė, I., Ek, C. H., & Campbell, N. D. (2018). Gaussian Process Latent Variable Alignment Learning. *Arxiv Preprint Arxiv:1803.02603*.
- Kent, R. D., & Vorperian, H. K. (2018). Static Measurements of Vowel Formant Frequencies and Bandwidths: A Review. *Journal of Communication Disorders*, 74, 74–97. <https://doi.org/10.1016/j.jcomdis.2018.05.004>
- Klatt, D. H. (1986). Representation of the First Formant in Speech Recognition and in Models of the Auditory Periphery. *Canadian Acoustics*, 14(3bis), 5–7.
- Klatt, D. H., & Klatt, L. C. (1990). Analysis, Synthesis, and Perception of Voice Quality Variations among Female and Male Talkers. *The Journal of the Acoustical Society of America*, 87(2), 820–857. <https://doi.org/10.1121/1.398894>
- Knuth, K. H., & Skilling, J. (2012). Foundations of Inference. *Axioms*, 1(1), 38–73. <https://doi.org/10.3390/axioms1010038>
- Kominek, J., & Black, A. W. (2004,). The CMU Arctic Speech Databases. *Fifth ISCA Workshop on Speech Synthesis*.
- Kreiman, J., Gerratt, B. R., & Antónanzas-Barroso, N. (2007). Measures of the Glottal Source Spectrum. *Journal of Speech, Language, And Hearing Research*, 50(3), 595–610. [https://doi.org/10.1044/1092-4388\(2007/042\)](https://doi.org/10.1044/1092-4388(2007/042))
- Krug, P. K., Stone, S., & Birkholz, P. (2021). Intelligibility and Naturalness of Articulatory Synthesis with VocalTractLab Compared to Established Speech Synthesis Technologies. *11th ISCA Speech Synthesis Workshop (SSW 11)*, 102–107. <https://doi.org/10.21437/SSW.2021-18>
- Lalchand, V., Ravuri, A., & Lawrence, N. D. (2022, April). *Generalised Gaussian Process Latent Variable Models (GPLVM) with Stochastic Variational Inference* (Issue arXiv:2202.12979). arXiv. <https://doi.org/10.48550/arXiv.2202.12979>
- Landau, H. J. (1983). The Inverse Problem for the Vocal Tract and the Moment Problem. *SIAM Journal on Mathematical Analysis*, 14(5), 1019–1035. <https://doi.org/10.1137/0514082>
- Le Besnerais, G., & Giovannelli, J.-F. (2010). Inverse Filtering and Other Linear Methods. In J. Idier (Ed.), *Bayesian Approach to Inverse Problems: Bayesian Approach to Inverse Problems* (pp. 80–116). ISTE. <https://doi.org/10.1002/9780470611197.ch4>

- Lighthill, M. J. (1958). *An Introduction to Fourier Analysis and Generalised Functions*. Cambridge University Press.
- Lindeberg, T. (2013). *Scale-Space Theory in Computer Vision* (Vol. 256). Springer Science & Business Media.
- Little, M. A. (2011). Mathematical Foundations of Nonlinear, Non-Gaussian, and Time-Varying Digital Speech Signal Processing. In C. M. Travieso-González & J. B. Alonso-Hernández (Eds.), *Advances in Nonlinear Speech Processing: Vol. 7015. Advances in Nonlinear Speech Processing* (pp. 9–16). Springer Berlin Heidelberg. https://doi.org/10.1007/978-3-642-25020-0_2
- Little, M. A., McSharry, P. E., Roberts, S. J., Costello, D. A., & Moroz, I. M. (2007). Exploiting Nonlinear Recurrence and Fractal Scaling Properties for Voice Disorder Detection. *Bio-medical Engineering Online*, 6(1), 23. <https://doi.org/10.1186/1475-925X-6-23>
- Ma, A., Lau, K. K., & Thyagarajan, D. (2020). Voice Changes in Parkinson's Disease: What Are They Telling Us?. *Journal of Clinical Neuroscience*, 72, 1–7. <https://doi.org/10.1016/j.jocn.2019.12.029>
- Ma, C., Kamp, Y., & Willems, L. F. (1993). Robust Signal Selection for Linear Prediction Analysis of Voiced Speech. *Speech Communication*, 12(1), 69–81. [https://doi.org/10.1016/0167-6393\(93\)90019-H](https://doi.org/10.1016/0167-6393(93)90019-H)
- MacKay, D. J. (1998). Introduction to Gaussian Processes. *NATO ASI Series F Computer and Systems Sciences*, 168, 133–166.
- MacKay, D. J. C. (2005). *Information Theory, Inference, and Learning Algorithms*. <https://doi.org/10.1198/jasa.2005.s54>
- MacKay, D. J. C., & Bauman Peto, L. C. (1994). *A Hierarchical Dirichlet Language Model*. 1(1), 1–19.
- Markel, J. D., & Gray, A. H. (1976). *Linear Prediction of Speech: 12* (Vol. 12). Springer Berlin Heidelberg. <https://doi.org/10.1007/978-3-642-66286-7>
- Maurer, D. (2016). *Acoustics of the Vowel*. Peter Lang.
- Mehta, D. D., Rudoy, D., & Wolfe, P. J. (2012). Kalman-Based Autoregressive Moving Average Modeling and Inference for Formant and Antiformant Tracking. *The Journal of the Acoustical Society of America*, 132(3), 1732–1746. <https://doi.org/10.1121/1.4739462>
- Milenkovic, P. (1986). Glottal Inverse Filtering by Joint Estimation of an AR System with a Linear Input Model. *IEEE Transactions on Acoustics, Speech, And Signal Processing*, 34(1), 28–42. <https://doi.org/10.1109/TASSP.1986.1164778>
- Miller, R. L. (1959). Nature of the Vocal Cord Wave. *The Journal of the Acoustical Society of America*, 31(6), 667–677. <https://doi.org/10.1121/1.1907771>

- Monahan, J. F. (1984). A Note on Enforcing Stationarity in Autoregressive-Moving Average Models. *Biometrika*, 71(2), 403–404. <https://doi.org/10.1093/biomet/71.2.403>
- Murphy, K. P. (2022). *Probabilistic Machine Learning: An Introduction*. MIT Press.
- Naylor, P. A., Kounoudes, A., Gudnason, J., & Brookes, M. (2007). Estimation of Glottal Closure Instants in Voiced Speech Using the DYPSA Algorithm. *IEEE Transactions on Audio, Speech and Language Processing*, 15(1), 34–43. <https://doi.org/10.1109/TASL.2006.876878>
- Neal, R. M. (1996). Priors for Infinite Networks. In *Bayesian Learning for Neural Networks: Bayesian Learning for Neural Networks* (pp. 29–53). Springer.
- O Cinneide, A. (2012). *Phase-Distortion-Robust Voice-Source Analysis*. <https://doi.org/10.21427/D7FW3B>
- Orbanz, P., & Teh, Y. W. (2010). Bayesian Nonparametric Models. *Encyclopedia of Machine Learning*, 1.
- Orlikoff, R. F., Golla, M. E., & Deliyski, D. D. (2012). Analysis of Longitudinal Phase Differences in Vocal-Fold Vibration Using Synchronous High-Speed Videoendoscopy and Electroglossography. *Journal of Voice : Official Journal of the Voice Foundation*, 26(6), 816.e13–816.e20. <https://doi.org/10.1016/j.jvoice.2012.04.009>
- Orvieto, A., & Gower, R. M. (2025, December). *In Search of Adam's Secret Sauce* (Issue arXiv:2505.21829). arXiv. <https://doi.org/10.48550/arXiv.2505.21829>
- Pandey, G., & Dukkipati, A. (2014). *To Go Deep or Wide in Learning?*.
- Perrotin, O., Feugère, L., & d'Alessandro, C. (2021). Perceptual Equivalence of the Liljencrants–Fant and Linear-Filter Glottal Flow Models. *The Journal of the Acoustical Society of America*, 150(2), 1273–1285. <https://doi.org/10.1121/10.0005879>
- Peterson, G. E., & Barney, H. L. (1952). Control Methods Used in a Study of the Vowels. *The Journal of the Acoustical Society of America*, 24(2), 175–184.
- Plumpe, M., Quatieri, T., & Reynolds, D. (1999). Modeling of the Glottal Flow Derivative Waveform with Application to Speaker Identification. *IEEE Transactions on Speech and Audio Processing*, 7(5), 569–586. <https://doi.org/10.1109/89.784109>
- Press, W. H., Teukolsky, S. A., Vetterling, W. T., & Flannery, B. P. (1992). *Numerical Recipes in C (2nd Ed.): The Art of Scientific Computing*. Cambridge University Press.
- Rabiner, L. R., & Schafer, R. W. (2007). *Introduction to Digital Speech Processing* (Vol. 1). <https://doi.org/10.1561/20000000001>
- Rabiner, L. R., Juang, B.-H., & Rutledge, J. C. (1993). *Fundamentals of Speech Recognition* (Vol. 14). PTR Prentice Hall Englewood Cliffs.
- Rahimi, A., & Recht, B. (2008). Weighted Sums of Random Kitchen Sinks: Replacing Minimization with Randomization in Learning. In D. Koller, D. Schuurmans, Y. Bengio, & L.

- Bottou (Eds.), *Advances in Neural Information Processing Systems: Vol. 21. Advances in Neural Information Processing Systems*.
- Ramsay, J. O., & Silverman, B. W. (1997). *Functional Data Analysis*. Springer.
- Rao, A., & Ghosh, P. K. (2018). Glottal Inverse Filtering Using Probabilistic Weighted Linear Prediction. *IEEE/ACM Transactions on Audio, Speech, And Language Processing*, 27(1), 114–124.
- Rasmussen, C. E., & Williams, C. K. I. (2006). *Gaussian Processes for Machine Learning*. MIT Press.
- Riutort-Mayol, G., Bürkner, P.-C., Andersen, M. R., Solin, A., & Vehtari, A. (2020). Practical Hilbert Space Approximate Bayesian Gaussian Processes for Probabilistic Programming. *Arxiv:2004.11408 [Stat]*.
- Rosenberg, A. E. (1971). Effect of Glottal Pulse Shape on the Quality of Natural Vowels. *The Journal of the Acoustical Society of America*, 49(2B), 583–590. <https://doi.org/10.1121/1.1912389>
- Sasou, A., & Chen, Y. (2023). Comparison of GIF- and SSL-based Features in Pathological-voice Detection. *INTERSPEECH 2023*, 2893–2897. <https://doi.org/10.21437/Interspeech.2023-1418>
- Sawada, H., Kameoka, H., Araki, S., & Ueda, N. (2012). Efficient Algorithms for Multichannel Extensions of Itakura-Saito Nonnegative Matrix Factorization. *2012 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 261–264. <https://doi.org/10.1109/ICASSP.2012.6287867>
- Schleusing, O. (2012). *Multi-Parametric Source-Filter Separation of Speech and Prosodic Voice Restoration* [Technical report].
- Schnupp, J., Nelken, I., & King, A. (2011). *Auditory Neuroscience: Making Sense of Sound*. MIT press.
- Schoder, S., Falk, S., Wurzing, A., Lodermeier, A., Becker, S., & Kniesburges, S. (2024). A Benchmark Case for Aeroacoustic Simulations Involving Fluid-Structure-Acoustic Interaction Transferred from the Process of Human Phonation. *Acta Acustica*, 8, 13. <https://doi.org/10.1051/aacus/2024005>
- Schroeder, M. R. (1999). *Computer Speech: Recognition, Compression, Synthesis*. Springer-Verlag. <https://doi.org/10.1007/978-3-662-03861-1>
- Serwy, R. D. (2017). *Hilbert Phase Methods for Glottal Activity Detection* [Doctoral dissertation].
- Simpson, F., Lalchand, V., & Rasmussen, C. E. (2020,). Marginalised Spectral Mixture Kernels with Nested Sampling. *Third Symposium on Advances in Approximate Bayesian Inference*.

- Sivia, D. S., & Skilling, J. (2006). *Data Analysis: A Bayesian Tutorial* (2nd ed). Oxford University Press.
- Skilling, J. (2005). Bayesics. *AIP Conference Proceedings*, 803, 3–24.
- Stevens, K. N. (2000). *Acoustic Phonetics*. MIT press.
- Stoica, P., & Moses, R. L. (2005). *Spectral Analysis of Signals*. Pearson/Prentice Hall.
- Sun, Y., Huang, Q., & Wu, X. (2022). Unsupervised Inference of Physiologically Meaningful Articulatory Trajectories with VocalTractLab. *Interspeech 2022*, 4661–4665. <https://doi.org/10.21437/Interspeech.2022-659>
- Svec, J. G., Schutte, H. K., Chen, C. J., & Titze, I. R. (2021). Integrative Insights into the Myoelastic-Aerodynamic Theory and Acoustics of Phonation. Scientific Tribute to Donald G. Miller. *Journal of Voice*. <https://doi.org/10.1016/j.jvoice.2021.01.023>
- Sørbye, S. H., & Rue, H. (2017). Penalised Complexity Priors for Stationary Autoregressive Processes. *Journal of Time Series Analysis*, 38(6), 923–935. <https://doi.org/10.1111/jtsa.12242>
- Taylor, A. M., & Reby, D. (2010). The Contribution of Source–Filter Theory to Mammal Vocal Communication Research. *Journal of Zoology*, 280(3), 221–236. <https://doi.org/10.1111/j.1469-7998.2009.00661.x>
- Thomas, M. R. P., Gudnason, J., & Naylor, P. A. (2009). *Data-Driven Voice Source Waveform Modelling*.
- Titsias, M. (2009). Variational Learning of Inducing Variables in Sparse Gaussian Processes. *Proceedings of the Twelfth International Conference on Artificial Intelligence and Statistics*, 567–574.
- Titsias, M. K. (2025, June). *New Bounds for Sparse Variational Gaussian Processes* (Issue arXiv:2502.08730). arXiv. <https://doi.org/10.48550/arXiv.2502.08730>
- Titsias, M., & Lawrence, N. D. (2010). Bayesian Gaussian Process Latent Variable Model. *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics*, 844–851.
- Titze, I. R. (1995). *Workshop on Acoustic Voice Analysis: Summary Statement*. National Center for Voice and Speech.
- Titze, I. R. (2000). *Principles of Voice Production* (Second Printing). Iowa City, IA: National Center for Voice and Speech.
- Titze, I. R., Baken, R. J., Bozeman, K. W., Granqvist, S., Henrich, N., Herbst, C. T., Howard, D. M., Hunter, E. J., Kaelin, D., Kent, R. D., Kreiman, J., Kob, M., Löfqvist, A., McCoy, S., Miller, D. G., Noé, H., Scherer, R. C., Smith, J. R., Story, B. H., ... Wolfe, J. (2015). Toward a Consensus on Symbolic Notation of Harmonics, Resonances, and Formants in

- Vocalization. *The Journal of the Acoustical Society of America*, 137(5), 3005–3007. <https://doi.org/10.1121/1.4919349>
- Turner, R. E., & Sahani, M. (2014). Time-Frequency Analysis as Probabilistic Inference. *IEEE Transactions on Signal Processing*, 62(23), 6171–6183. <https://doi.org/10.1109/TSP.2014.2362100>
- Van Soom, M., & de Boer, B. (2019). A New Approach to the Formant Measuring Problem. *Proceedings*, 33(1), 29. <https://doi.org/10.3390/proceedings2019033029>
- Van Soom, M., & de Boer, B. (2020). Detrending the Waveforms of Steady-State Vowels. *Entropy*, 22(3), 331. <https://doi.org/10.3390/e22030331>
- Van Soom, M., & de Boer, B. (2021). A Weakly Informative Prior for Resonance Frequencies. *Physical Sciences Forum*, 3(1), 2. <https://doi.org/10.3390/psf2021003002>
- Van Soom, M., & De Boer, B. (2022,). Softly Constrained Agent-Based Models. *Joint Conference on Language Evolution*.
- Veldhuis, R. (1998). A Computationally Efficient Alternative for the Liljencrants-Fant Model and Its Perceptual Evaluation. *The Journal of the Acoustical Society of America*, 103(1), 566–571. <https://doi.org/10.1121/1.421103>
- Verdolini, K., & Titze, I. R. (1995). The Application of Laboratory Formulas to Clinical Voice Management. *American Journal of Speech-Language Pathology*, 4(2), 62–69. <https://doi.org/10.1044/1058-0360.0402.62>
- Viterbi, A. (1967). Error Bounds for Convolutional Codes and an Asymptotically Optimum Decoding Algorithm. *IEEE Transactions on Information Theory*, 13(2), 260–269. <https://doi.org/10.1109/TIT.1967.1054010>
- Wang, R., Zhang, Y., Ou, Z., & Hasegawa-Johnson, M. (2016). Use of Particle Filtering and MCMC for Inference in Probabilistic Acoustic Tube Model. *2016 IEEE Statistical Signal Processing Workshop (SSP)*, 1–5. <https://doi.org/10.1109/SSP.2016.7551748>
- Whalen, D. H., Chen, W.-R., Shadle, C. H., & Fulop, S. A. (2022). Formants Are Easy to Measure; Resonances, Not so Much: Lessons from Klatt (1986). *The Journal of the Acoustical Society of America*, 152(2), 933–941. <https://doi.org/10.1121/10.0013410>
- Wilk, M. van der, Dutordoir, V., John, S. T., Artemev, A., Adam, V., & Hensman, J. (2020, March). *A Framework for Interdomain and Multioutput Gaussian Processes* (Issue arXiv:2003.01115). arXiv. <https://doi.org/10.48550/arXiv.2003.01115>
- Williams, C. K. I. (1998). Computation with Infinite Neural Networks. *Neural Computation*, 10(5), 1203–1216. <https://doi.org/10.1162/089976698300017412>
- Wong, D., Markel, J., & Gray, A. (1979). Least Squares Glottal Inverse Filtering from the Acoustic Speech Waveform. *IEEE Transactions on Acoustics, Speech, And Signal Processing*, 27(4), 350–355. <https://doi.org/10.1109/TASSP.1979.1163260>

- Xuan, J., Lu, J., & Zhang, G. (2021, September). *Bayesian Transfer Learning: An Overview of Probabilistic Graphical Models for Transfer Learning* (Issue arXiv:2109.13233). arXiv.
- Yoshii, K., & Goto, M. (2013). Infinite Kernel Linear Prediction for Joint Estimation of Spectral Envelope and Fundamental Frequency. *2013 IEEE International Conference on Acoustics, Speech and Signal Processing*, 463–467. <https://doi.org/10.1109/ICASSP.2013.6637690>
- Zhang, B., & Shin, Y. C. (2021). An Adaptive Gaussian Mixture Method for Nonlinear Uncertainty Propagation in Neural Networks. *Neurocomputing*. <https://doi.org/10.1016/j.neucom.2021.06.007>
- Zhang, B., & Shin, Y. C. (2022). A Gaussian Mixture Filter with Adaptive Refinement for Nonlinear State Estimation. *Signal Processing*, 201, 108677. <https://doi.org/10.1016/j.sigpro.2022.108677>
- Zhang, Y., Zheng, X., & Xue, Q. (2020). A Deep Neural Network Based Glottal Flow Model for Predicting Fluid-Structure Interactions during Voice Production. *Applied Sciences*, 10(2), 705. <https://doi.org/10.3390/app10020705>

Appendices

Appendix A. Regularization and priors

As its name suggests, GIF belongs to the class of *inverse problems* (Le Besnerais & Giovannelli, 2010). Inverse problems confront us when we have observations of some effects [noisy observations of $s(t)$ in our case] of which we want to infer the causes $u(t)$ and $h(t)$ within the context of some model $s(t) = u(t) * h(t) * r(t)$.

In the applied sciences it is very often the case for inverse problems that the observed data are quite consistent with a broad and sometimes colorful class of possible solutions rather than a unique and logically deductible solution. From the viewpoint of pure mathematics therefore such inverse problems are often deemed ‘ill-posed’, but perhaps a more productive viewpoint³¹ is that such problems are simply underdetermined.

Underdetermined problems are ubiquitous in a wide range of applications, from medical diagnosis (what could be causing the patient’s persistent cough) to astronomy (what is perturbing the orbit of Uranus). Bayesian theory by itself does not require any particular structure on the “feasible set” of plausible solutions and as such underdetermined problems have no special status within it, nor pose any special challenge to it (Jaynes, 2003). In practice however, we would like to have a more well determined solution to underdetermined problems as our computing budgets (and patience) are limited. For example, it would most certainly help if we can already discard those solutions that lack to some degree a set of required properties we established beforehand. This is where *regularization* and *priors* come into play.

In general, *regularization* is the process of constraining the solution space of some problem to obtain a more stable and perhaps unique solution; in other words, to convert an underdetermined problem into something more ‘well-determined’ and, hopefully, more amenable to computational optimization. For example, ridge regression is a widely used regularization technique to numerically stabilize linear regression in high dimensions.³²

In the Bayesian framework regularization is imposed simply through careful assignment of the probability distributions describing a specific problem – of which the bulk of the effort is traditionally concentrated on the prior probabilities, or *priors* (Sivia & Skilling, 2006). Assuming our prior information about the problem at hand is correct and the assignment of the priors accurately reflect that information (and we got away with the approximations made in our inference methods) we will have automatically implemented just the right degree of regularization for our problem; indeed, specially catered to the one data set we

³¹And closer to the historical origin of the epithet “mal posée” (Jaynes, 1984).

³²Ridge regression (also known as ℓ^2 regularization) can be derived from Bayesian theory as standard linear regression with Gaussian priors on the amplitudes (Murphy, 2022, Sec. 11.3). This principled view allows us to go much further than numerical stabilization. We will show in the next chapters that ridge regression can (approximately) express intricate pieces of prior information such as definite signal polarity, the differentiability class to which a function belongs, and power constraints on waveforms.

happened to observe, because we express everything in terms of posterior probabilities that are conditioned on the observed data.

For simple inverse problems where intuition can see the answer clearly, Bayesian regularization leads to answers that comply with common sense (MacKay, 2005). This is simply because Bayesian probability theory is a formalized theory of plausible reasoning itself (Skilling, 2005), which humans undertake every day. But regularization in the Bayesian framework can go beyond our intuition: it will of course continue to work for complex problems that require considerable analysis, and this consistency makes it a valuable tool for attacking inverse problems (Calvetti & Somersalo, 2018).

Appendix B. Parametric, nonparametric, semiparametric?

We clarify some of the confusing terminology associated with Bayesian nonparametric models.

Parametric models are simply models with a fixed and prespecified amount of unknown parameters R , independent of the size of the dataset N .

Examples are linear regression and neural networks. In BNGIF the filter priors π_{PZ} and π_{AP} priors are parametric as they have a fixed amount of parameters R depending on their order K , similar to linear prediction (LP) methods. And of course the π_{LF} prior with $R = 4$ is parametric as well.

Nonparametric models in a Bayesian context are models for which the effective number of parameters $\tilde{R}$ is capable of adapting to N (Orbanz & Teh, 2010).

They are also referred to as infinite-dimensional parametric models, which is just a shorthand way of saying that they are parametric models which have been mathematically constructed to behave consistently under inference as $R \rightarrow \infty$. That is, given the data of size N , a nonparametric model roughly behaves as an ordinary parametric model with $\tilde{R} < \infty$ parameters, with the exception that $\tilde{R}$ is determined automatically from the data rather than specified beforehand.

Examples are Dirichlet process mixture models, for which the effective number of mixture components does not need to be known beforehand, and, of course, GPs, which are infinite-dimensional distributions over functions. In BNGIF the π_{GP} and π_{VS} priors are nonparametric because they are GPs.³³

The name ‘nonparametric’ is a bit of a misnomer because nonparametric models almost always have a set of Q hyperparameters which typically determine either the scales present in the data or the degree to which sparsity or smoothness is encouraged. In theory, inference of these Q hyperparameters is no different from inference with a parametric model of order $R = Q + \tilde{R}$ in which $\tilde{R}$ parameters have been marginalized over analytically.

Semiparametric models in a Bayesian context are models with both parametric and nonparametric components for which the latter is considered a nuisance parameter (Ghosal & van der Vaart, 2017, p. 368). This means that we are mainly interested in conclusions about the parametric component of the model, which require marginalizing over the nonparametric component.

In BNGIF the parametric and nonparametric components are very neatly divided into models for the filter $h(t)$ and the source $u(t)$, respectively, reflecting the qualitative difference between the information available from acoustic phonetics about the two. One might be

³³Strictly, these are not fully nonparametric because they are only reduced-rank GPs and not full-rank GPs. Their capacity to adapt their complexity to growing amounts of data N is limited by the number of basisfunctions M ; nonparametric behavior is recovered in the limit $M \rightarrow \infty$.

tempted to call BNGIF semiparametric for this reason if it weren't for the fact that the nonparametric component describing $u(t)$ is actually the main object of interest for the GIF problem. If anything the nuisance bit for GIF is in the parametric component describing $h(t)$, not the nonparametric component for $u(t)$.

Appendix C. JAX implementation of the Liljencrants-Fant model

```
1 import jax
2 import jax.numpy as jnp
3 import jaxopt
4
5
6 def _get_initial_bracket(T0, tol, UPPER_FACTOR=1e3):
7     lower = (1 / T0) * tol
8     upper = (1 / T0) * UPPER_FACTOR
9     return lower, upper
10
11
12 def _bisect_exponential_rate(
13     f, tol, initial_bracket, maxiter, fail_val=jnp.nan, **kwargs
14 ):
15     lower, upper = initial_bracket
16     bisec = jaxopt.Bisection(
17         optimality_fun=f,
18         lower=lower,
19         upper=upper,
20         maxiter=maxiter,
21         tol=tol,
22         check_bracket=False,
23     )
24     result = bisec.run(**kwargs)
25     success = result.state.error < tol
26     exponential_rate = result.params
27     return jax.lax.cond(success, lambda: exponential_rate, lambda:
fail_val)
28
29
30 def _bisect_epsilon(p, tol, initial_bracket, maxiter):
31     def f(epsilon, p):
32         LHS = epsilon * p["Ta"]
33         RHS = 1 - jnp.exp(-epsilon * (p["T0"] - p["Te"]))
34         return LHS - RHS
35
36     return _bisect_exponential_rate(f, tol, initial_bracket, maxiter,
p=p)
37
38
39 def _bisect_a(p, epsilon, tol, initial_bracket, maxiter):
40     def f(a, p, epsilon):
41         factor = 1 / (a**2 + (jnp.pi / p["Tp"]) ** 2)
42         term1 = (
43             jnp.exp(-a * p["Te"])
44             * (jnp.pi / p["Tp"])
45             / jnp.sin(jnp.pi * p["Te"] / p["Tp"])
46         )
47         term2 = a
48         term3 = (
49             jnp.pi / p["Tp"] * (1 / jnp.tan(jnp.pi * p["Te"] / p["Tp"]))
50         ) # cotg() is 1/tan()
```

```

51     LHS = factor * (term1 + term2 - term3)
52
53     RHS = (p["T0"] - p["Te"]) / (
54         jnp.exp(epsilon * (p["T0"] - p["Te"])) - 1
55     ) - 1 / epsilon
56
57     return LHS - RHS
58
59     # Assume that the bisection will only fail when `a -> 0`
60     return _bisect_exponential_rate(
61         f, tol, initial_bracket, maxiter, fail_val=0.0, p=p,
62         epsilon=epsilon
63     )
64
65 def _dgf(t, p, epsilon, a, **kwargs):
66     """Implement the LF model per Doval et al. (2006) Section A1.4."""
67
68     def rise(t):
69         return (
70             -jnp.exp(a * (t - p["Te"]))
71             * jnp.sin(jnp.pi * t / p["Tp"])
72             / jnp.sin(jnp.pi * p["Te"] / p["Tp"])
73         )
74
75     def decrease(t):
76         return (
77             -1
78             / (epsilon * p["Ta"])
79             * (
80                 jnp.exp(-epsilon * (t - p["Te"]))
81                 - jnp.exp(-epsilon * (p["T0"] - p["Te"]))
82             )
83         )
84
85     return jnp.pieceswise(
86         t,
87         [t < 0, (0 <= t) & (t < p["Te"]), (p["Te"] <= t) & (t <=
88         p["T0"])],
89         [0.0, rise, decrease, 0.0],
90     )
91
92 def _nans_like(a):
93     return jnp.full_like(a, jnp.nan)
94
95
96 def dgf(t, p, offset=0.0, tol=1e-6, initial_bracket=None, maxiter=100):
97     """Calculate the glottal flow derivative (DGF) using the LF
98     (Liljencrants-Fant) model
99
100     This JAX implementation can be jitted and is differentiable with
101     respect to `p`.

```

```

100     Jitting this function will gain several a speedup of several orders
101     of magnitude.
102     The LF model [1,2] is nonzero only for values of `t` in `[offset,
103     offset + T0]`, where
104     `T0 = p['T0']`. If the LF parameters in `p` are inconsistent, such
105     that the
106     implicit equations cannot be solved, the function's output `u` will
107     be all `nan`s.
108
109     Args:
110     `t` (array): Time points at which to evaluate the DGF `u(t)`
111     `p` (dict): Dictionary with the LF parameters. It must have the
112     following keys:
113     `['T0', 'Te', 'Tp', 'Ta']`. `Ee` is always assumed to be 1.
114     `offset` (scalar): The DGF waveform starts at this offset and is
115     nonzero in
116     `[offset, offset + T0]`.
117     `tol`, `initial_bracket`, `maxiter` (scalar): Parameters
118     controlling the bisection
119     routine `_bisect_exponential_rate()` to solve the two
120     equations implicit in
121     the LF model.
122
123     Returns:
124     `u` (DeviceArray): The glottal flow derivative evaluated at
125     times `t` if the LF
126     parameters in `p` are consistent, otherwise an array of
127     `nan`s shaped as `t`.
128
129     References:
130     [1] G. Fant, J. Liljencrants, and Q. Lin, "A four-parameter
131     model of glottal flow",
132     STL-QPSR, vol. 4, no. 1985, pp. 1–13, 1985.
133     [2] Doval, Boris, Christophe d'Alessandro, and Nathalie Henrich.
134     "The spectrum of
135     glottal flow models." Acta acustica united with acustica
136     92.6 (2006): 1026-1046.
137     """
138     if initial_bracket is None:
139         initial_bracket = _get_initial_bracket(p["T0"], tol)
140
141     epsilon = _bisect_epsilon(
142         p, tol=tol, initial_bracket=initial_bracket, maxiter=maxiter
143     )
144     a = _bisect_a(
145         p, epsilon=epsilon, tol=tol, initial_bracket=initial_bracket,
146         maxiter=maxiter
147     )
148     u = _dgf(t - offset, p=p, epsilon=epsilon, a=a)
149     inconsistent_parameters = jnp.any(jnp.isnan(u))

```

```

137     return jax.lax.cond(inconsistent_parameters, lambda: _nans_like(u),
138                          lambda: u)
139
140 def consistent_lf_params(p):
141     """Check whether the implicit LF equations can be solved given the
142     LF parameters in `p`"""
143     return ~jnp.any(jnp.isnan(dgf(0.0, p)))
144
145 def _select_keys(q, *keys):
146     return {k: q[k] for k in keys if k in q}
147
148
149 def convert_lf_params(p, s, join=True):
150     if s == "R -> generic":
151         # Perrotin et al. (2021) Eq. (A1)
152         Ra, Rk, Rg = p["Ra"], p["Rk"], p["Rg"]
153         Oq = (1 + Rk) / (2 * Rg)
154         am = 1 / (1 + Rk)
155         Qa = Ra / (1 - Oq) # Doval (2006) p. 5
156         q = dict(Oq=Oq, am=am, Qa=Qa)
157     elif s == "generic -> R":
158         Oq, am, Qa = p["Oq"], p["am"], p["Qa"]
159         Rk = 1 / am - 1
160         Rg = (1 + Rk) / (2 * Oq)
161         Ra = Qa * (1 - Oq)
162         q = dict(Ra=Ra, Rk=Rk, Rg=Rg)
163     elif s == "T -> R":
164         # Fant (1994)
165         T0, Te, Tp, Ta = p["T0"], p["Te"], p["Tp"], p["Ta"]
166         Ra = Ta / T0
167         Rk = (Te - Tp) / Tp
168         Rg = T0 / (2 * Tp)
169         q = dict(Ra=Ra, Rk=Rk, Rg=Rg)
170     elif s == "R -> T":
171         T0, Ra, Rk, Rg = p["T0"], p["Ra"], p["Rk"], p["Rg"]
172         Ta = T0 * Ra
173         Tp = T0 / (2 * Rg)
174         Te = Tp * (1 + Rk)
175         q = dict(Te=Te, Tp=Tp, Ta=Ta)
176     elif s == "generic -> T":
177         q = convert_lf_params(convert_lf_params(p, "generic -> R"), "R -
178 > T")
179         q = _select_keys(q, "Te", "Tp", "Ta")
180     elif s == "T -> generic":
181         q = convert_lf_params(convert_lf_params(p, "T -> R"), "R ->
182 generic")
183         q = _select_keys(q, "Oq", "am", "Qa")
184     elif s == "Rd -> R":
185         # Perrotin et al. (2021) Eq. (A1)
186         # See also Fant (1994) for the meaning of these dimensionless
187         parameters

```

```

185     Rd = p["Rd"]
186     Ra = (-1 + 4.8 * Rd) / 100
187     Rk = (22.4 + 11.8 * Rd) / 100
188     Rg = Rk * (0.5 + 1.2 * Rk) / (0.44 * Rd - 4 * Ra * (0.5 + 1.2 *
Rk))
189     q = dict(Ra=Ra, Rk=Rk, Rg=Rg)
190     elif s == "Rd -> T":
191         q = convert_lf_params(convert_lf_params(p, "Rd -> R"), "R -> T")
192         q = _select_keys(q, "Te", "Tp", "Ta")
193     else:
194         raise ValueError(f"Unknown conversion: {s}")
195     return {**p, **q} if join else q
196
197
198 def fant_params(s):
199     """Generate typical LF model parameters for a male or female vowel
200
201     Note that the unit of time is `msec`. Values are taken from [1, p.
121].
202
203     [1] Fant, Gunnar. "The LF-model revisited. Transformations and
frequency domain analysis."
204         Speech Trans. Lab. Q. Rep., Royal Inst. of Tech. Stockholm 2.3
(1995): 40.
205     """
206     if s == "female vowel":
207         return {"T0": 5.0, "Te": 3.25, "Tp": 2.5, "Ta":
0.3978873577297384}
208     elif s == "male vowel":
209         return {
210             "T0": 8.333333333333334,
211             "Te": 4.513888888888889,
212             "Tp": 3.472222222222228,
213             "Ta": 0.22736420441699334,
214         }
215     else:
216         raise ValueError(s)

```

Appendix D. More comments on the nonparametric glottal flow model

Some aspects of the nonparametric GFM derived in Chapter 4 are discussed in more detail here.

Infinite precision? In the GP view, increasing H increases the *Monte Carlo resolution* with which the finite neural network model samples the arc cosine kernel — the deviation from the limiting kernel vanishes as $\mathcal{O}(1/\sqrt{H})$ by (154). The nonparametric limit replaces explicit random changepoints with their continuous Gaussian measure, and the prior over $u'(t)$ becomes a GP whose covariance is the arc cosine kernel restricted to the open phase.

There is a caveat though. We gained full rank — the prior now has support over an infinite-dimensional function space — but we lost the ability to represent true discontinuities. Arc cosine GP sample paths are continuous; they cannot jump.³⁴ This is the analogue of the amplitude prior constraining the ellipsoid in the finite- H case (footnote 17 of Chapter 4): there we traded a large discrete family of possible changepoint configurations for a continuous Gaussian prior that concentrates mass inside a smooth ellipsoid; here we trade a distribution over Poisson-style jump processes for a Gaussian prior that concentrates mass on smooth function classes. The best the arc cosine GP can do is approximate a sharp GCI with a steep ramp.

Whether this is a problem in practice is an empirical question. The closed-phase GCI is physiologically sharp but not infinitely so, and the DGF data from VocalTractLab are sampled at finite rate, so the distinction between “true jump” and “ramp steep enough that it spans one sample” may not be detectable.

Why this matters for the choice of modeled signal The continuity constraint has a direct implication for which signal we should model with the arc cosine GP. We model $u'(t)$ — the DGF — because the GCI appears as a sharp negative peak there, which is easier to detect and align than the integral $u(t)$. But if the arc cosine GP is continuous, then $u'(t)$ is smooth, which means $u(t)$ is differentiable — and real glottal flow does not have a differentiable closure.

One way around this is to model $u(t)$ directly with the arc cosine GP instead of $u'(t)$. Then the GP sample paths are continuous (which $u(t)$ should be) and discontinuities can be pushed into the *derivative* implicitly via the radiation factor: the combined vocal tract and radiation impulse response is modeled as the AR filter, and the combined source signal becomes $u(t) * r(t)$ rather than $u'(t)$ alone. This is, at bottom, what classical LPC does — it posits $s(t) = e(t) * h(t)$ and allows $h(t)$ to absorb the radiation effect $r(t)$, since differentiation corresponds to a zero in the transfer function, not a pole, so $h(t)$ can remain all-pole. From this viewpoint the present approach generalizes LPC in exactly that way.

³⁴Specifically, sample paths of a GP with a kernel of degree d are d -times continuously differentiable with probability one (Rasmussen & Williams, 2006). For $d = 0$ this gives continuous but nowhere-differentiable paths; for $d \geq 1$, paths are smoother still. In either case there are no jumps.

The alternative of obtaining $u'(t)$ implicitly — from the spectral domain or from the AR pole structure — is also possible and may allow the sharp-closure structure to be recovered without committing the GP to model it directly.

Why not go infinite depth? Increasing depth has a different character from increasing width. For a deep GP with Gaussian kernels at each layer, letting the depth diverge drives the process toward one that is independent of its inputs (Diaconis & Freedman, 1999) — a known pathology in practice, where deep GPs tend to “forget” their inputs and produce nearly constant sample paths. Finite but large depth mitigates this, but one might note that infinite width carries a vaguely analogous cost: in the GP limit, the basis functions are frozen at their prior expectation and no longer adapt to the data as drawn features. What is learned is entirely encoded in the kernel hyperparameters. This makes kernel learning — and in particular the multiple-kernel structure of IKLP — the mechanism through which the model does acquire data-driven spectral features, rather than learning them through random feature adaptation.

Why not spline models? Piecewise spline models are well-studied and effective in low-dimensional nonparametric regression. The issue here is resolution dependence. As GP priors, spline kernels produce posterior means that are splines with knots fixed at the observed input locations (MacKay, 1998, p. 6): add more data, add more knots. The arc cosine kernel, by contrast, learns the effective number of hinge points from the data through its hyperparameters, and this number can remain $\mathcal{O}(1)$ even as the number of observations grows. For glottal flow this matters: a DGF waveform sampled at 16 kHz and one sampled at 44 kHz should both require $\mathcal{O}(1)$ effective changepoints to describe the open phase, not a number proportional to the sampling rate.

Why not Lévy processes? These encode Poisson-style jumps, $\mathcal{O}(1)$ in number, which is exactly the right prior for a signal with a small number of hard discontinuities. The problem is inference: the posterior over a Lévy process is indexed by the jump locations, and marginalizing those out requires MCMC or particle methods that scale with the number of jumps. The GP approach achieves fast closed-form posterior inference precisely because everything — changepoints, amplitudes, timing — has been absorbed into the kernel. Trading exact jumps for steep smooth ramps is the price of that tractability.

Appendix E. On derivative time scales

The EGIFA and CommonVoice datasets present us with ground-truth $u(t)$ but we need $u'(t)$ for learning. Naive finite difference of $u(t)$,

$$u'_\Delta(t) = \frac{u(t + \Delta) - u(t)}{\Delta}, \quad \Delta = \frac{1}{f_s}, \quad (287)$$

often yields spikes in $u'(t)$ which effectively defeat kernel modeling because the increased observation noise will prevent learning fine structure. Luckily, we are saved by the physiology of the speech production process.

Glottal closure is not an instantaneous discontinuity but a short transition extending over a finite time interval. High-speed imaging and synchronized EGG measurements show that opening and closing events occupy $O(0.1)$ ms durations and may exhibit small temporal offsets depending on how closure propagates along the folds (e.g. zipper-like anterior-posterior contact), as described by Herbst et al. (2014) and Orlikoff et al. (2012). Consequently, a derivative operator that reacts to arbitrarily small sample-scale fluctuations does not reflect the underlying physiology; it instead amplifies numerical artefacts or simulator-specific microstructure. A meaningful derivative therefore requires an explicit temporal scale reflecting the duration over which closure dynamics remain physically interpretable.

Since we already model $u(t)$ by Gaussian processes,³⁵ we are implicitly asserting that first and second moments alone carry most of the important information. Therefore, we propose to use a *Gaussian-smoothed derivative* to measure the slope of a *locally averaged signal*

$$u'_\sigma(t) := \partial_t(G_\sigma * u)(t), \quad G_\sigma(t) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{t^2}{\sigma^2}\right), \quad (288)$$

rather than differentiating raw samples directly.³⁶ This type of differentiation results in another GP, since differentiation and convolution (smearing) are closed operators in GP space.

Selecting σ such that the effective smoothing width corresponds to approximately 0.1 ms (about four to five samples at $f_s = 44.1$ kHz) aligns the operator with the expected duration of glottal closure transitions. At smaller scales the derivative is dominated by high-frequency noise and discretization effects; at larger scales the transition itself becomes blurred. The resulting derivative can therefore be interpreted as a slope measured at the physically relevant temporal resolution rather than as a raw numerical difference, and this removes the occurrence of large spikes in $u'(t)$ almost completely; an example is shown in Figure 44.

³⁵More precisely, we model its derivative $u'(t)$ as a GP, but the integration in (8) preserves GP-ness.

³⁶Gaussian smoothing also turns out to be the unique linear kernel that preserves causality in scale-space, meaning that increasing scale removes fine structure without introducing spurious extrema (Lindeberg, 2013).




Figure 44 | **Smooth differentiation dampens spikes in $u'(t)$.** The top panel shows the ground truth glottal flow $u(t)$ extracted from a random EGIFA file. The middle panel contains its derivative $u'_\Delta(t)$ obtained by finite differences (287), and the bottom panel the Gaussian-smoothed derivative $u'_\sigma(t)$ proposed in (288). It can be seen that raw differentiation amplifies sample-scale fluctuations into large spikes, whereas the smoothed derivative retains the closure structure at the intended temporal resolution.

Appendix F. Predicting period-to-period jitter

Classical local jitter is defined as the expected absolute change in consecutive periods, normalized by the mean period,

$$\text{jitter} := \frac{\langle |T_{k+1} - T_k| \rangle}{\langle T_k \rangle}. \quad (289)$$

In the present framework, jitter is not introduced as an independent noise source. Instead, it emerges deterministically from a smooth stochastic model of phase warping. All variability originates from curvature of the time warp $t(\tau)$, induced by a Gaussian process prior on the log-period

$$g(\tau) = \log T(\tau). \quad (290)$$

Sampling the process at integer phases $\tau = k$ yields a stationary Gaussian sequence

$$g_k := g(k), \quad (291)$$

with mean μ , marginal variance σ^2 , and lag-one correlation

$$\rho(\ell) := \frac{k_{\ell(1)}}{k_{\ell(0)}}, \quad (292)$$

which depends only on the kernel family and the lengthscale ℓ measured in cycles.

For the kernels considered here, the lag-one correlation is given by³⁷

The observed periods are lognormal variables,

$$T_k = \exp(g_k), \quad (293)$$

and the pair (T_k, T_{k+1}) is bivariate lognormal. Because g_k and g_{k+1} have equal variance, the sum and difference variables

$$S = \frac{g_{k+1} + g_k}{2}, \quad D = g_{k+1} - g_k \quad (294)$$

are independent Gaussians, with

$$S \sim \text{Normal}\left(\mu, \sigma^2 \frac{1 + \rho}{2}\right), \quad D \sim \text{Normal}(0, 2\sigma^2(1 - \rho)). \quad (295)$$

Using the identity

³⁷Matérn 1/2: $\rho(\ell) = \exp(-\frac{1}{\ell})$. Matérn 3/2: $\rho(\ell) = \left(1 + \frac{\sqrt{3}}{\ell}\right) \exp\left(-\frac{\sqrt{3}}{\ell}\right)$. Matérn 5/2: $\rho(\ell) = \left(1 + \frac{\sqrt{5}}{\ell} + \frac{5}{3\ell^2}\right) \exp\left(-\frac{\sqrt{5}}{\ell}\right)$. Squared exponential: $\rho(\ell) = \exp\left(-\frac{1}{2\ell^2}\right)$.

$$|\exp(a) - \exp(b)| = 2 \exp\left(\frac{a+b}{2}\right) \left| \sinh\left(\frac{a-b}{2}\right) \right|, \quad (296)$$

the expected absolute period difference factorizes as

$$\langle |T_{k+1} - T_k| \rangle = 2 \langle \exp(S) \rangle \left\langle \left| \sinh\left(\frac{D}{2}\right) \right| \right\rangle. \quad (297)$$

The first factor evaluates to

$$\langle \exp(S) \rangle = \exp\left(\mu + \sigma^2 \frac{1+\rho}{4}\right), \quad (298)$$

while the second admits a closed-form expression in terms of the Gaussian error function. After normalization by $\langle T_k \rangle = \exp\left(\mu + \frac{\sigma^2}{2}\right)$, all dependence on μ cancels, yielding the exact prediction

$$\langle \text{jitter} \mid \ell \rangle = 4\Phi\left(\sigma\sqrt{\frac{1-\rho(\ell)}{2}}\right) - 2, \quad (299)$$

where Φ denotes the standard normal cumulative distribution function.

This expression constitutes a quantitative theory of jitter. Given the marginal log-period variability σ and the kernel-dependent correlation $\rho(\ell)$, it predicts the expected period-to-period jitter without further assumptions.

For small arguments, the Gaussian CDF admits the expansion

$$\Phi(x) \approx \frac{1}{2} + \frac{x}{\sqrt{2\pi}}, \quad (300)$$

which yields the asymptotic scaling law

$$\langle \text{jitter} \mid \ell \rangle \approx \left(2\frac{\sigma}{\sqrt{\pi}}\right) \sqrt{1-\rho(\ell)}, \quad (301)$$

recovering the intuitive result that jitter vanishes as $\rho(\ell) \rightarrow 1$, i.e. as the phase warp becomes locally affine.

To make the prediction concrete, we fix a representative adult pitch distribution with median period $T_0 \approx 6.7$ ms (150 Hz) and log-period standard deviation $\sigma = 0.25$. The table below reports the predicted mean jitter for several kernel families and lengthscales.

Normal adult jitter values around 0.5% are only attained for lengthscales of several tens of cycles under smooth kernels. Rougher kernels require substantially larger ℓ to suppress cycle-to-cycle variation. In this sense, ℓ admits a direct physiological interpretation as a cycle-to-cycle smoothness parameter governing the temporal regularity of voicing.

ℓ (cycles)	$\nu = 1/2$	$\nu = 3/2$	$\nu = 5/2$	$\nu = \infty$
5	12.0%	6.2%	5.0%	4.0%
10	8.7%	3.3%	2.6%	2.0%
20	6.2%	1.7%	1.3%	1.0%
40	4.4%	0.85%	0.64%	0.50%
80	3.1%	0.43%	0.32%	0.25%

Table 9 | **Predicted mean period-to-period jitter** as a function of lengthscale ℓ for Matérn kernels with $\nu \in \{\frac{1}{2}, \frac{3}{2}, \frac{5}{2}\}$ and the squared exponential kernel ($\nu = \infty$). Predictions use $\sigma = 0.25$ and are obtained from the exact closed-form expression

$$\langle \text{jitter} \mid \ell \rangle = 4\Phi\left(\sigma\sqrt{\frac{1-\rho(\ell)}{2}}\right) - 2. \quad (302)$$

This framework yields falsifiable predictions linking observed jitter statistics to latent smoothness of the underlying pitch trajectory, which can subsequently be tested by fitting the model to data.

Appendix G. Gauge invariance in glottal inverse filtering evaluation

Every evaluation metric carries an implicit assumption about what constitutes a meaningful difference between two signals. Spectral metrics make this assumption explicit: envelope-based measures deliberately ignore phase, ignore absolute timing, and ignore excitation fine structure. The question is whether source-domain evaluation — comparing an estimated DGF waveform against a ground-truth one — should be held to the same standard, or whether it is legitimate to compare raw waveforms sample by sample.

This chapter argues that it is not legitimate, and that doing so systematically penalizes algorithms for choices that are physically unobservable — choices about what physicists call the *gauge*. The argument is not a matter of taste. It follows directly from the structure of the source-filter model and from the physics of acoustic propagation. The practical consequence is a specific evaluation procedure, aligned at frame level with an affine correction, which we use throughout the experimental chapters.

G.1. The source-filter gauge group

The speech production model is convolutional:

$$y(t) = (u'(t) * h(t)) + \varepsilon(t), \quad (303)$$

where $u'(t)$ is the glottal flow derivative, $h(t)$ is the vocal tract impulse response including radiation, and $\varepsilon(t)$ is noise. The inverse filtering problem asks: given $y(t)$, recover $u'(t)$ and $h(t)$.

The trouble is that the factorization is not unique. Two transformations leave $y(t)$ invariant.

The first is *scale*: for any nonzero constant α , the pair $(\alpha u'(t), \frac{h(t)}{\alpha})$ produces exactly the same observed signal as $(u'(t), h(t))$. Absolute source amplitude is therefore not recoverable from acoustics alone; it is determined by convention.

The second is *time-translation*: for any constant delay τ , the pair $(u'(t - \tau), h(t + \tau))$ again produces the same convolution. The two transformations together form a gauge group acting on (u', h) : the observable signal depends only on the equivalence class

$$[(u', h)] = \{(\alpha u'(t - \tau), h(t + \tau)/\alpha) \mid \alpha \in \mathbb{R} \setminus \{0\}, \tau \in \mathbb{R}\}. \quad (304)$$

Inverse filtering does not recover a unique (u', h) pair. It selects one representative from this class.

The implication for evaluation is immediate: different algorithms can legitimately choose different representatives while producing identical acoustic reconstructions. An evaluation that compares representatives directly — rather than comparing the equivalence classes they represent — is measuring gauge choice, not signal quality. This is not a hypothetical concern. The closed-phase LP-based algorithms (CPCA, SLP, WLP) anchor the filter by minimizing LP residual energy in the closed phase, which drives the gauge toward a

particular delay convention. Complex cepstrum decomposition (CCD) (Drugman et al., 2011) does something entirely different: it separates the maximum-phase anticausal component of the cepstrum, which selects a different gauge altogether. The two families are not wrong in different ways; they are right in different gauges. Evaluating them against the same raw waveform reference therefore introduces a systematic bias.

G.2. Physical dispersion reinforces the ambiguity

The gauge ambiguity from the source-filter model is compounded by a physical one. The vocal tract is a resonant system with a frequency-dependent phase response. Its group delay

$$\tau_{g(\omega)} = -\frac{d}{d\omega} \arg H(e^{j\omega}) \quad (305)$$

is not constant across frequency: narrow formants store and release energy over extended time windows, smearing the temporal relationship between the excitation event and the observed pressure waveform. This smearing varies across vowels — close rounded vowels like /o/ and /u/ with their narrow first-formant bandwidths are empirically the most challenging cases for GIF (Chien et al., 2017) — and there is no single “correct” delay that compensates it.

The standard evaluation benchmark of Chien et al. (2017) accounts for this with a fixed 0.65 ms delay applied uniformly to all ground-truth waveforms, motivated by an assumed 22 cm radiation distance. This is a reasonable first-order correction, but it is a single number applied to a phenomenon that varies with vowel, pitch, voice quality, and fundamental frequency. More importantly, it corrects for the mean delay of one physical effect while leaving the gauge ambiguity from the source-filter decomposition entirely unaddressed. Applying a fixed shift does not constitute gauge-fixing; it shifts the entire evaluation into a particular gauge that may match the CPCA convention reasonably well but will systematically misalign with CCD estimates.

G.3. Per-cycle alignment is the wrong granularity

The same benchmark evaluates performance in a pitch-synchronous, per-cycle fashion: ground-truth and estimated waveforms are aligned within each glottal cycle separately, using glottal closure instants extracted from the synthesizer. This procedure is motivated by the desire to isolate the pulse shape from pitch-period-to-pitch-period variability.

The problem is that it introduces a reference frame — the glottal cycle — that is not implied by the generative model. The source-filter interaction in (303) occurs at frame scale: the filter $h(t)$ is estimated from a windowed segment spanning many cycles, and the gauge it selects is therefore a frame-level property, not a cycle-level one. Different LPC-based algorithms apply their weighting functions (the closed-phase window in CPCA, the Gaussian weights around glottal closure in SLP, the piecewise-linear weights in WLP) at frame scale, and the effective delay they induce drifts slowly across frames as the filter estimate adapts. Cycle-level alignment obscures this drift by correcting it separately at each cycle, which has the

unintended effect of partially correcting for LPC-induced gauge variation on behalf of the algorithm.

Frame-level alignment is the granularity that matches the statistical assumptions of the model and the practical structure of the algorithms.

G.4. Affine alignment at frame level

The physically motivated evaluation procedure aligns each estimated frame to the corresponding ground-truth frame by fitting an affine map and a nonnegative integer lag jointly,

$$(a^*, b^*, \tau^*) = \underset{a, b, \tau}{\operatorname{argmin}} \| \mathbf{y}_{\text{est}[\tau:]} \cdot a + b - \mathbf{y}_{\text{true}[:n-\tau]} \|^2, \quad (306)$$

where $\tau \in \{0, 1, \dots, \tau_{\max}\}$ is a nonnegative integer lag measured in samples.³⁸ The scalar a absorbs scale and polarity; the scalar b absorbs any mean offset introduced by pre-emphasis or measurement convention. The lag τ handles the frame-level delay that varies across algorithms and physical conditions.

The aligned estimate is

$$\tilde{\mathbf{y}}_{\text{est}} = a^* \mathbf{y}_{\text{est}[\tau^*]} + b^*, \quad (307)$$

defined on the overlap $[:n - \tau^*]$. Performance is then measured by centered cosine similarity between $\tilde{\mathbf{y}}_{\text{est}}$ and $\mathbf{y}_{\text{true}[:n-\tau^*]}$:

$$\text{score} = \frac{\langle \tilde{\mathbf{y}}_{\text{est}} - \overline{\tilde{\mathbf{y}}_{\text{est}}}, \mathbf{y}_{\text{true}} - \overline{\mathbf{y}_{\text{true}}} \rangle}{\| \tilde{\mathbf{y}}_{\text{est}} - \overline{\tilde{\mathbf{y}}_{\text{est}}} \| \| \mathbf{y}_{\text{true}} - \overline{\mathbf{y}_{\text{true}}} \|}. \quad (308)$$

Centering before taking the inner product removes the b degree of freedom redundantly, making the score invariant to the bias correction. Cosine similarity is a monotone function of RMSE after affine alignment and is bounded in $[-1, 1]$, which makes it convenient for visualization and for averaging across frames of different lengths.

The procedure is summarized in Table 10, alongside the corresponding treatment in the EGIFA benchmark (Chien et al., 2017) for comparison.

G.5. Summary

The source-filter model has a gauge group: scale and time-translation transformations that leave the observed signal invariant and therefore make the corresponding source properties unobservable from acoustics alone. Different inverse filtering algorithms choose different representatives from this group — LP-based methods anchor on closed-phase energy, CCD

³⁸The lag is constrained to be nonnegative and to shift the estimate earlier in time, reflecting the physical direction of propagation delay: the acoustic observation $y(t)$ always lags the source, so a positive lag on the estimate corresponds to shifting it left to match the ground-truth timing. Allowing both directions would conflate physical delay correction with arbitrary time-warping.

Gauge source	Origin	EGIFA treatment	This work
Scale / polarity	Source-filter non-identifiability	Orthogonal projection per cycle	Affine fit per frame (306)
Absolute delay	Source-filter non-identifiability	Fixed 0.65 ms shift	Lag search per frame (306)
Dispersive group delay	Vocal tract resonances	Fixed 0.65 ms shift (partial)	Lag search per frame
Algorithmic gauge	LP vs. cepstral estimation	Not corrected	Corrected uniformly
Mean offset	Pre-emphasis measurement	/ Not corrected	Affine fit (b term)

Table 10 | **Gauge sources in GIF evaluation** and their treatment in the EGIFA benchmark of Chien et al. (2017) versus this work. The EGIFA fixed delay of 0.65 ms addresses one physical effect; it does not address the gauge ambiguity from the source-filter decomposition or the algorithm-dependent gauge drift. Affine alignment at frame level handles all sources uniformly.

anchors on cepstral phase — and any evaluation that compares representatives directly rather than correcting for this choice conflates algorithmic style with algorithmic error.

The same argument applies to the fixed-delay and per-cycle alignment choices in the EGIFA benchmark (Chien et al., 2017). A fixed 0.65 ms delay corrects for one specific physical effect at one specific propagation distance; it does not address the gauge freedom of the model, and it does not adapt to the algorithm being evaluated. Per-cycle alignment introduces a reference frame that the generative model does not privilege and that partially absorbs algorithm-induced gauge drift.

Frame-level affine alignment with a lag search, followed by centered cosine similarity, addresses all of these uniformly. It is the evaluation procedure used throughout the experimental chapters of this thesis.

Appendix H. Related contributions

Van Soom & de Boer (2019) Measuring formant frequency $\mathbf{F}$ and bandwidth $\mathbf{B}$ from steady-state vowel waveforms, taking into account the low-frequency effects of the DGF $u'(t)$ during the open phase of the pitch period. Inference is implemented using a maximum-a-posteriori (MAP) approximation. In BNGIF these low-frequency effects are taken into account automatically.

Van Soom & de Boer (2020) Same as the above, but now using robust nested sampling inference instead of a MAP approximation. A new heuristic derivation of the model function from source-filter theory is added. In BNGIF variational inference is the preferred inference method.

Van Soom & De Boer (2022) Theoretical derivation of a new weakly informative prior for resonance frequencies. Shown to have superior performance in predicting steady-state vowel waveforms compared to the usual priors in inferring the pole frequency $\mathbf{x}$ and bandwidth $\mathbf{y}$, while being roughly equally (un)informative.